

Heaping on the Complexity

Adventures in Ruby's GC Compactor



Matt Valentine-House
Shopify - Ruby Infrastructure team
@eightbitraptor



Outline

- Background

Outline

- Background
- The Problem and Solutions

Outline

- Background
- The Problem and Solutions
- What happened next & Memory Allocation update

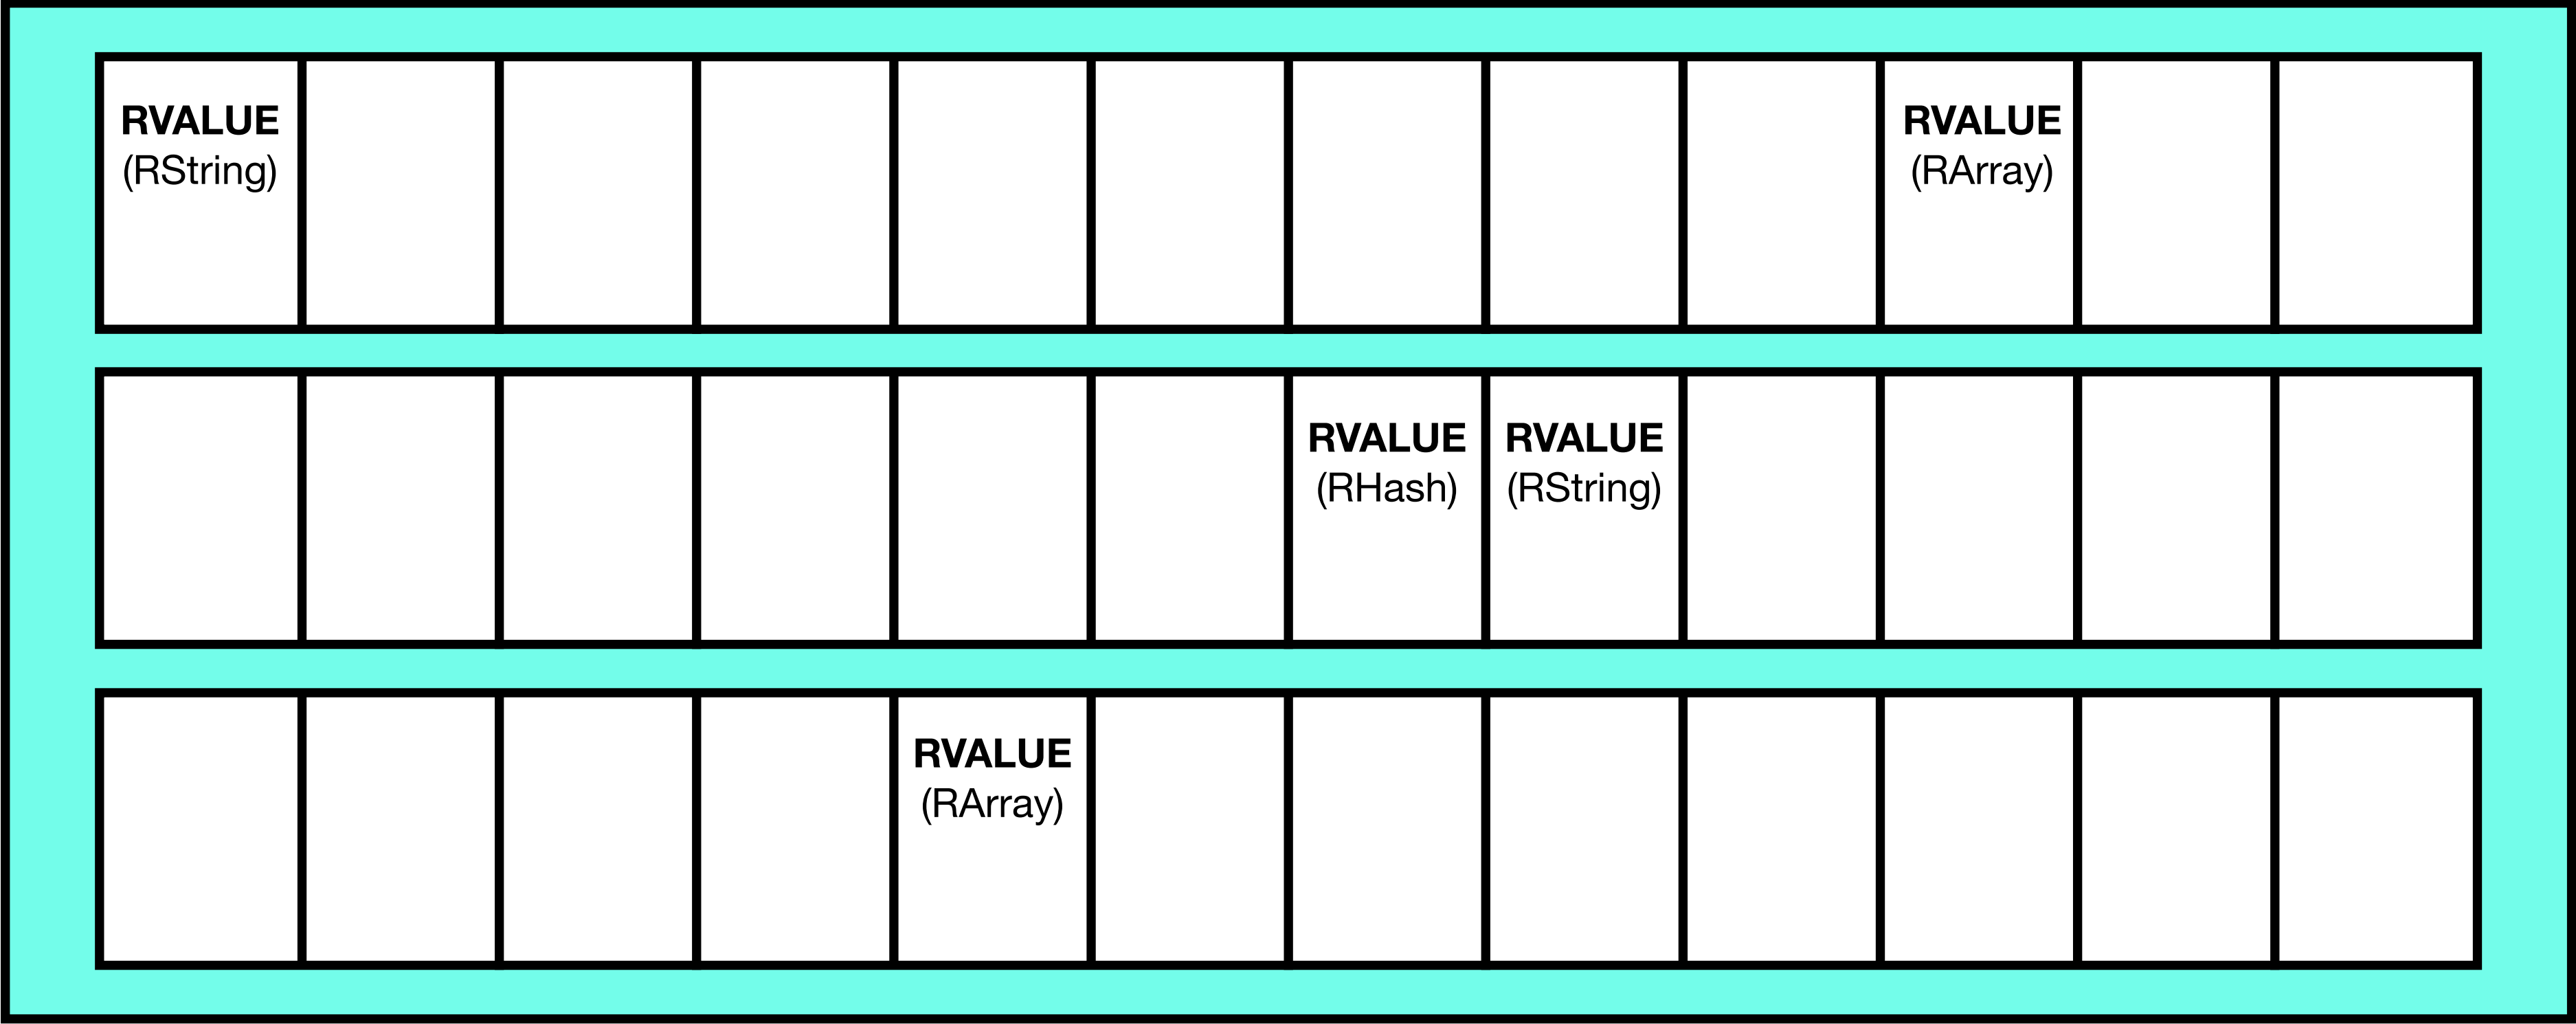
Outline

- Background
- The Problem and Solutions
- What happened next?
- Introduce Heapviz

Background

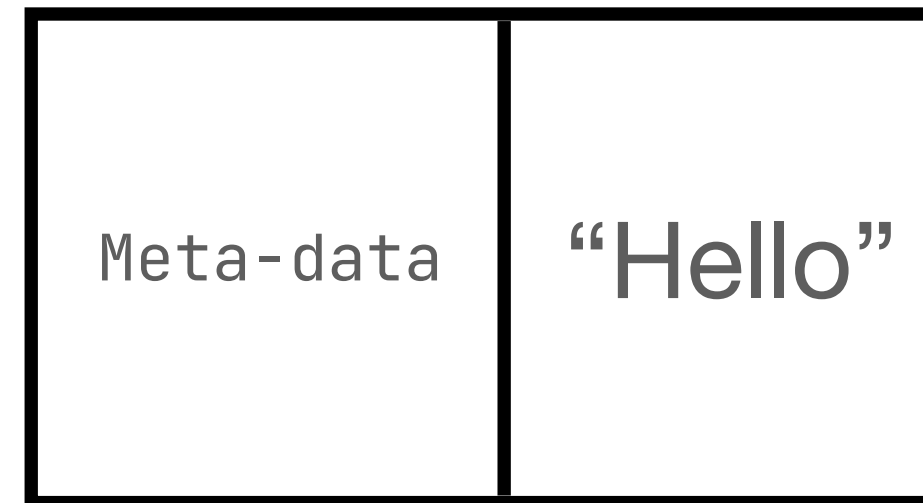
Memory

Background - The Heap

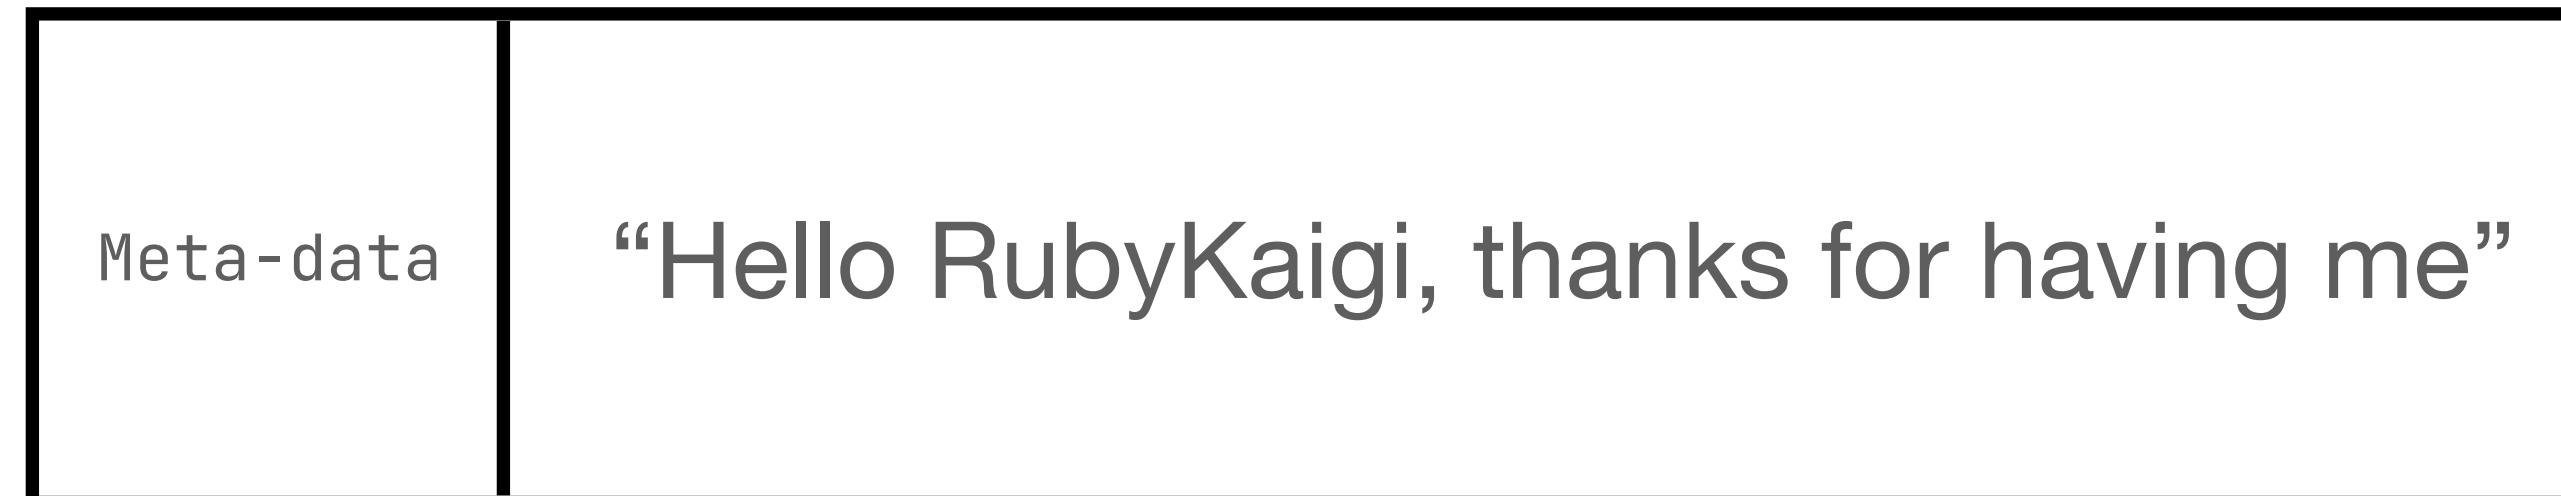


Background - Heap & Embedded Allocation

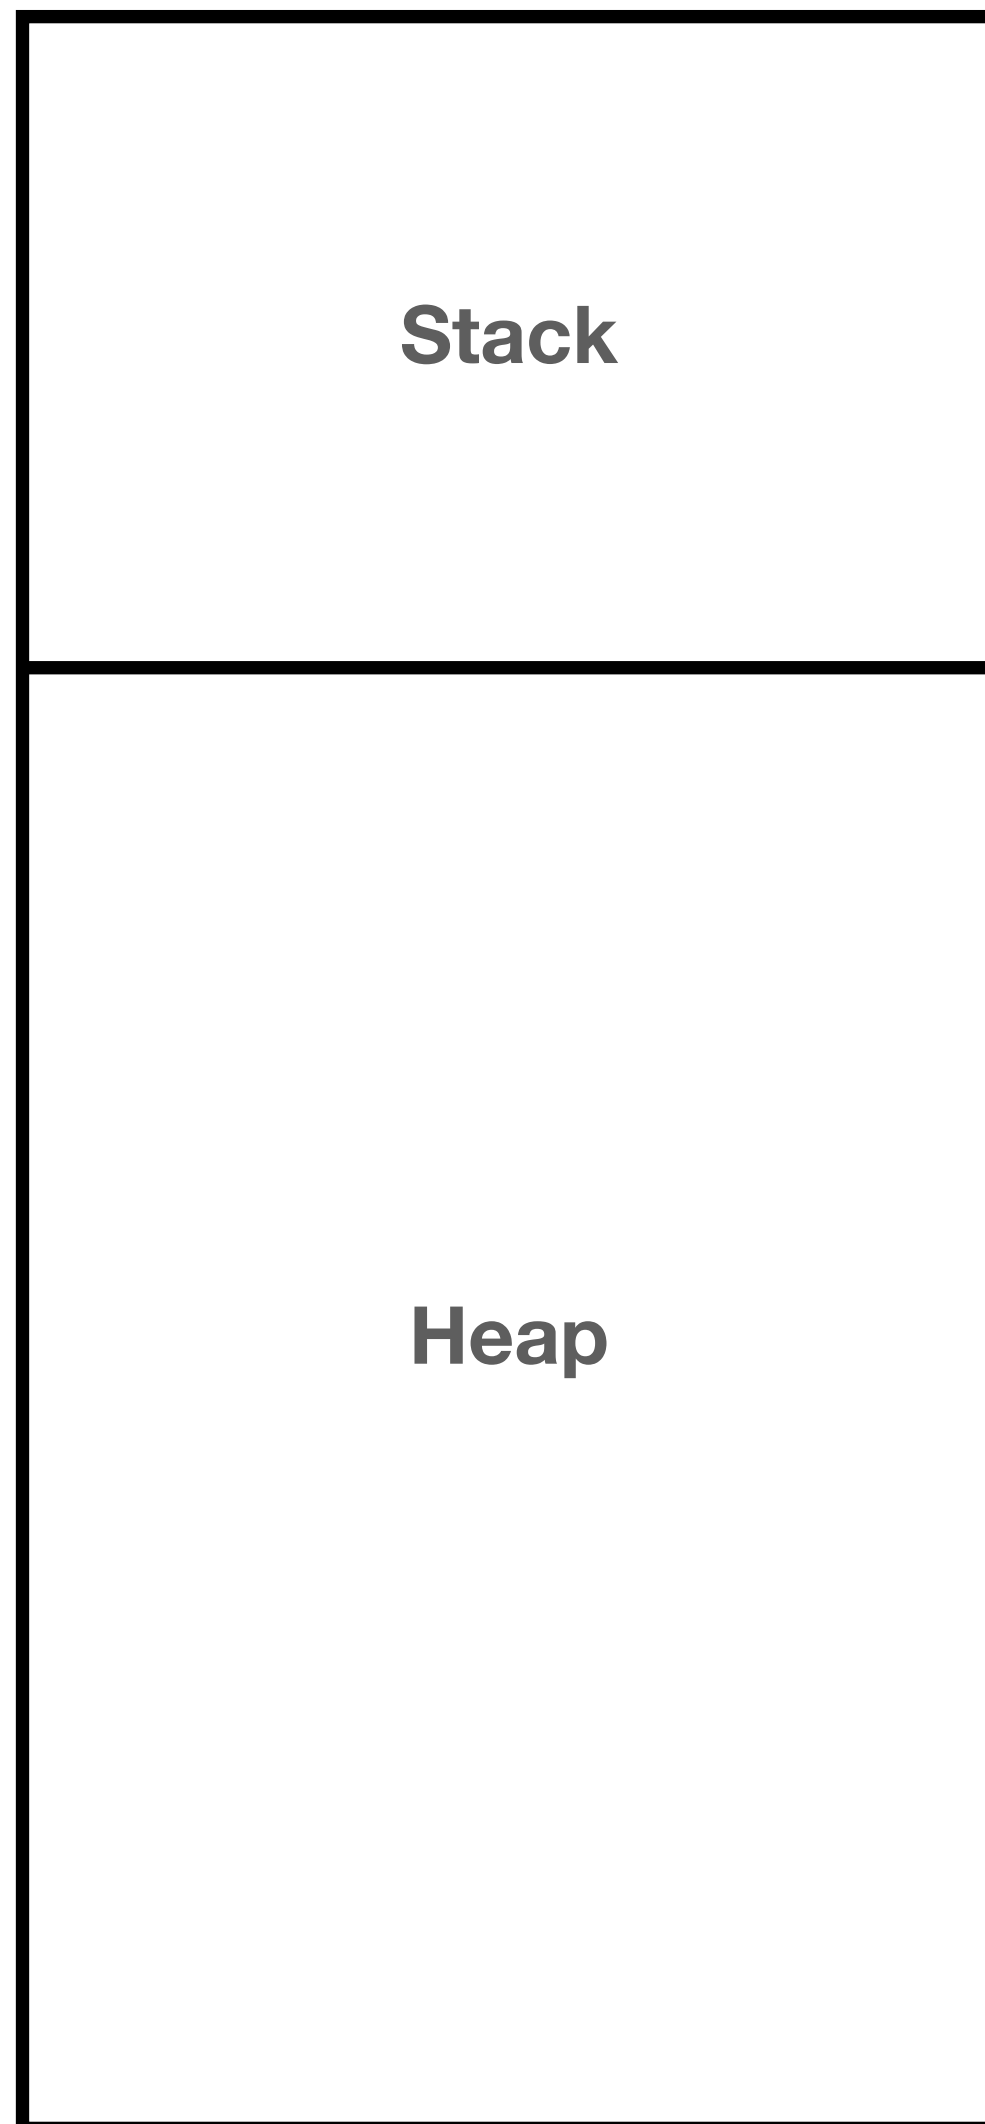
RString



RString



Background - Memory, an OS perspective

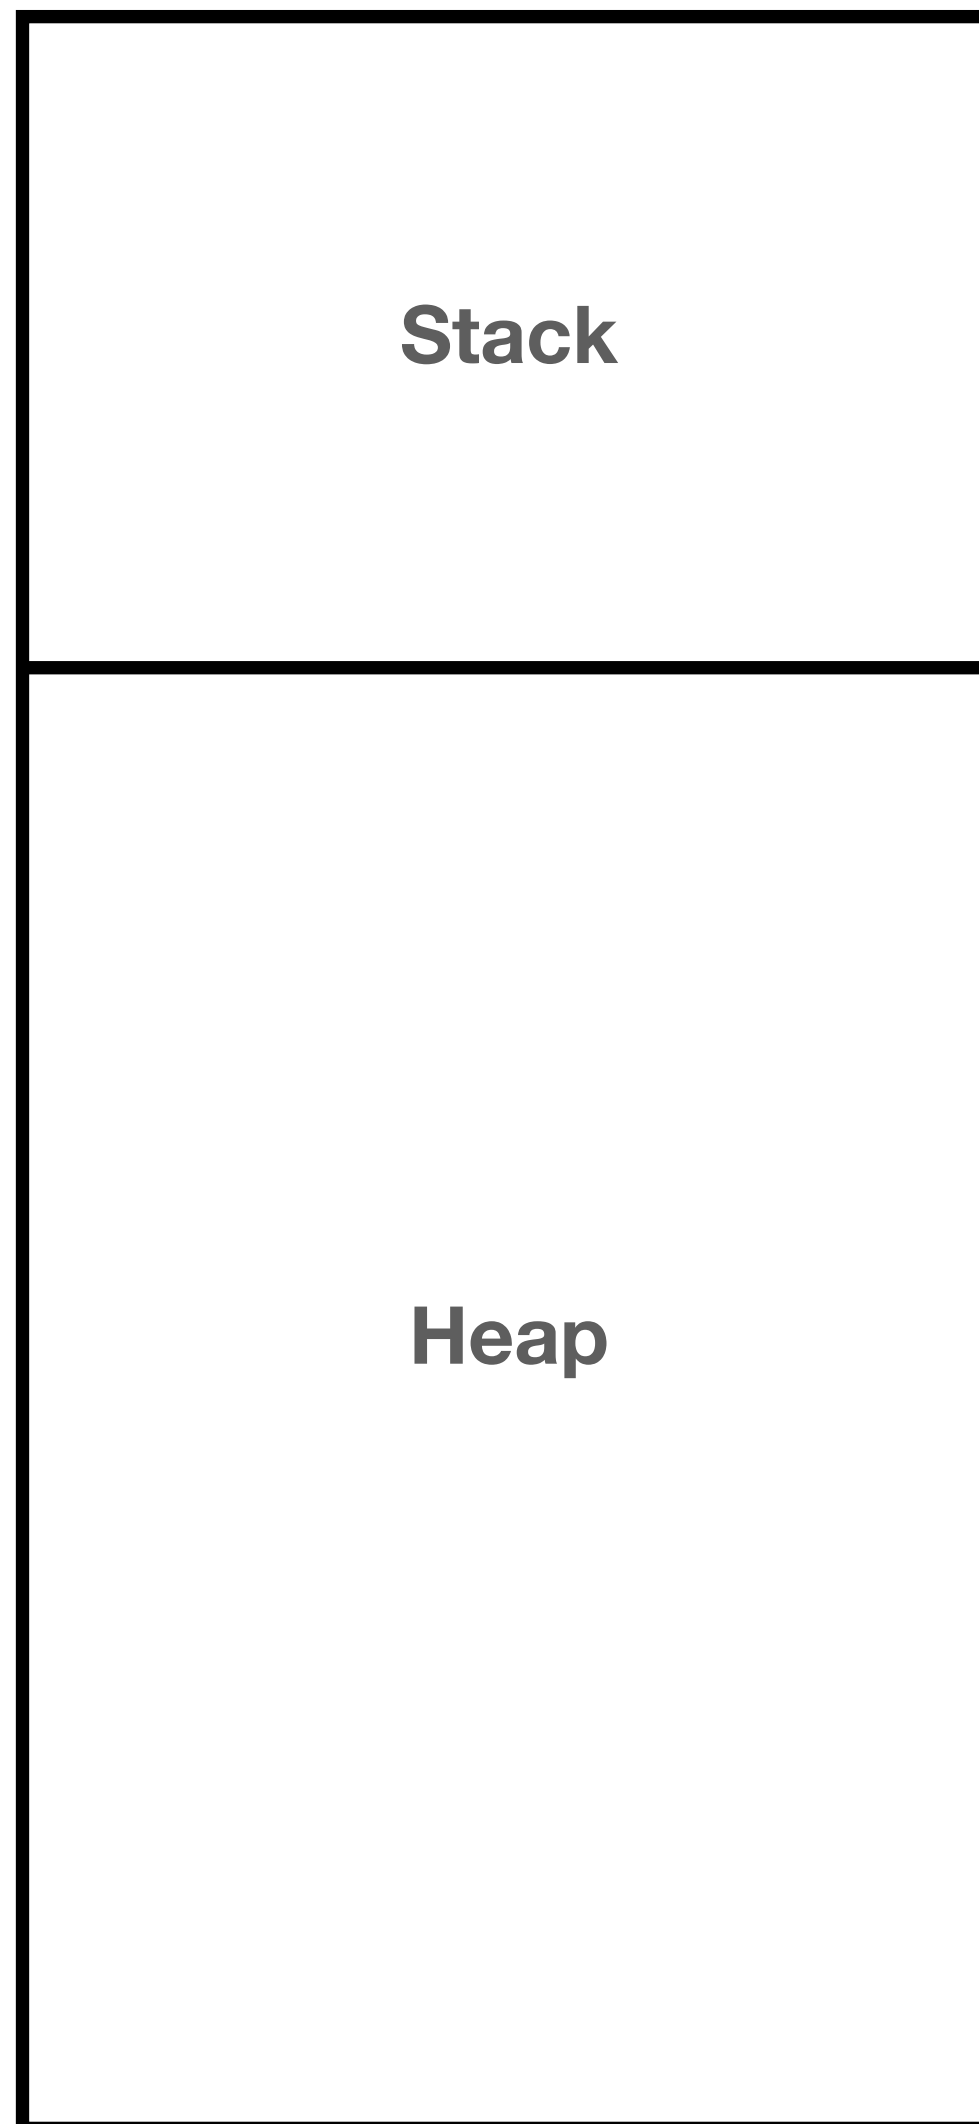


```
#include <stdlib.h>
```

```
int
main(int argc, char *argv)
{

}
```

Background - Memory, an OS perspective



```
#include <stdlib.h>
```

```
int
```

```
main(int argc, char *argv)
```

```
{
```

```
    // allocate space for a string of 100 characters in the heap
```

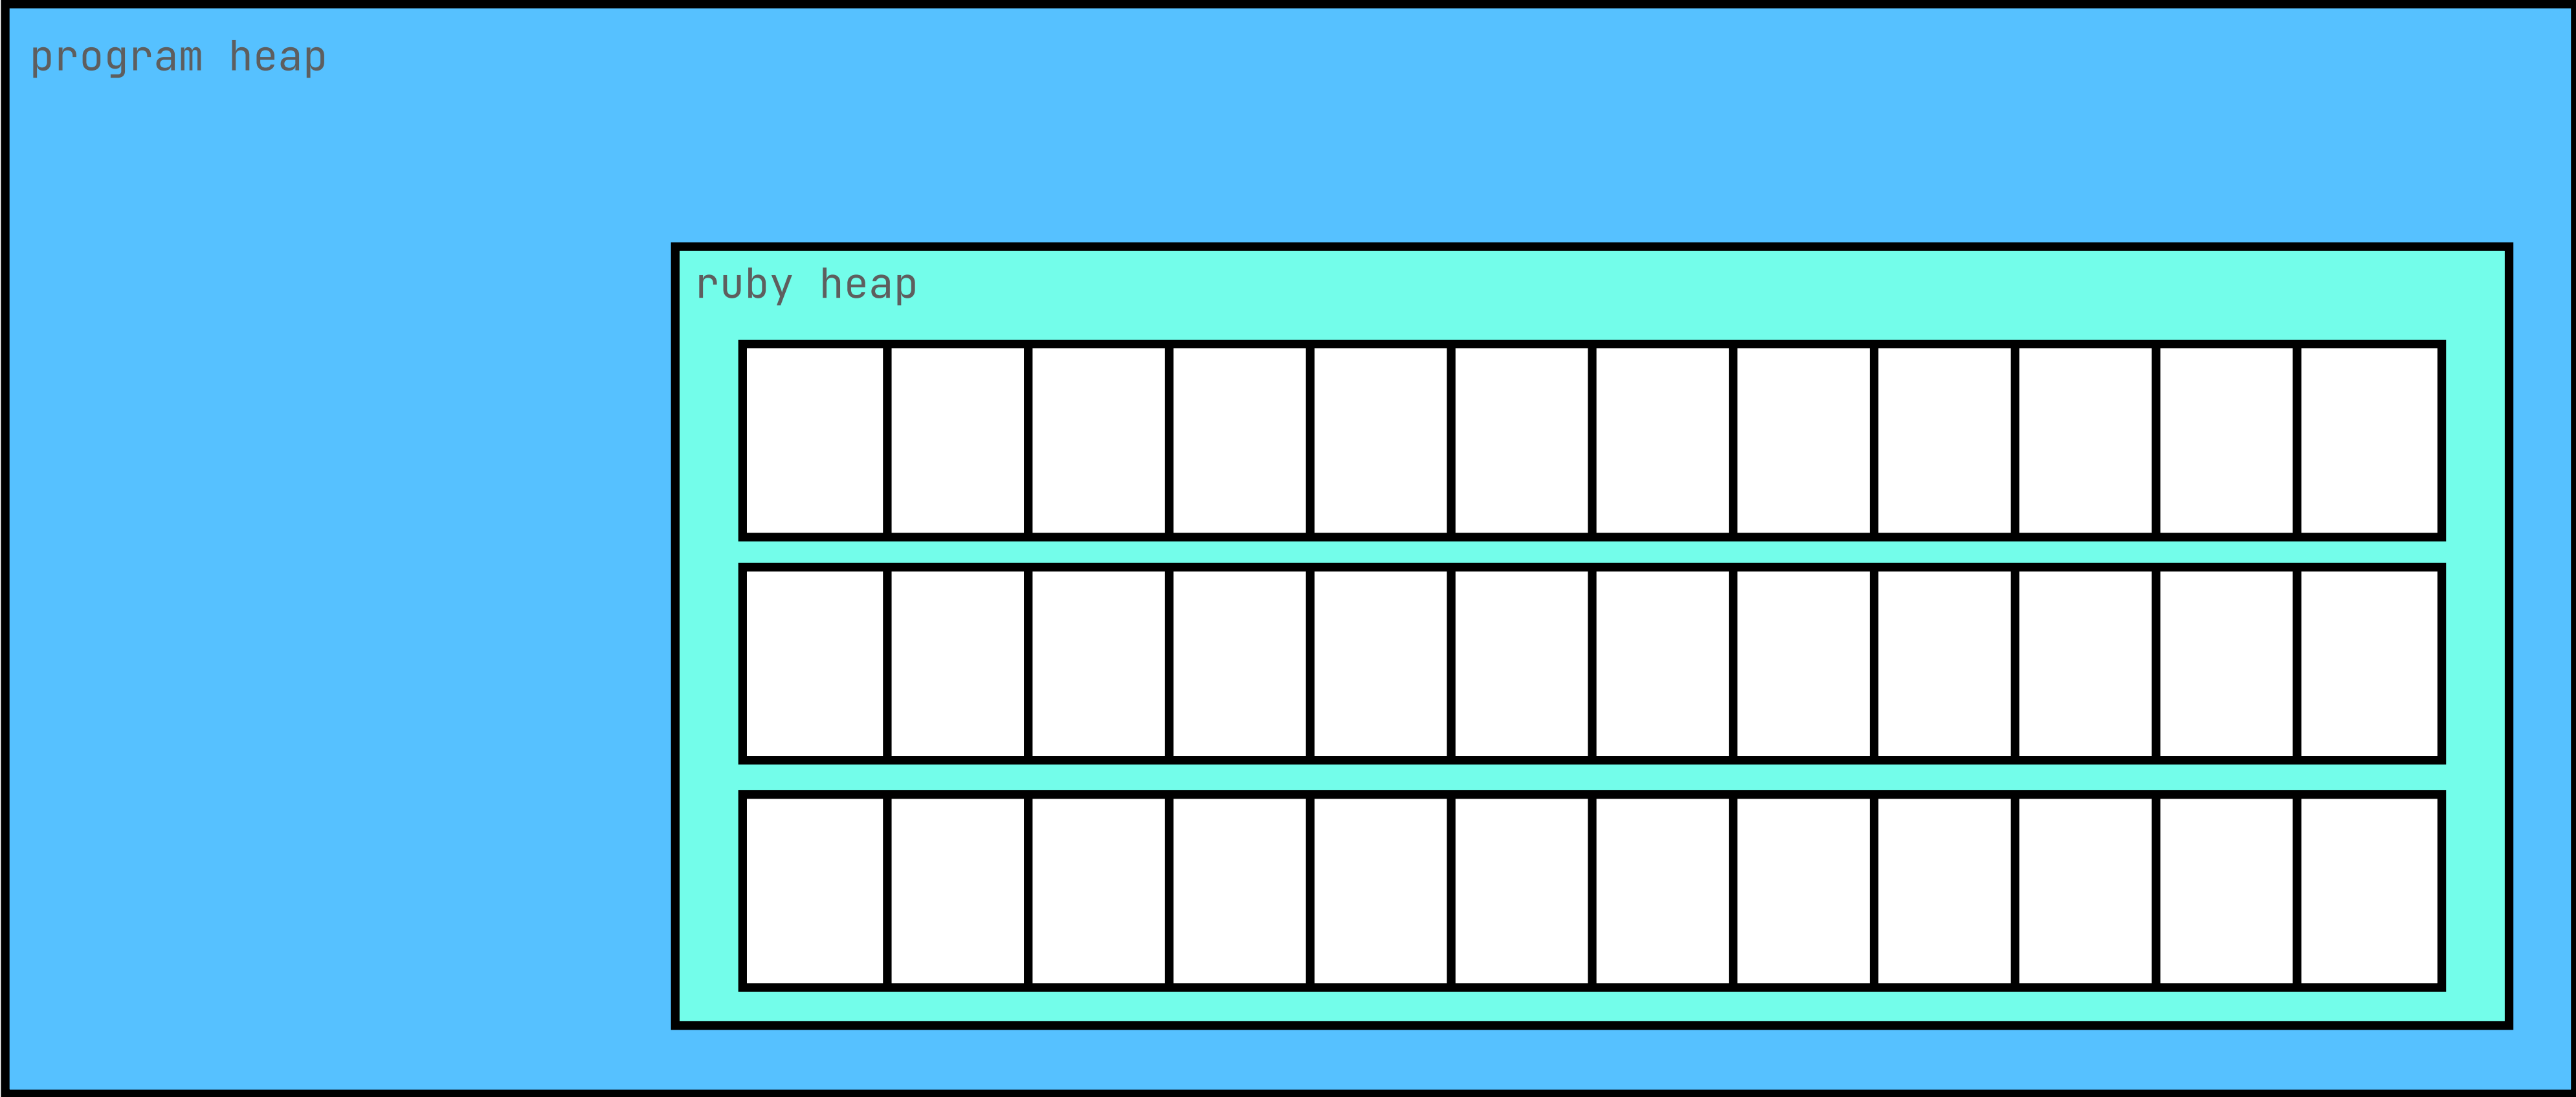
```
    char * dynamic_string = (char *)malloc(100);
```

```
    // release the memory back to the OS when we're done
```

```
    free(dynamic_string);
```

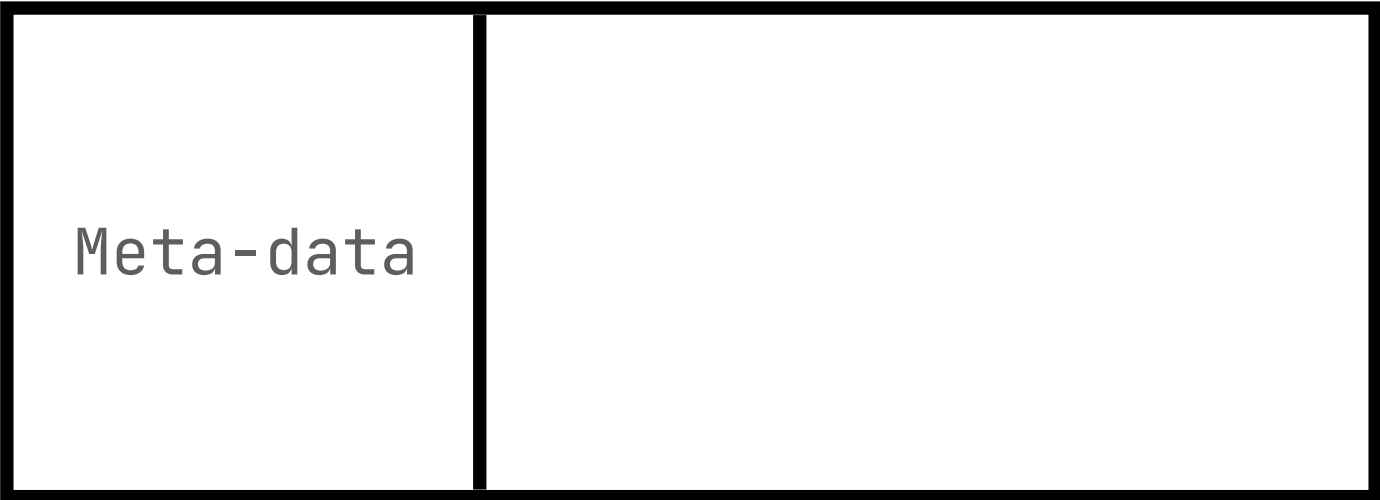
```
}
```

Background - The Ruby Heap



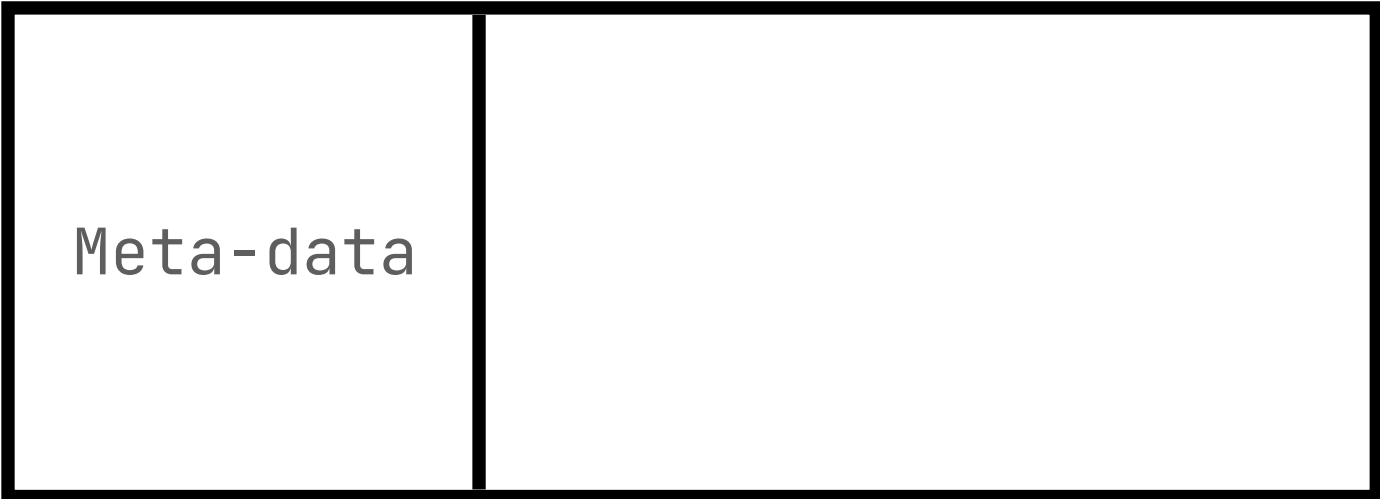
Background - Heap & Embedded Allocation

RString



“Hello”

RString



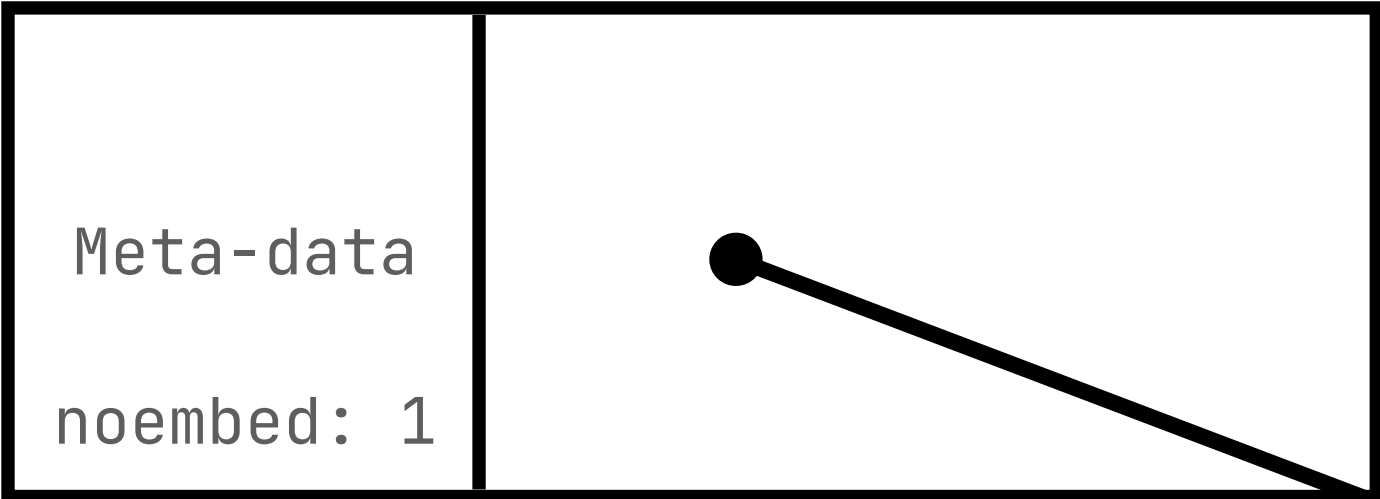
“Hello RubyKaigi, thanks for having me”

Background - Heap & Embedded Allocation

RString



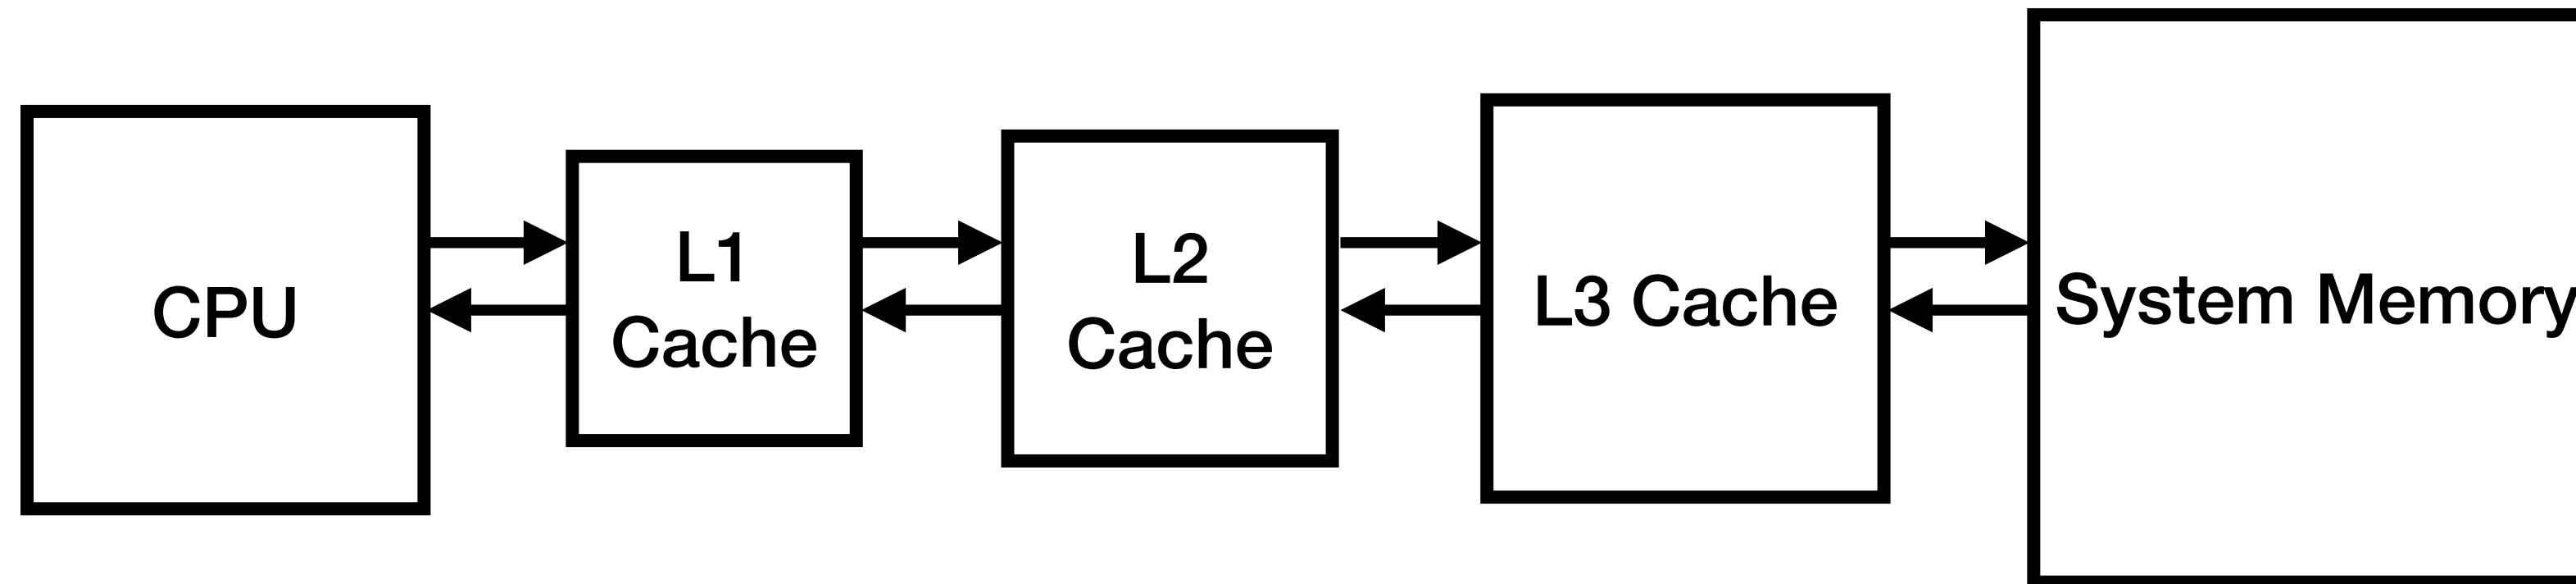
RString



"Hello RubyKaigi, thanks for having me"

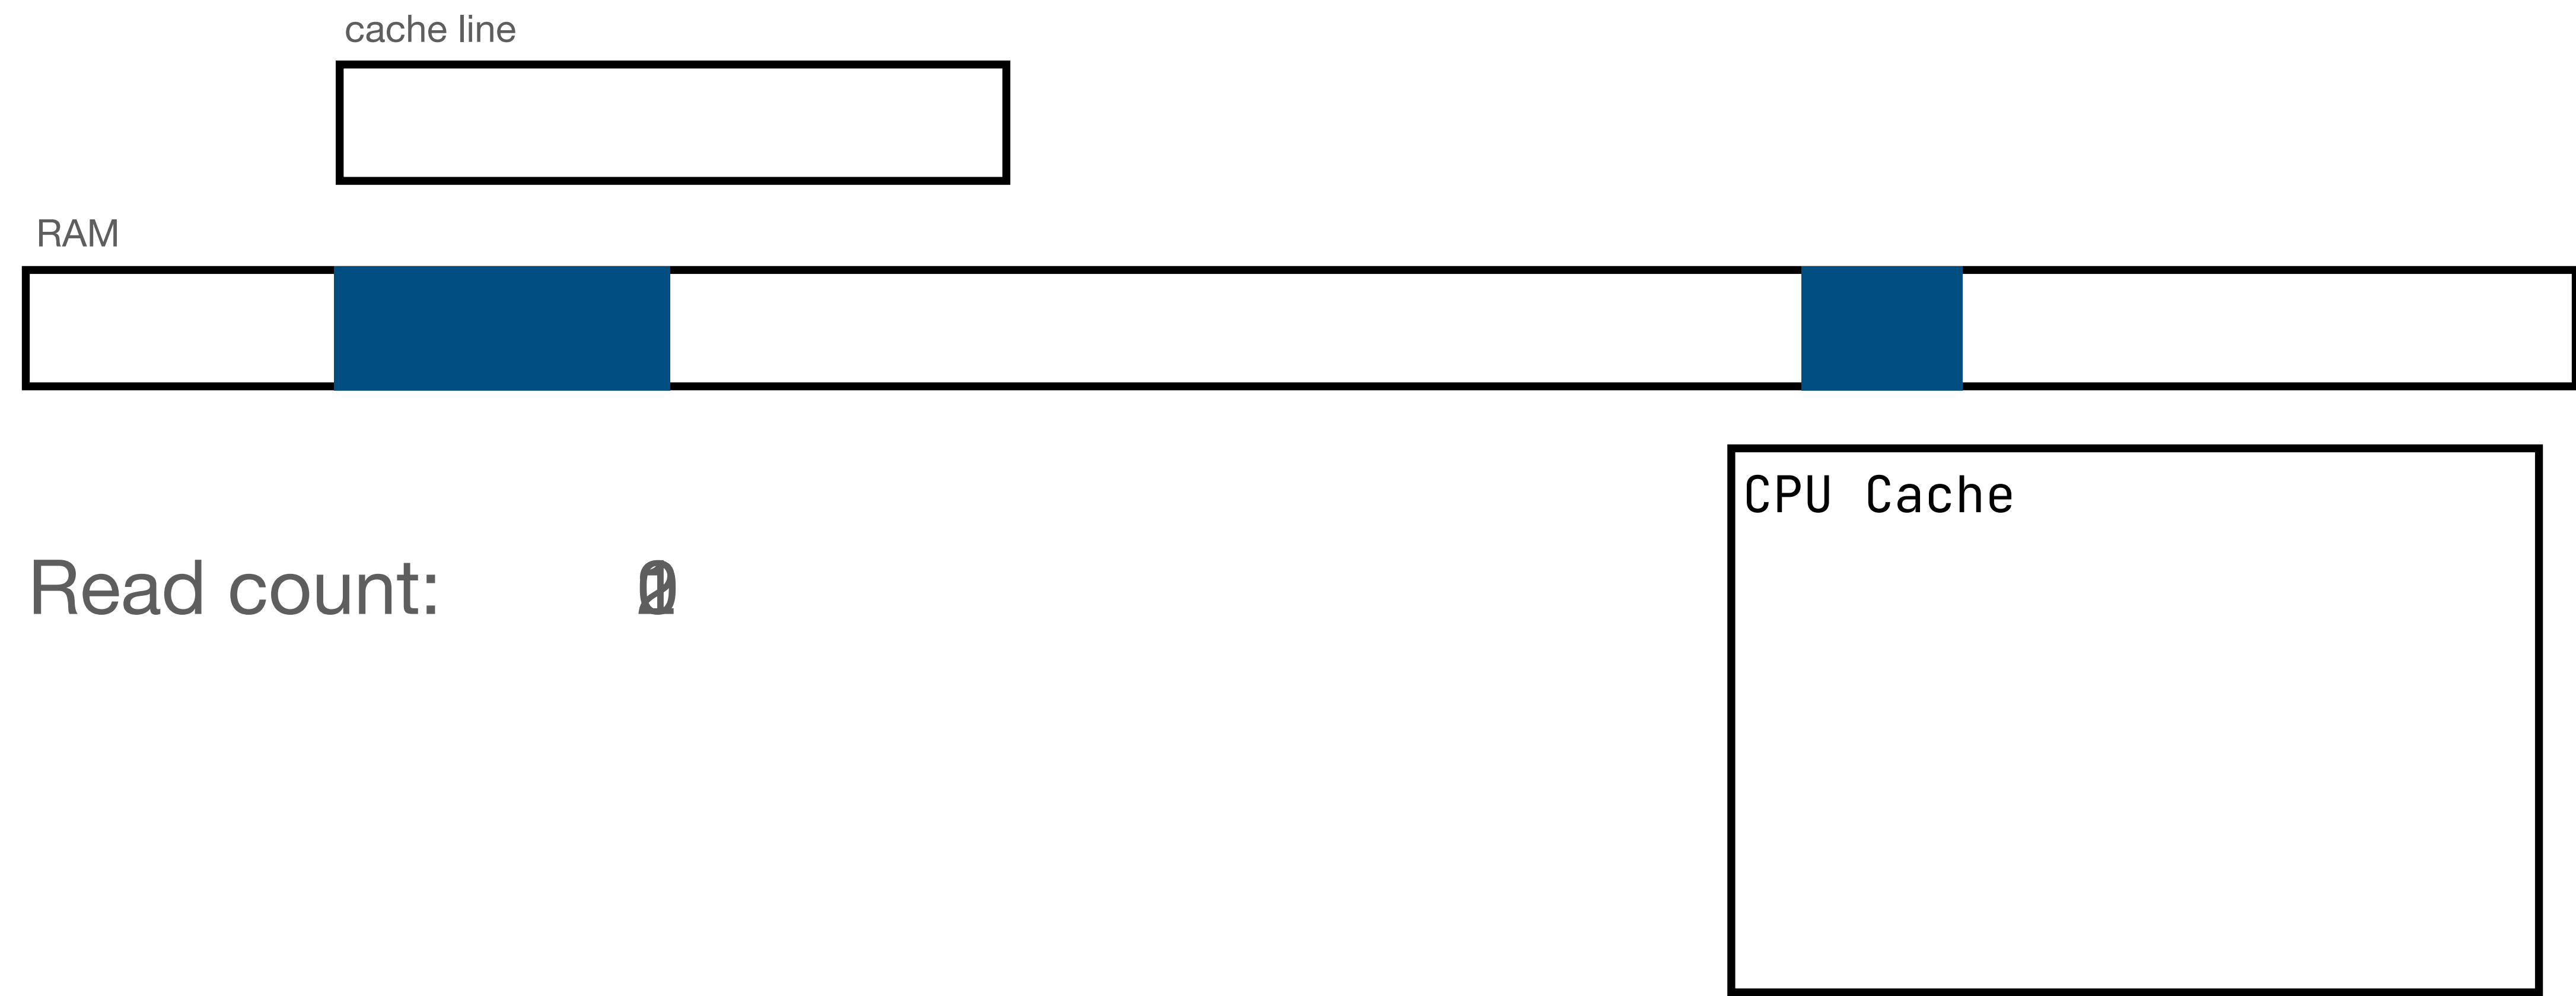


Background - CPU caches

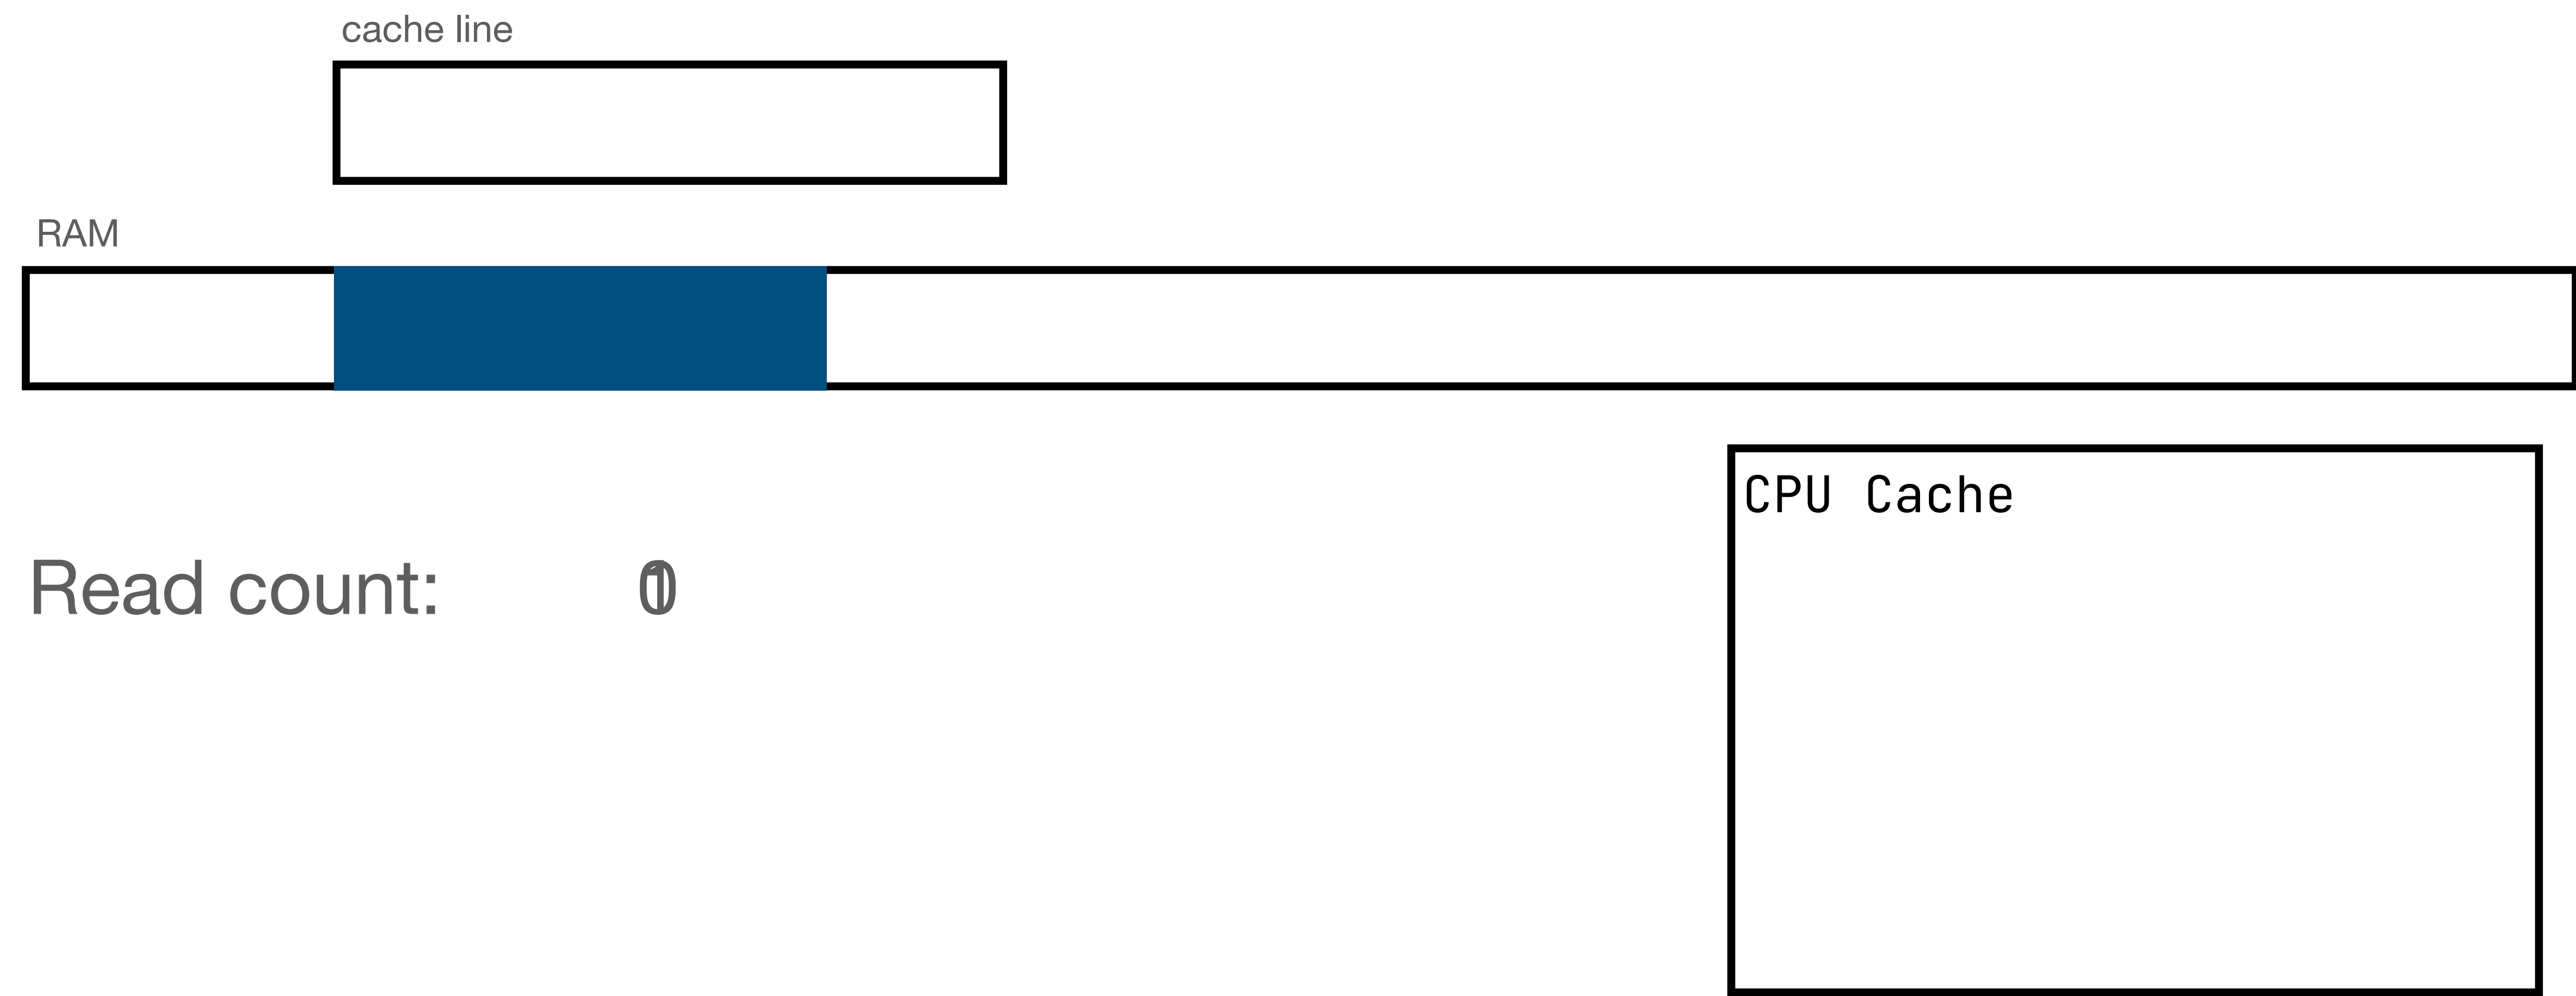


- **CPU's have multiple levels of cache**
- **The CPU cores read data from the closest cache**
- **A “cache miss” is when the data requested isn't in the desired cache.**
- **Poor cache performance can impact the performance of our programs**

Background - CPU Caches



Background - CPU Caches

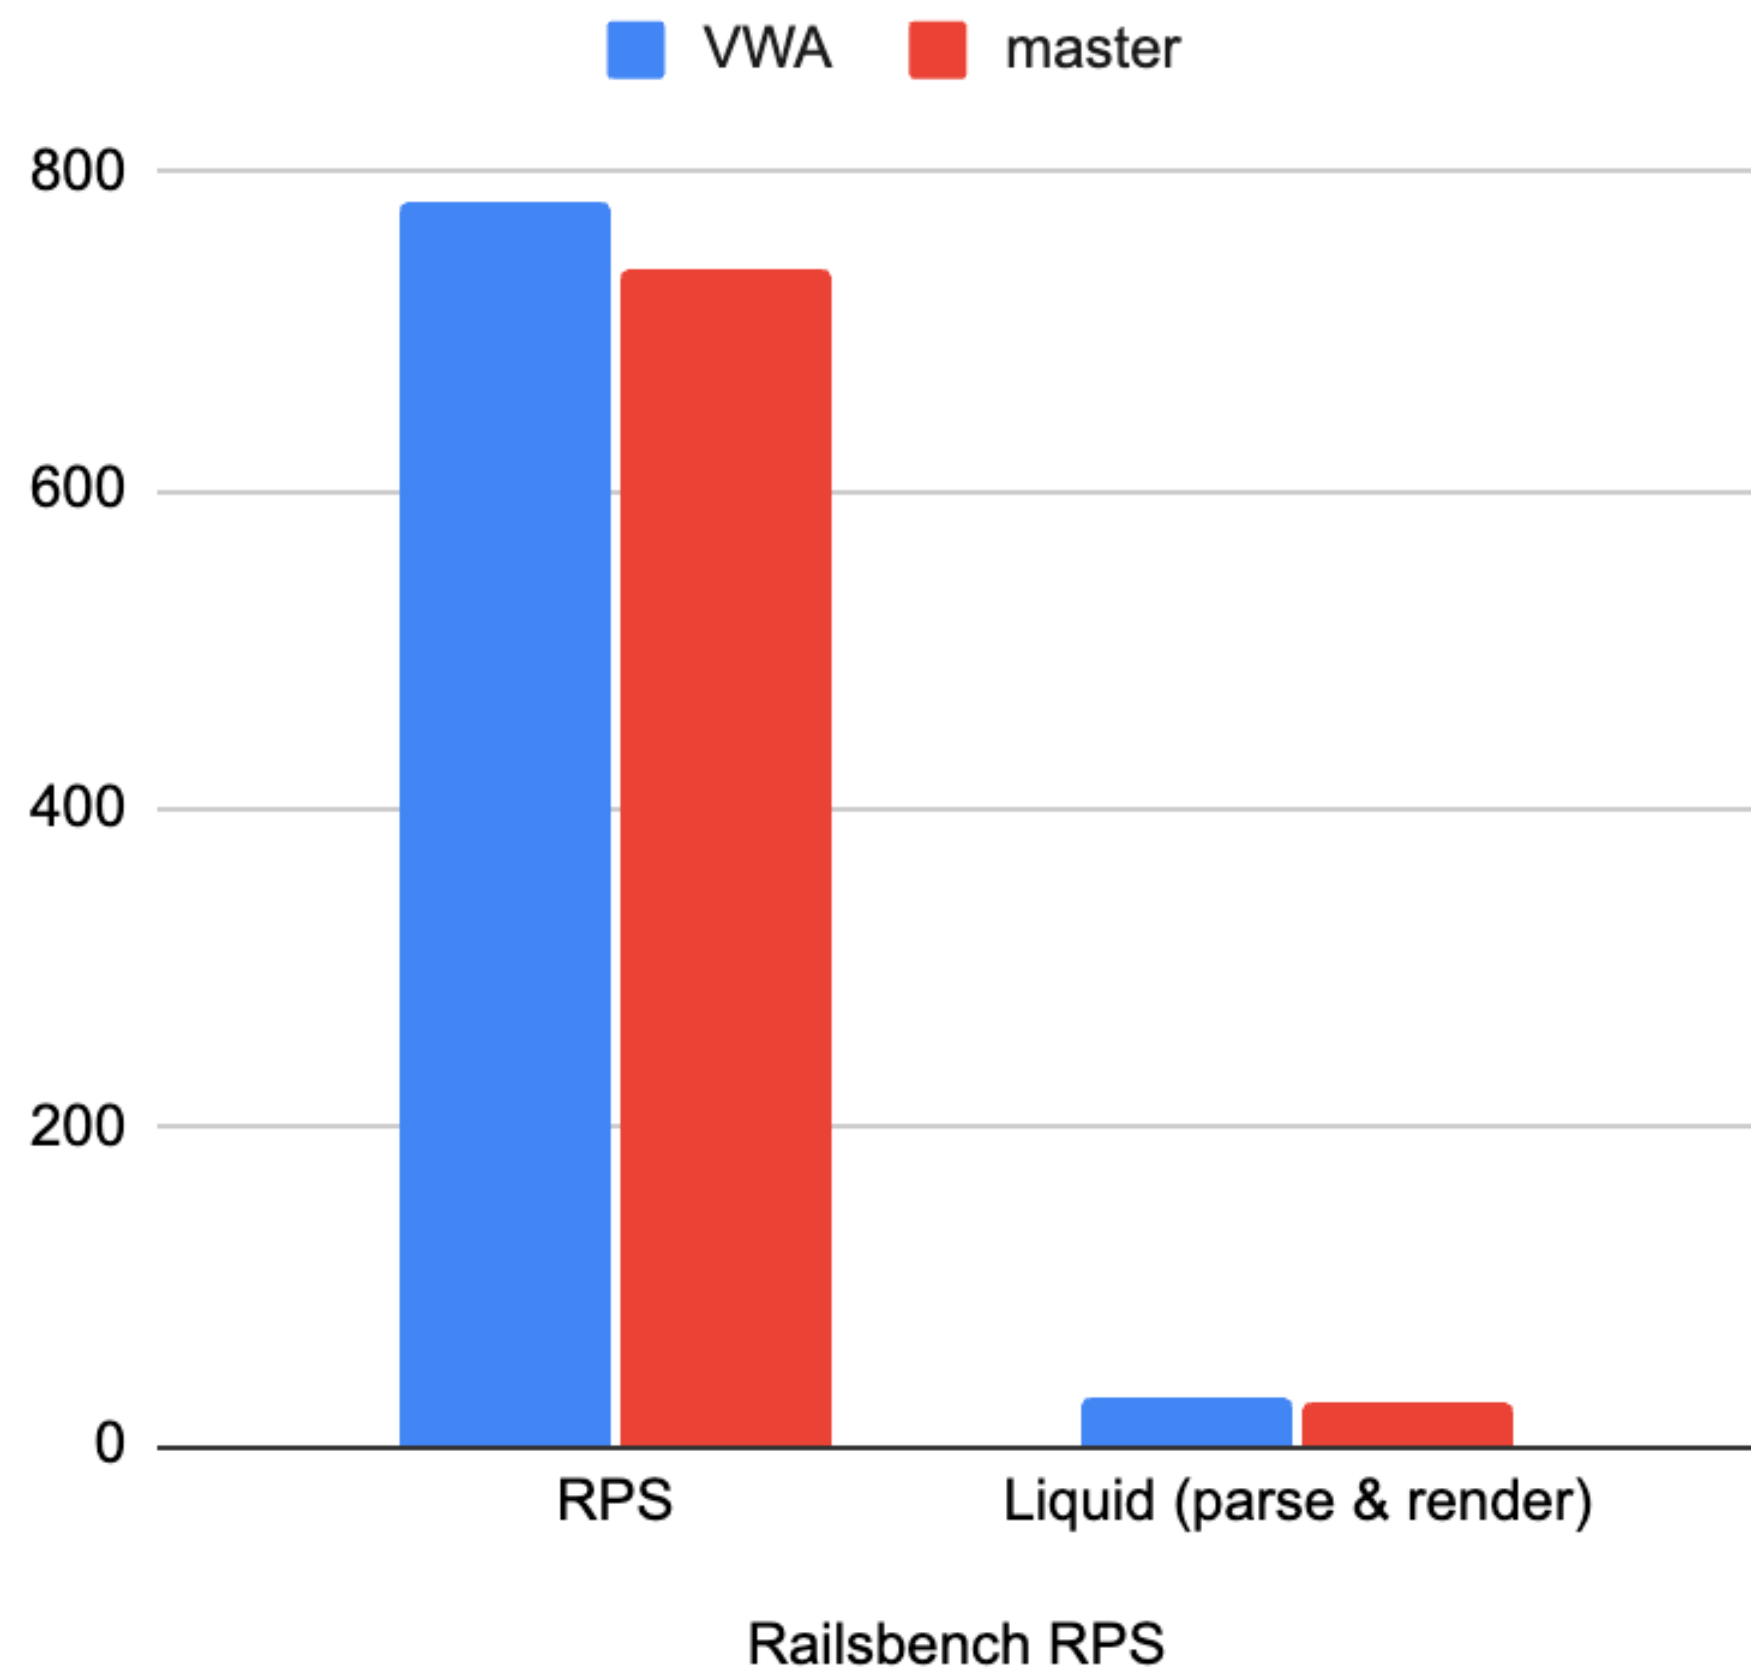


Background - CPU Caches

- **Heap allocated objects have poor locality.**
 - Objects with poor locality require **more reads when processing.**
 - Therefore poor locality has a **performance impact** on our programs.
- **Shopify are trying to solve this problem using Variable Width Allocation.**
 - VWA changes the heap structure to **enable more objects to be embedded.**
 - Embedded objects have **better data locality**, and are more **performant**

Background - Variable Width Allocation

Railsbench/Liquid - VWA vs master



VWA Results so far:

- **~5% improvement in Railsbench RPS**
- **~2-3% improvement in Liquid performance**

Background

Garbage Collection (briefly)

Background - Garbage Collection

“Garbage collection (GC) is a form of automatic memory management.

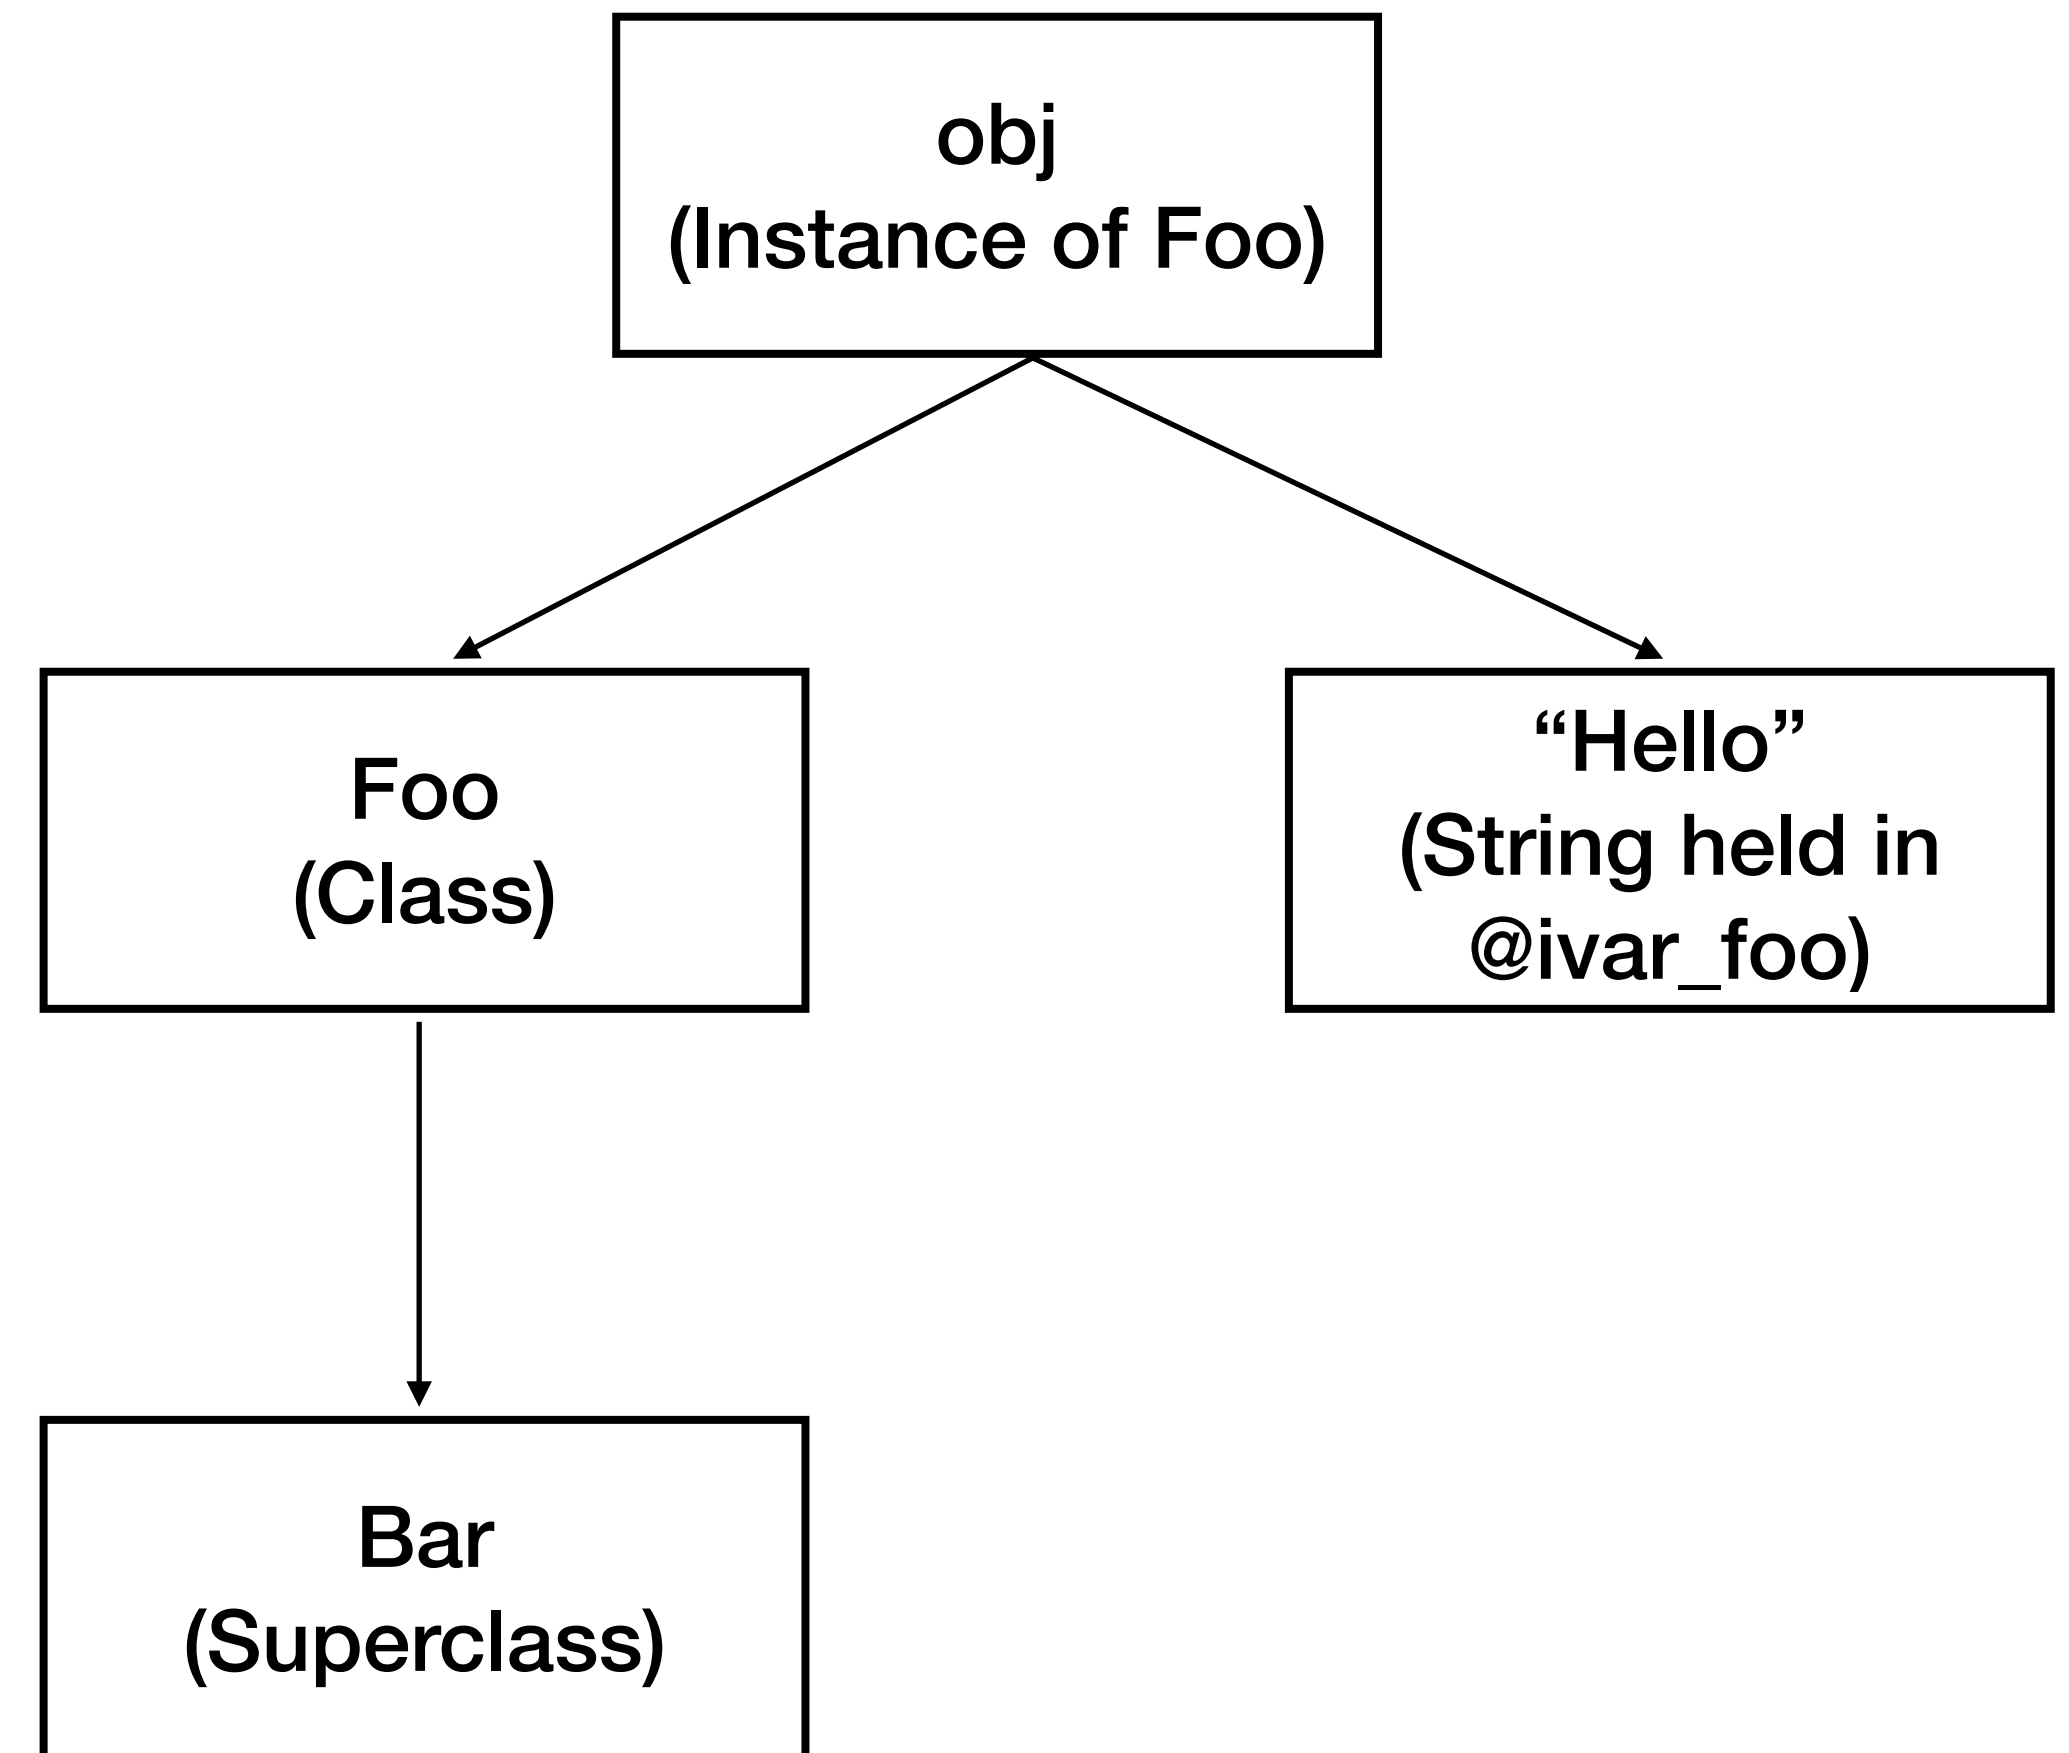
The garbage collector attempts to reclaim memory which was allocated by the program, but is no longer referenced”

—Wikipedia

Background - Garbage Collection

```
class Foo < Bar
  def initialize(my_foo)
    @ivar_foo = my_foo
  end
end
```

```
obj = Foo.new("Hello")
```



Background - Garbage Collection

Ruby's GC is:

- **Stop the World** - Execution of our program is paused while GC happens

Background - Garbage Collection

Ruby's GC is:

- **Stop the World** - Execution of our program is paused while GC happens
- **Mark-Sweep** - reachable objects are marked, unmarked objects are swept

Background - Garbage Collection

Ruby's GC is:

- **Stop the World** - Execution of our program is paused while GC happens
- **Mark-Sweep** - reachable objects are marked, unmarked objects are swept
- **Generational** - Young objects are collected more frequently than Old objects

Background - Garbage Collection

Ruby's GC is:

- **Stop the World** - Execution of our program is paused while GC happens
- **Mark-Sweep** - reachable objects are marked, unmarked objects are swept
- **Generational** - Young objects are collected more frequently than Old objects
- **Incremental** - GC pauses are spread over the runtime of the program

Background - Garbage Collection

Ruby's GC is:

- **Stop the World** - Execution of our program is paused while GC happens
- **Mark-Sweep** - reachable objects are marked, unmarked objects are swept
- **Generational** - Young objects are collected more frequently than Old objects
- **Incremental** - GC pauses are spread over the runtime of the program
- **Compacting** - Objects are moved closer together to avoid heap fragmentation

Background

GC Compaction (less briefly)

Background - GC Compaction

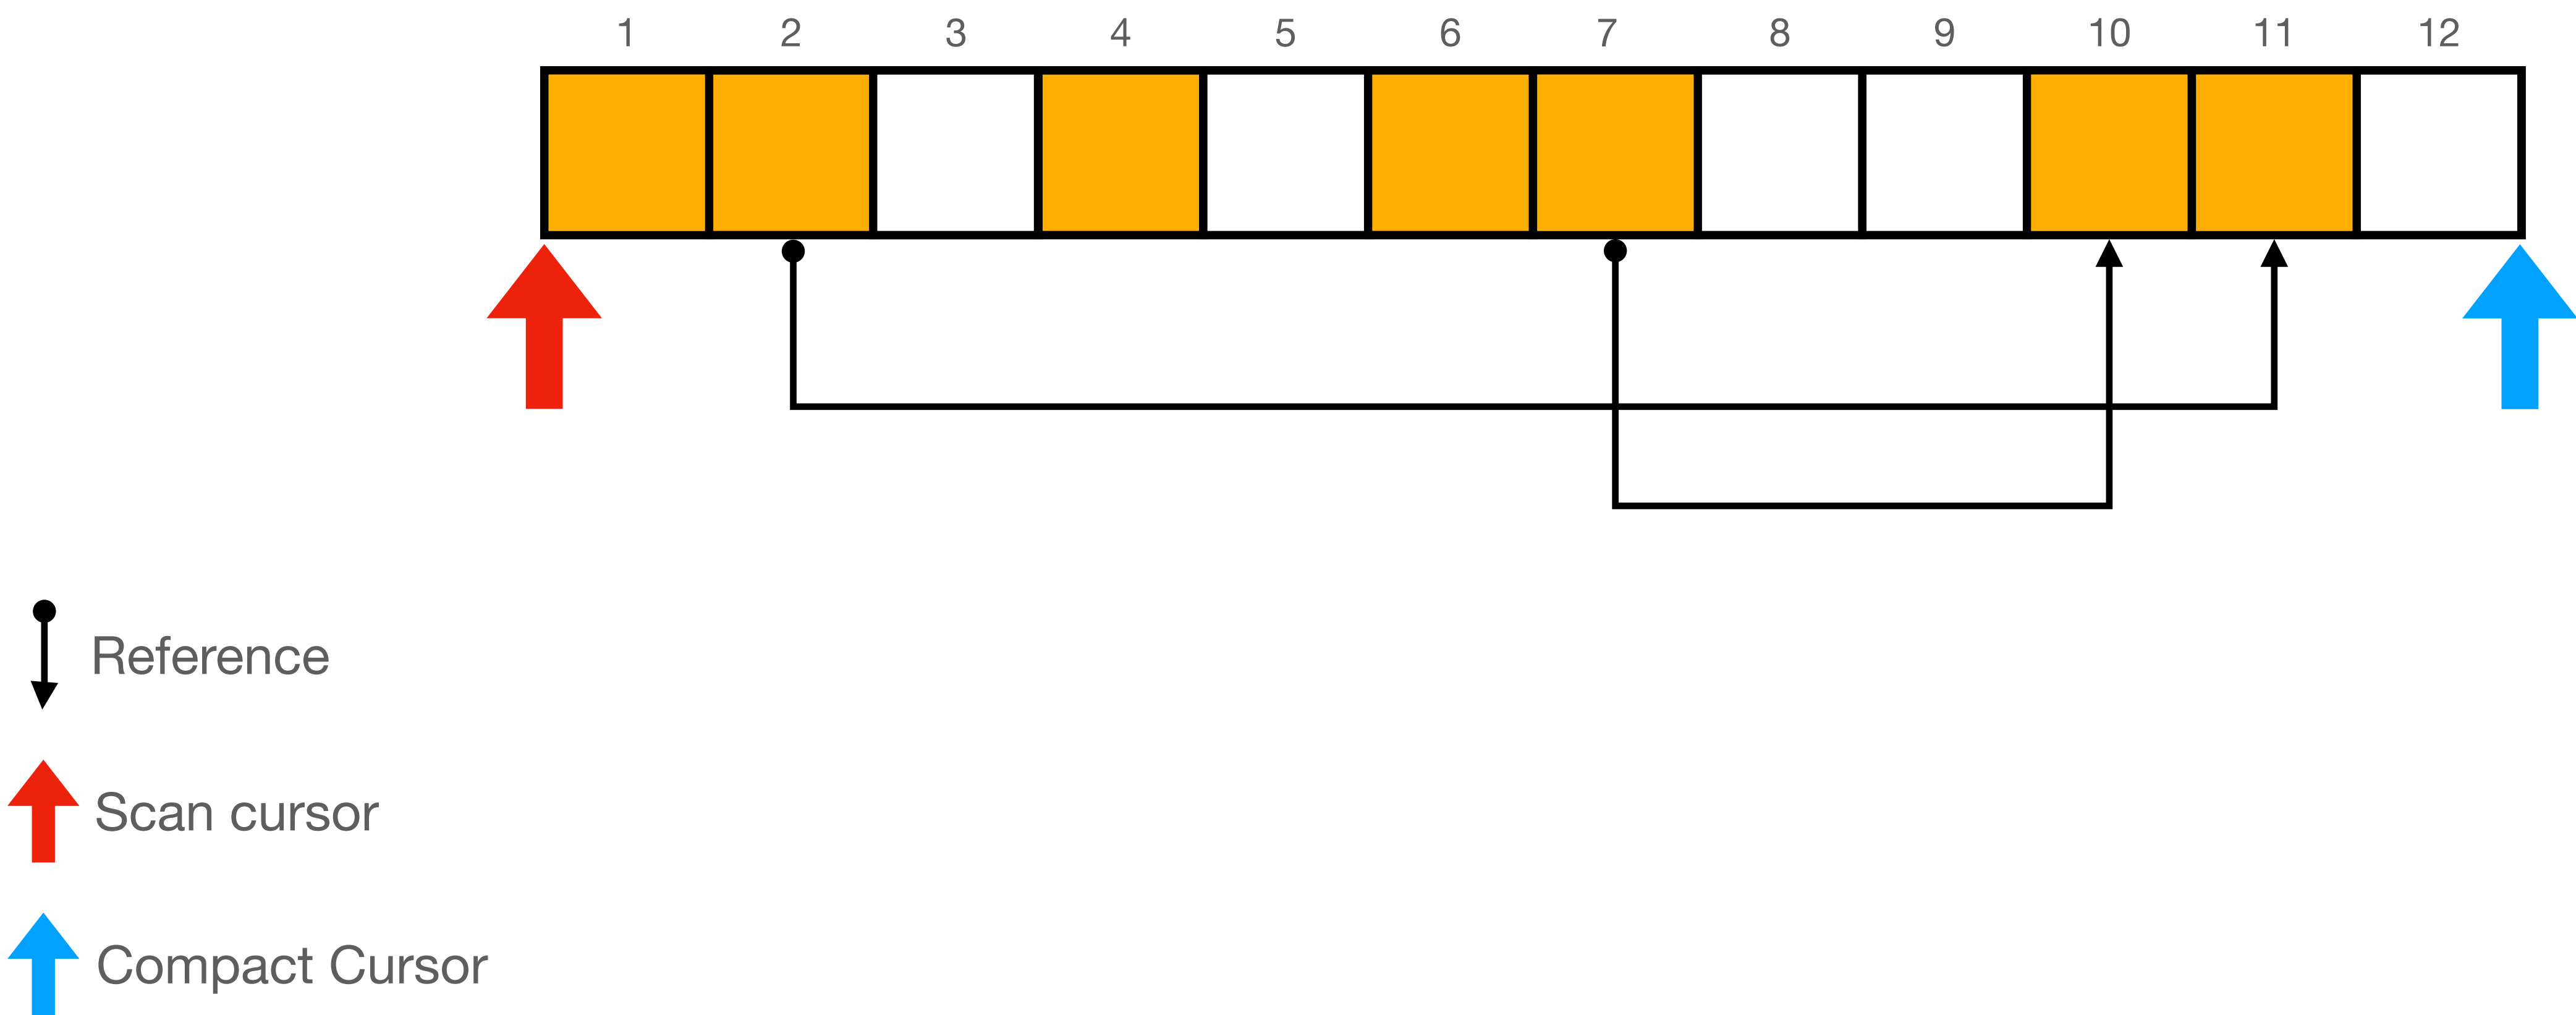
[illegible]

Background - GC Compaction

[illegible]

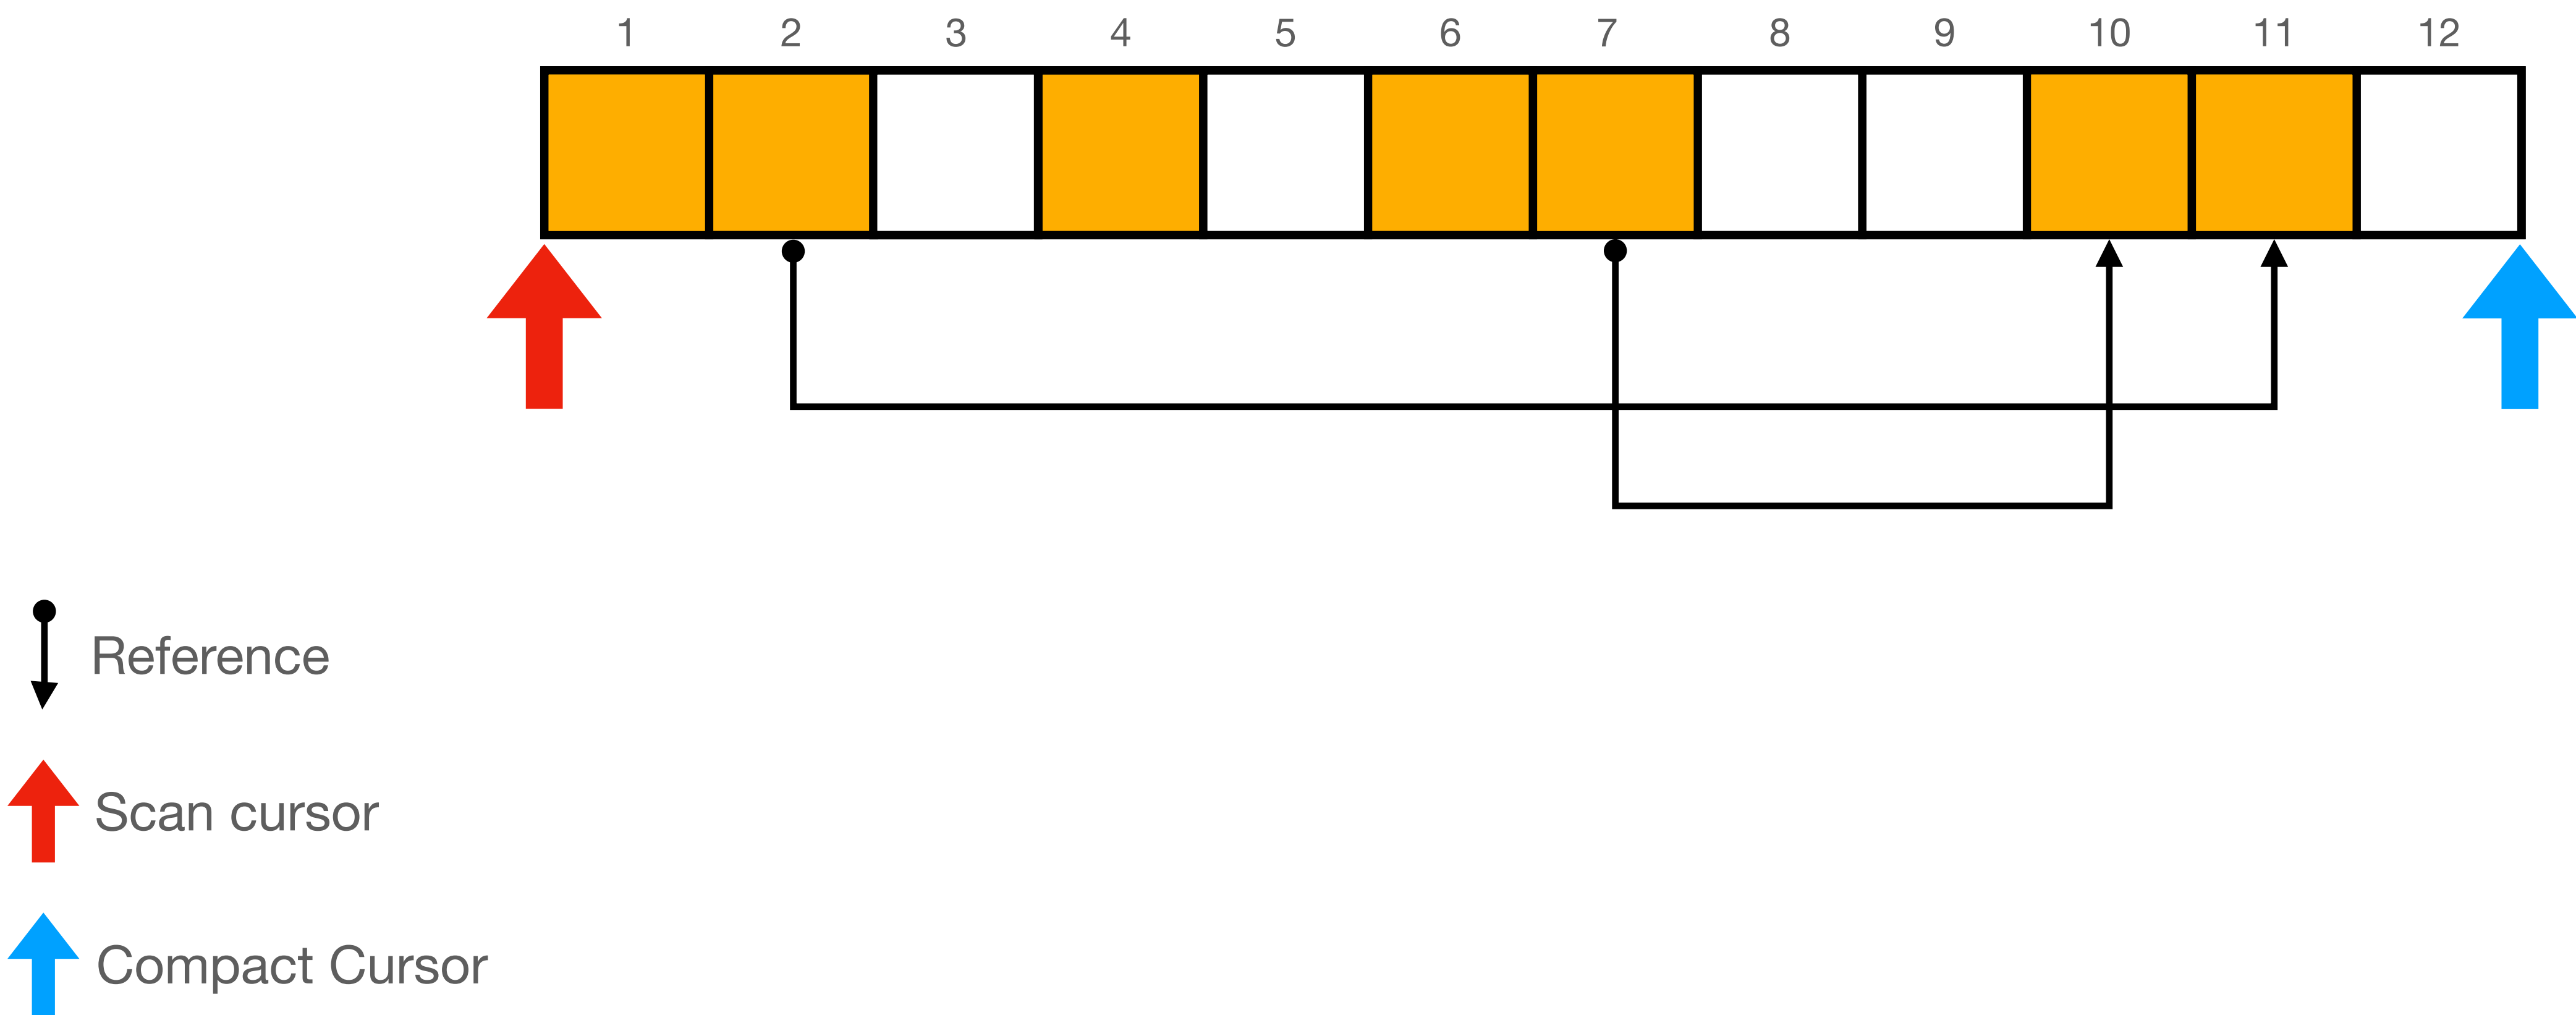
Background - The Two Finger Algorithm

Phase 1: Object Movement



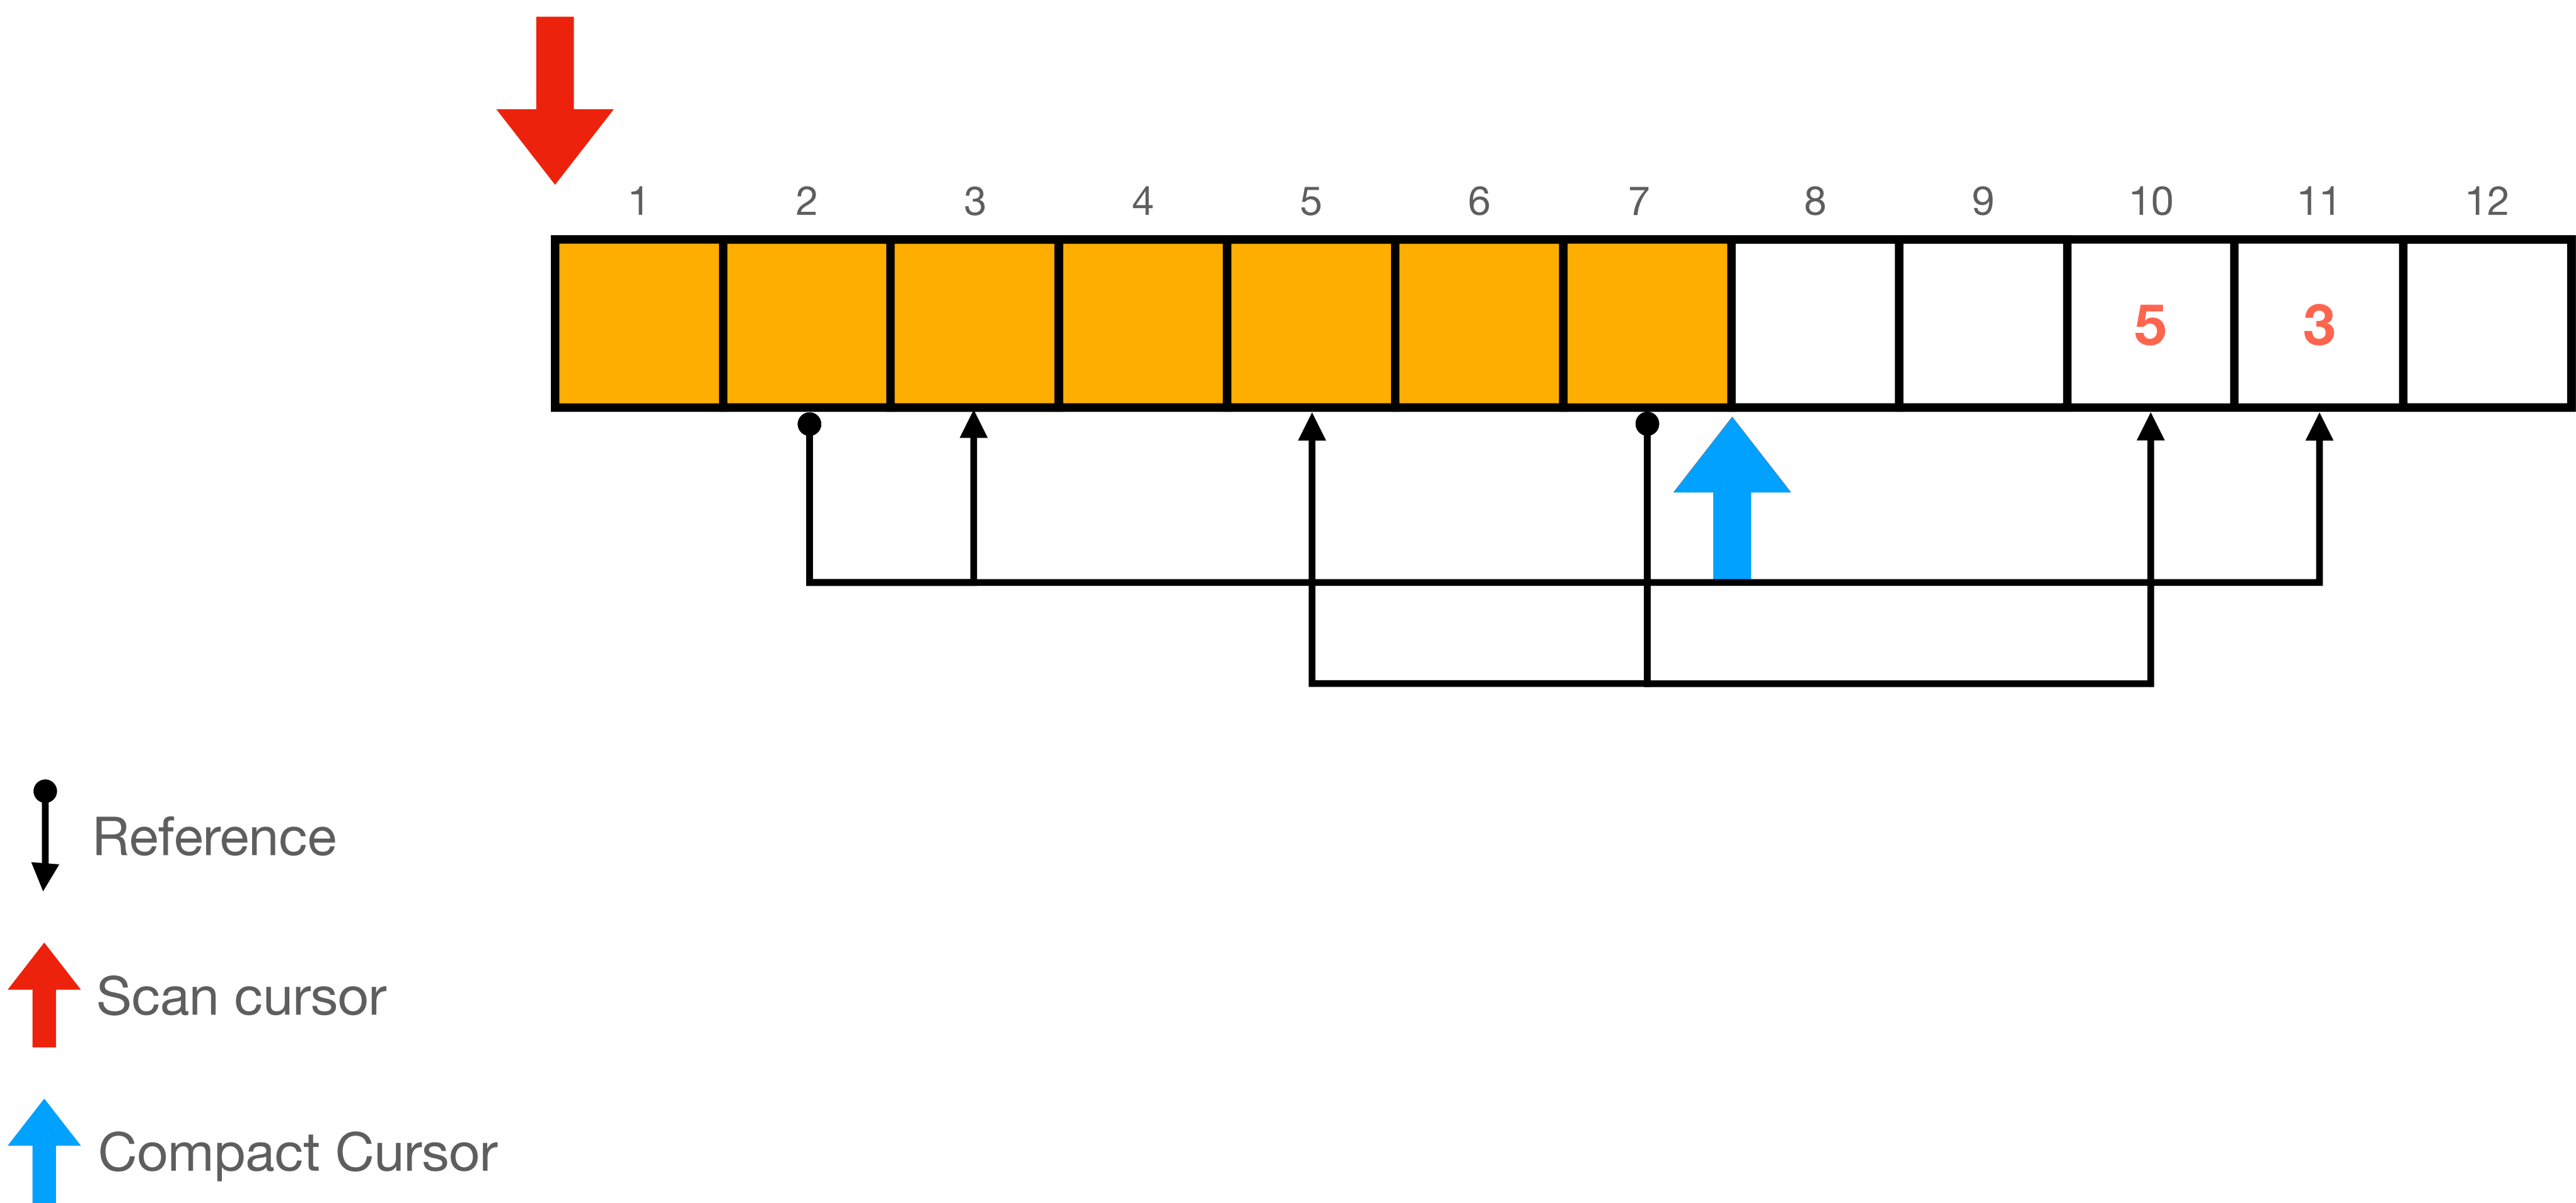
Background - The Two Finger Algorithm

Phase 1: Object Movement



Background - The Two Finger Algorithm

Phase 2: Update References



Background - GC Compaction

- Objects that cannot be moved are pinned
- `GC.compact` was introduced in Ruby 2.7
- Auto-compaction introduced in Ruby 3.0 but disabled by default
- Ruby 3.0 integrated sweep and compact

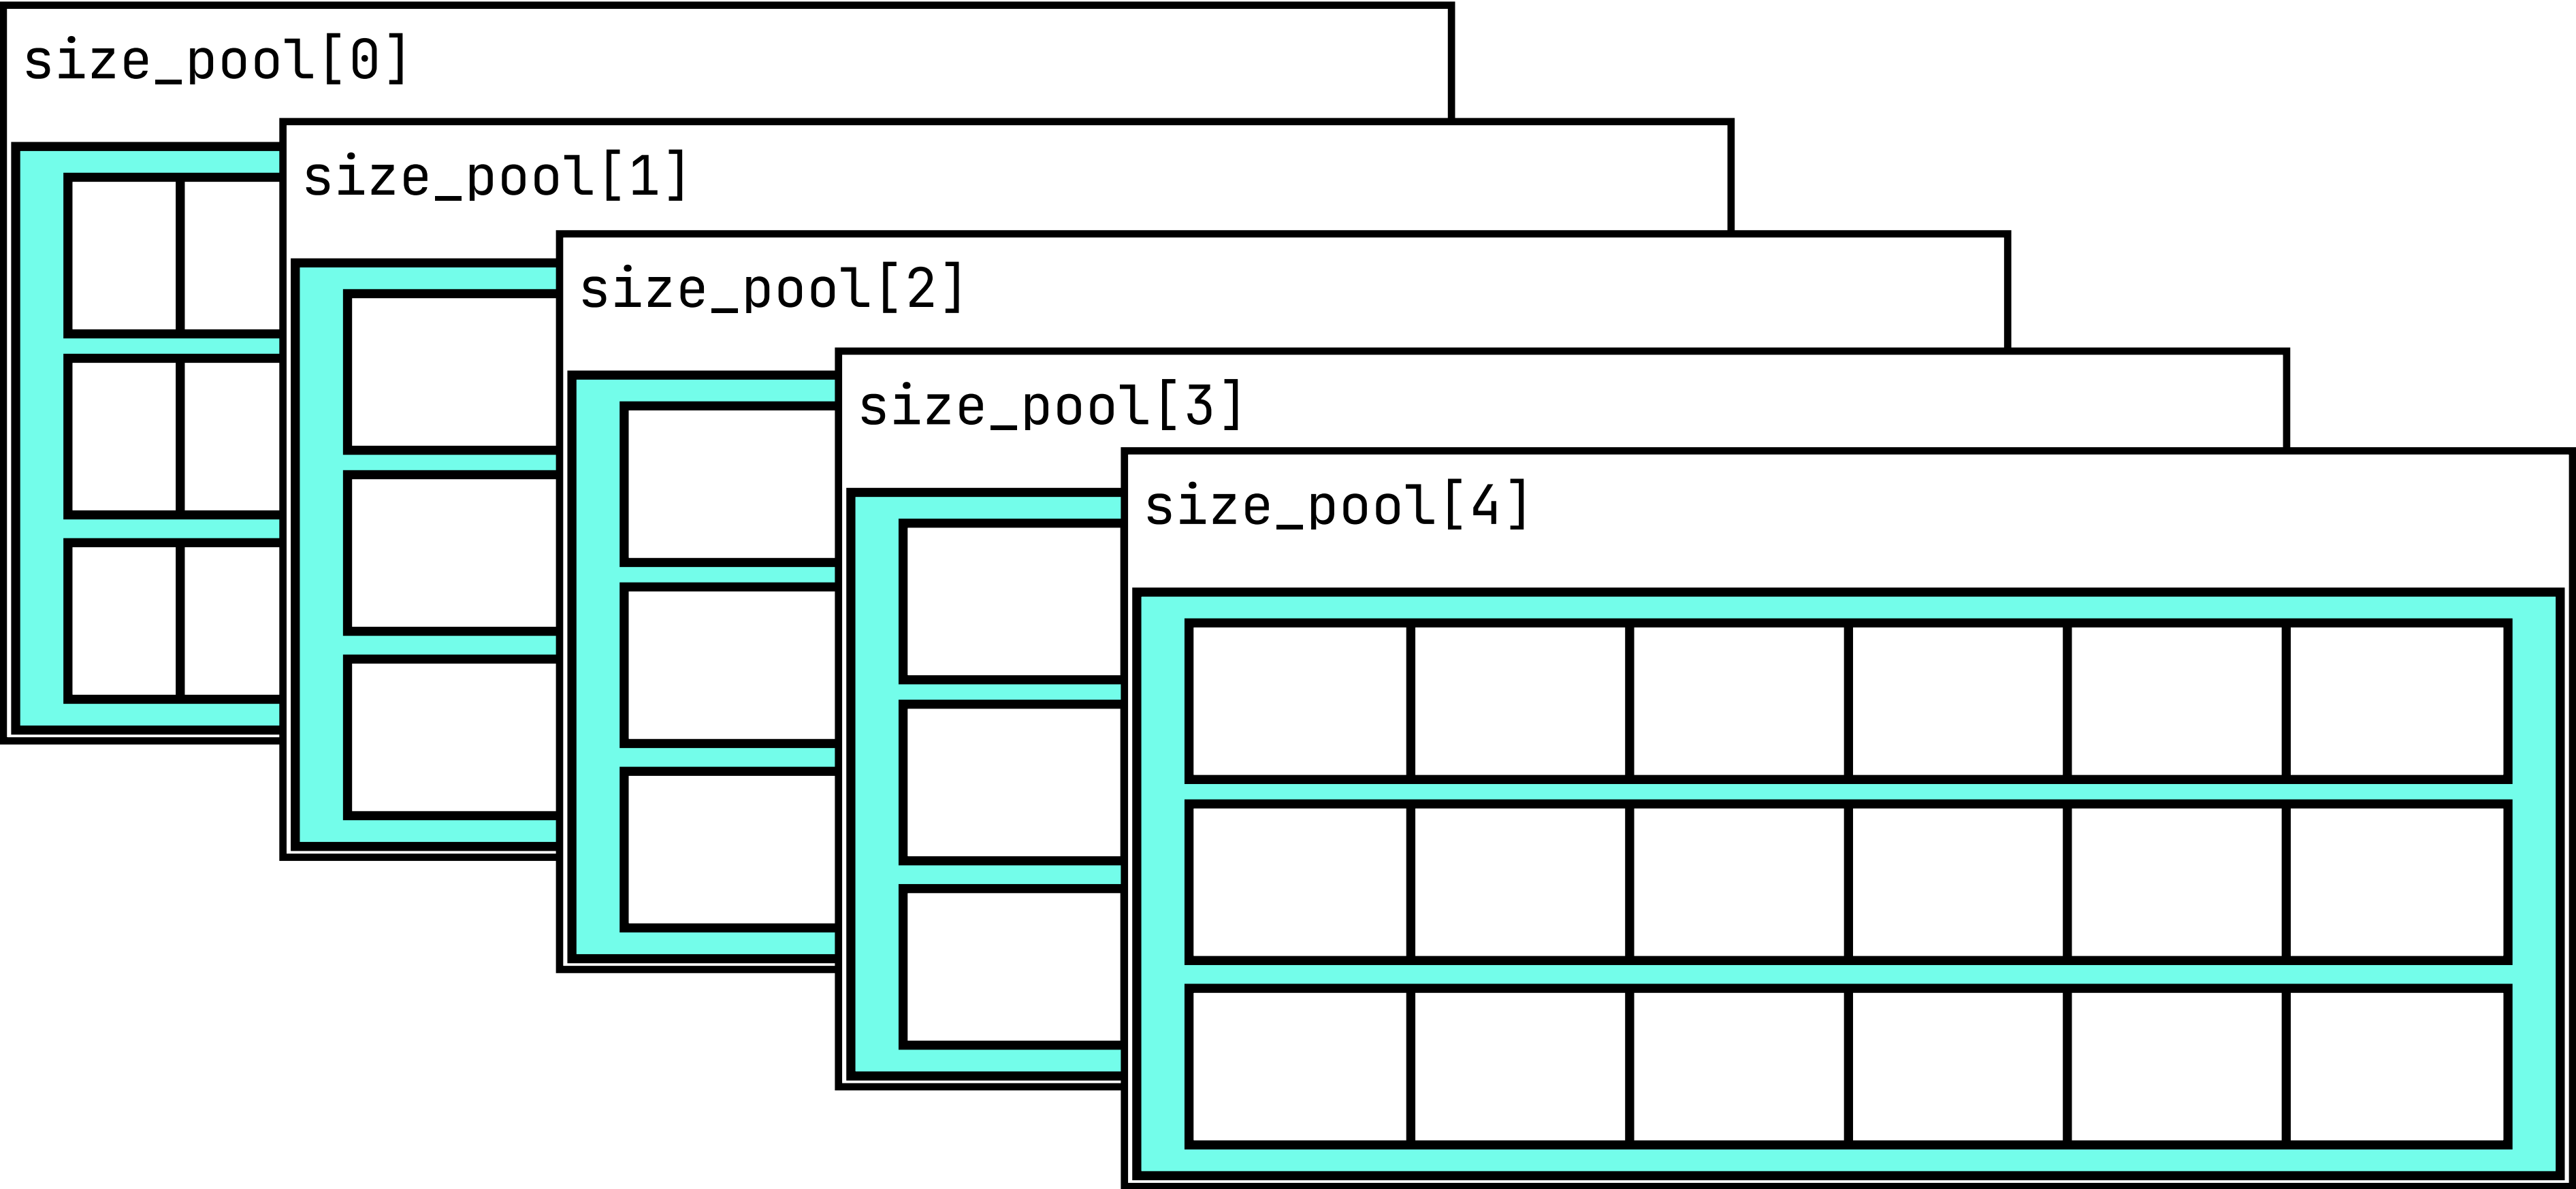
Background - Summary

- **Heap/Pages/RVALUES** - Ruby stores objects in fixed size slots in pages
- **Compacting GC** - Objects get moved during sweeping to improve fragmentation
- **Objects can have poor data locality** - related data is often far apart in memory
- **VWA is a project to improve data locality** - by restructuring Ruby's heap

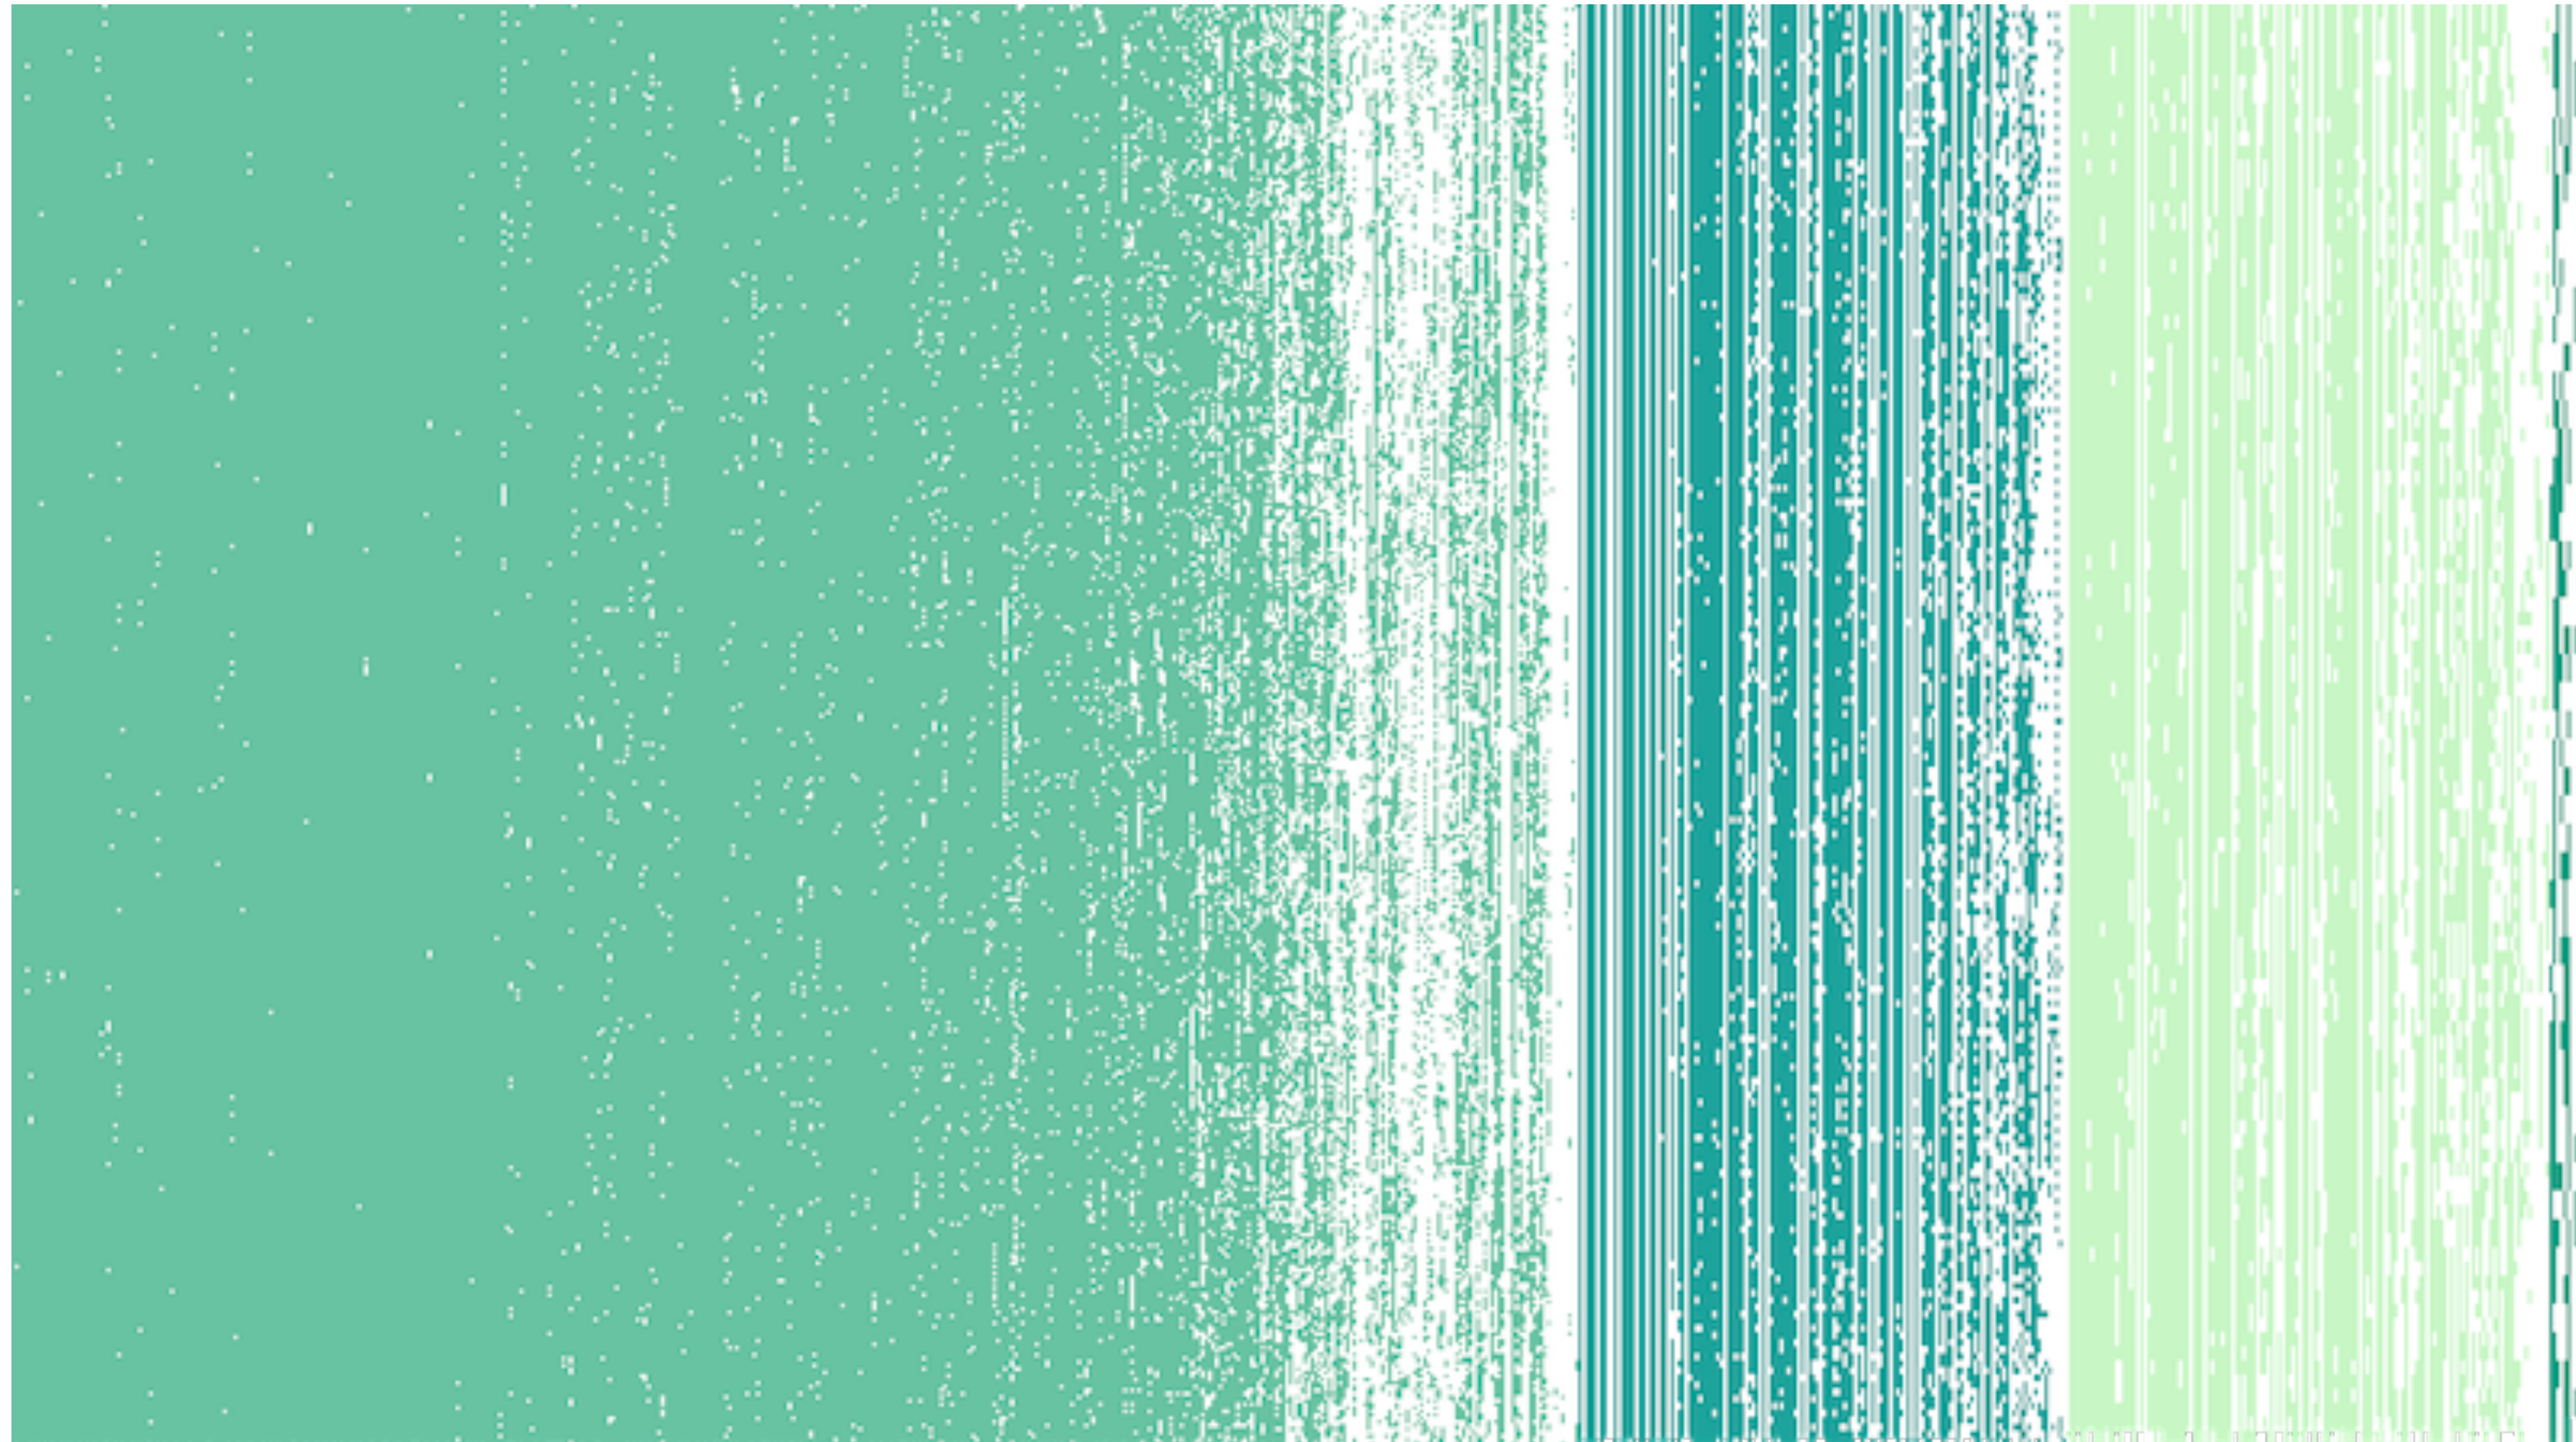
The Problem

Uh-Oh

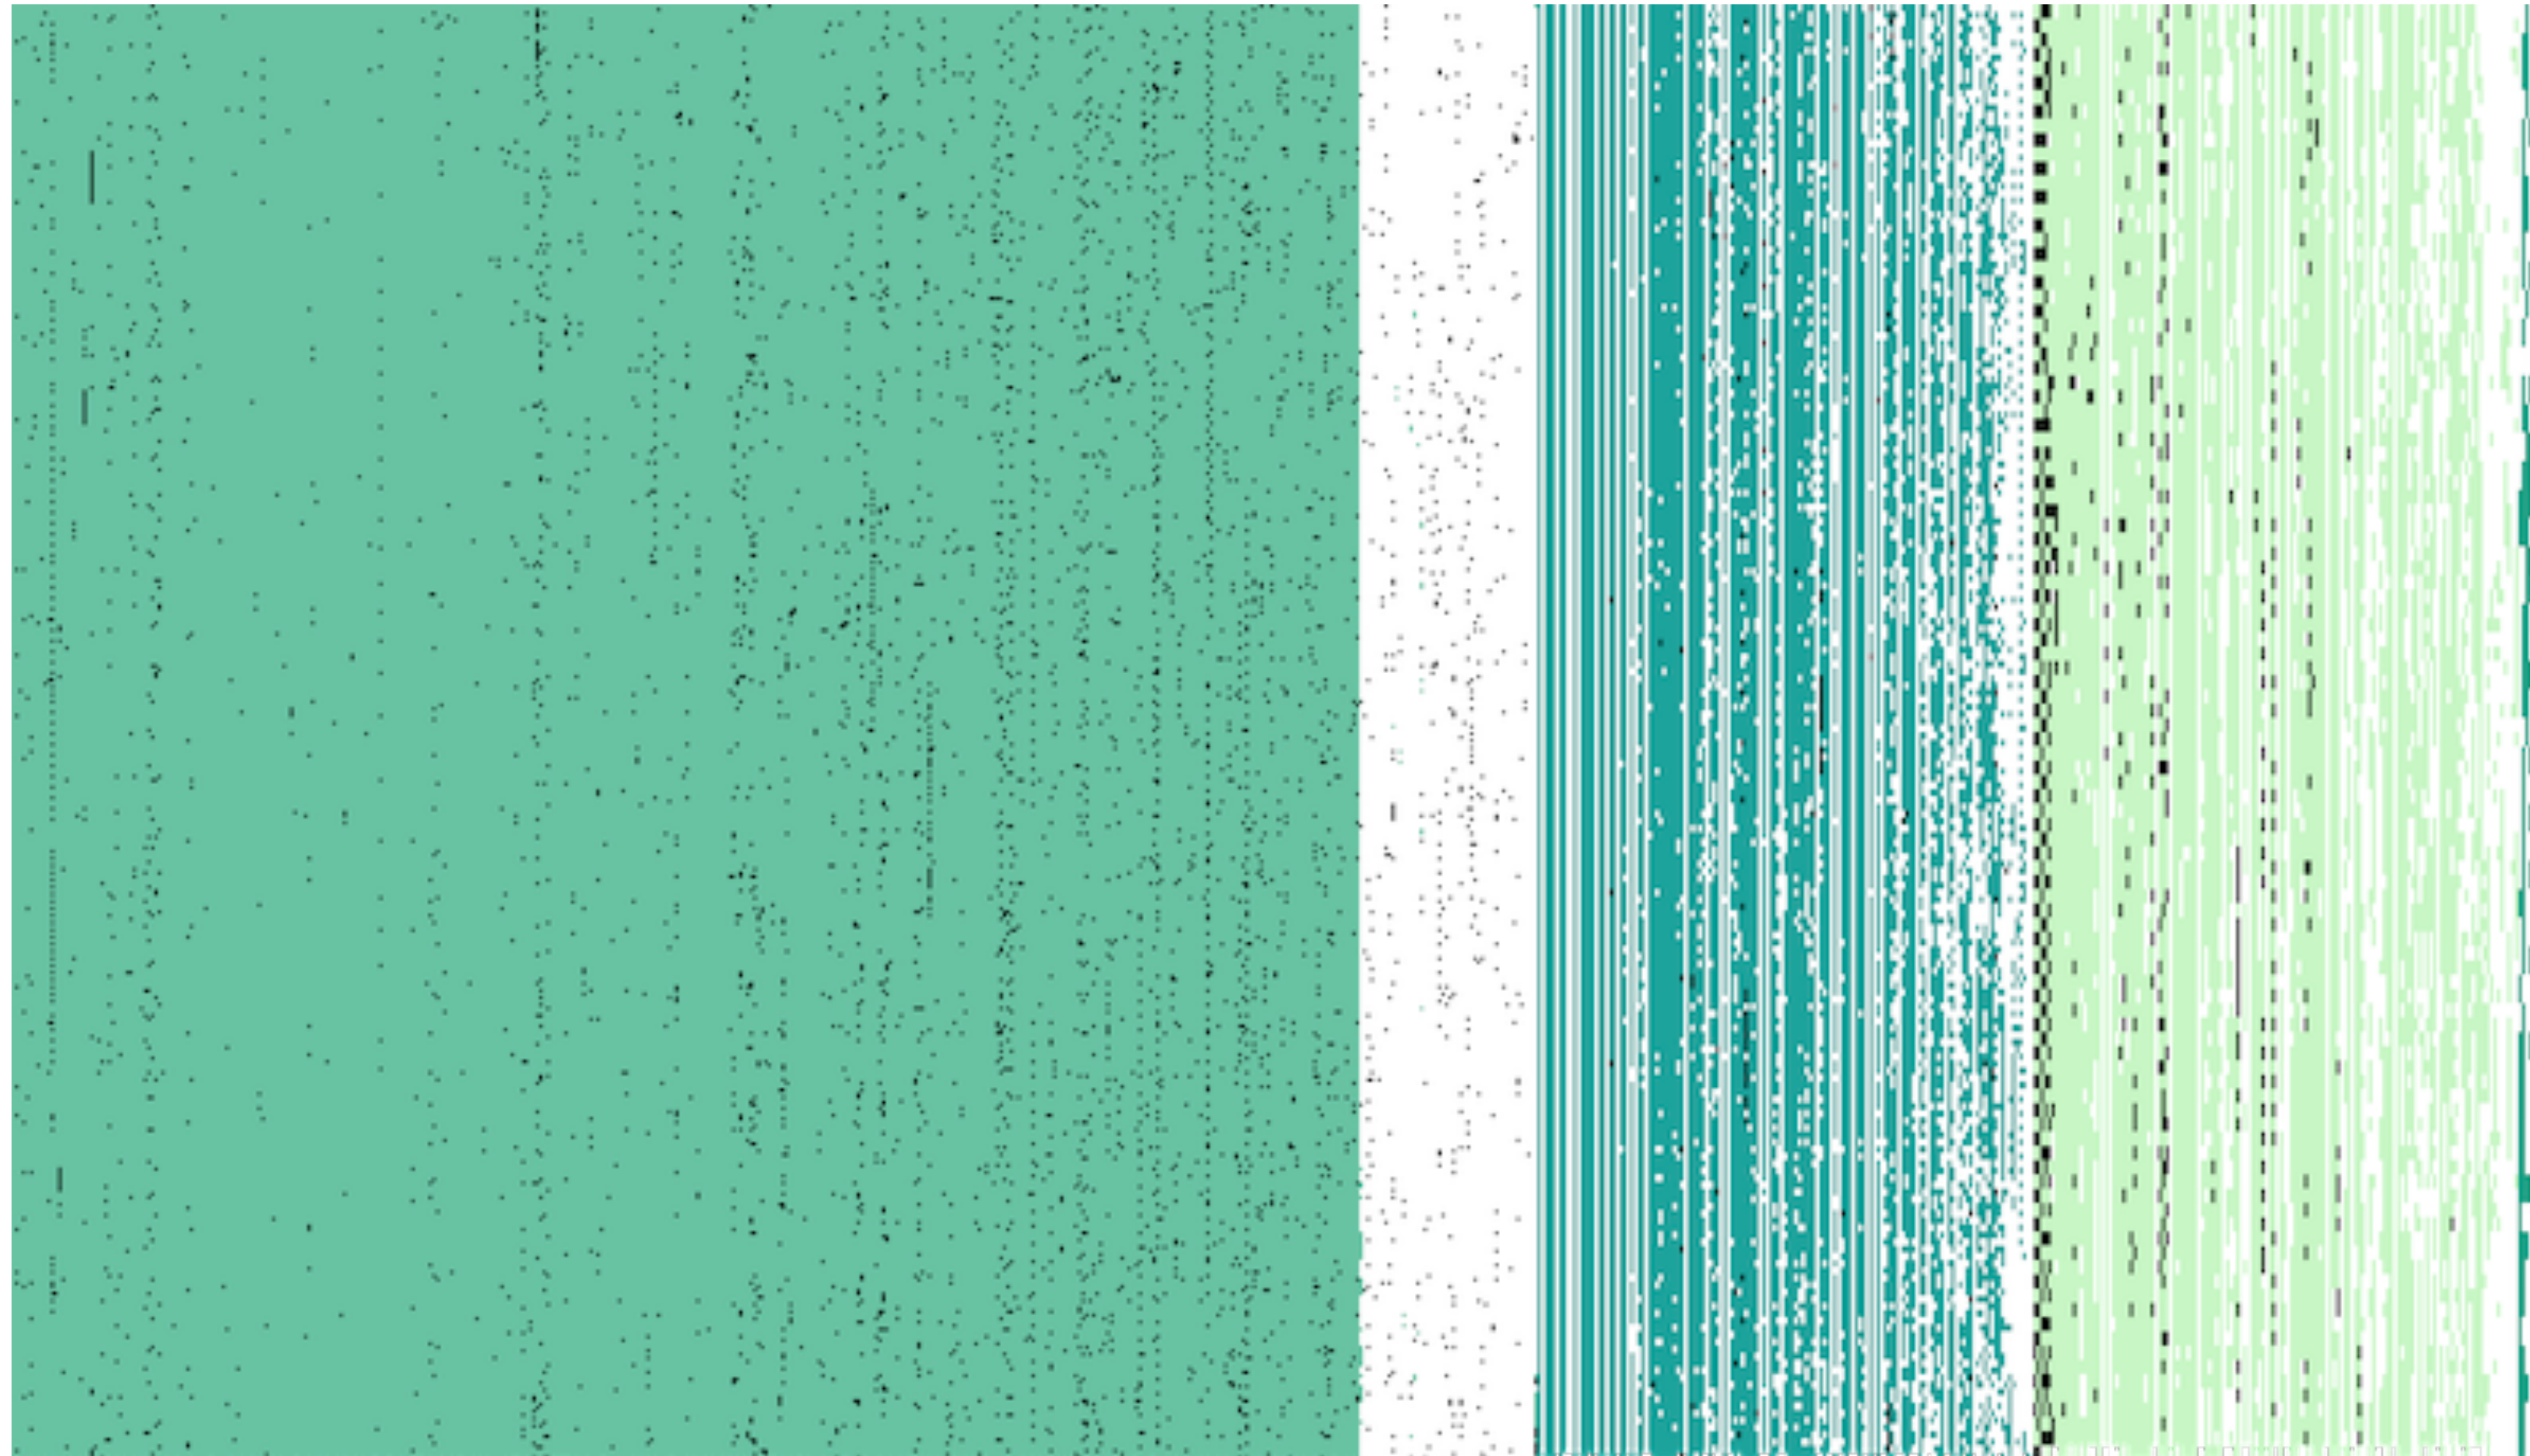
The Problem



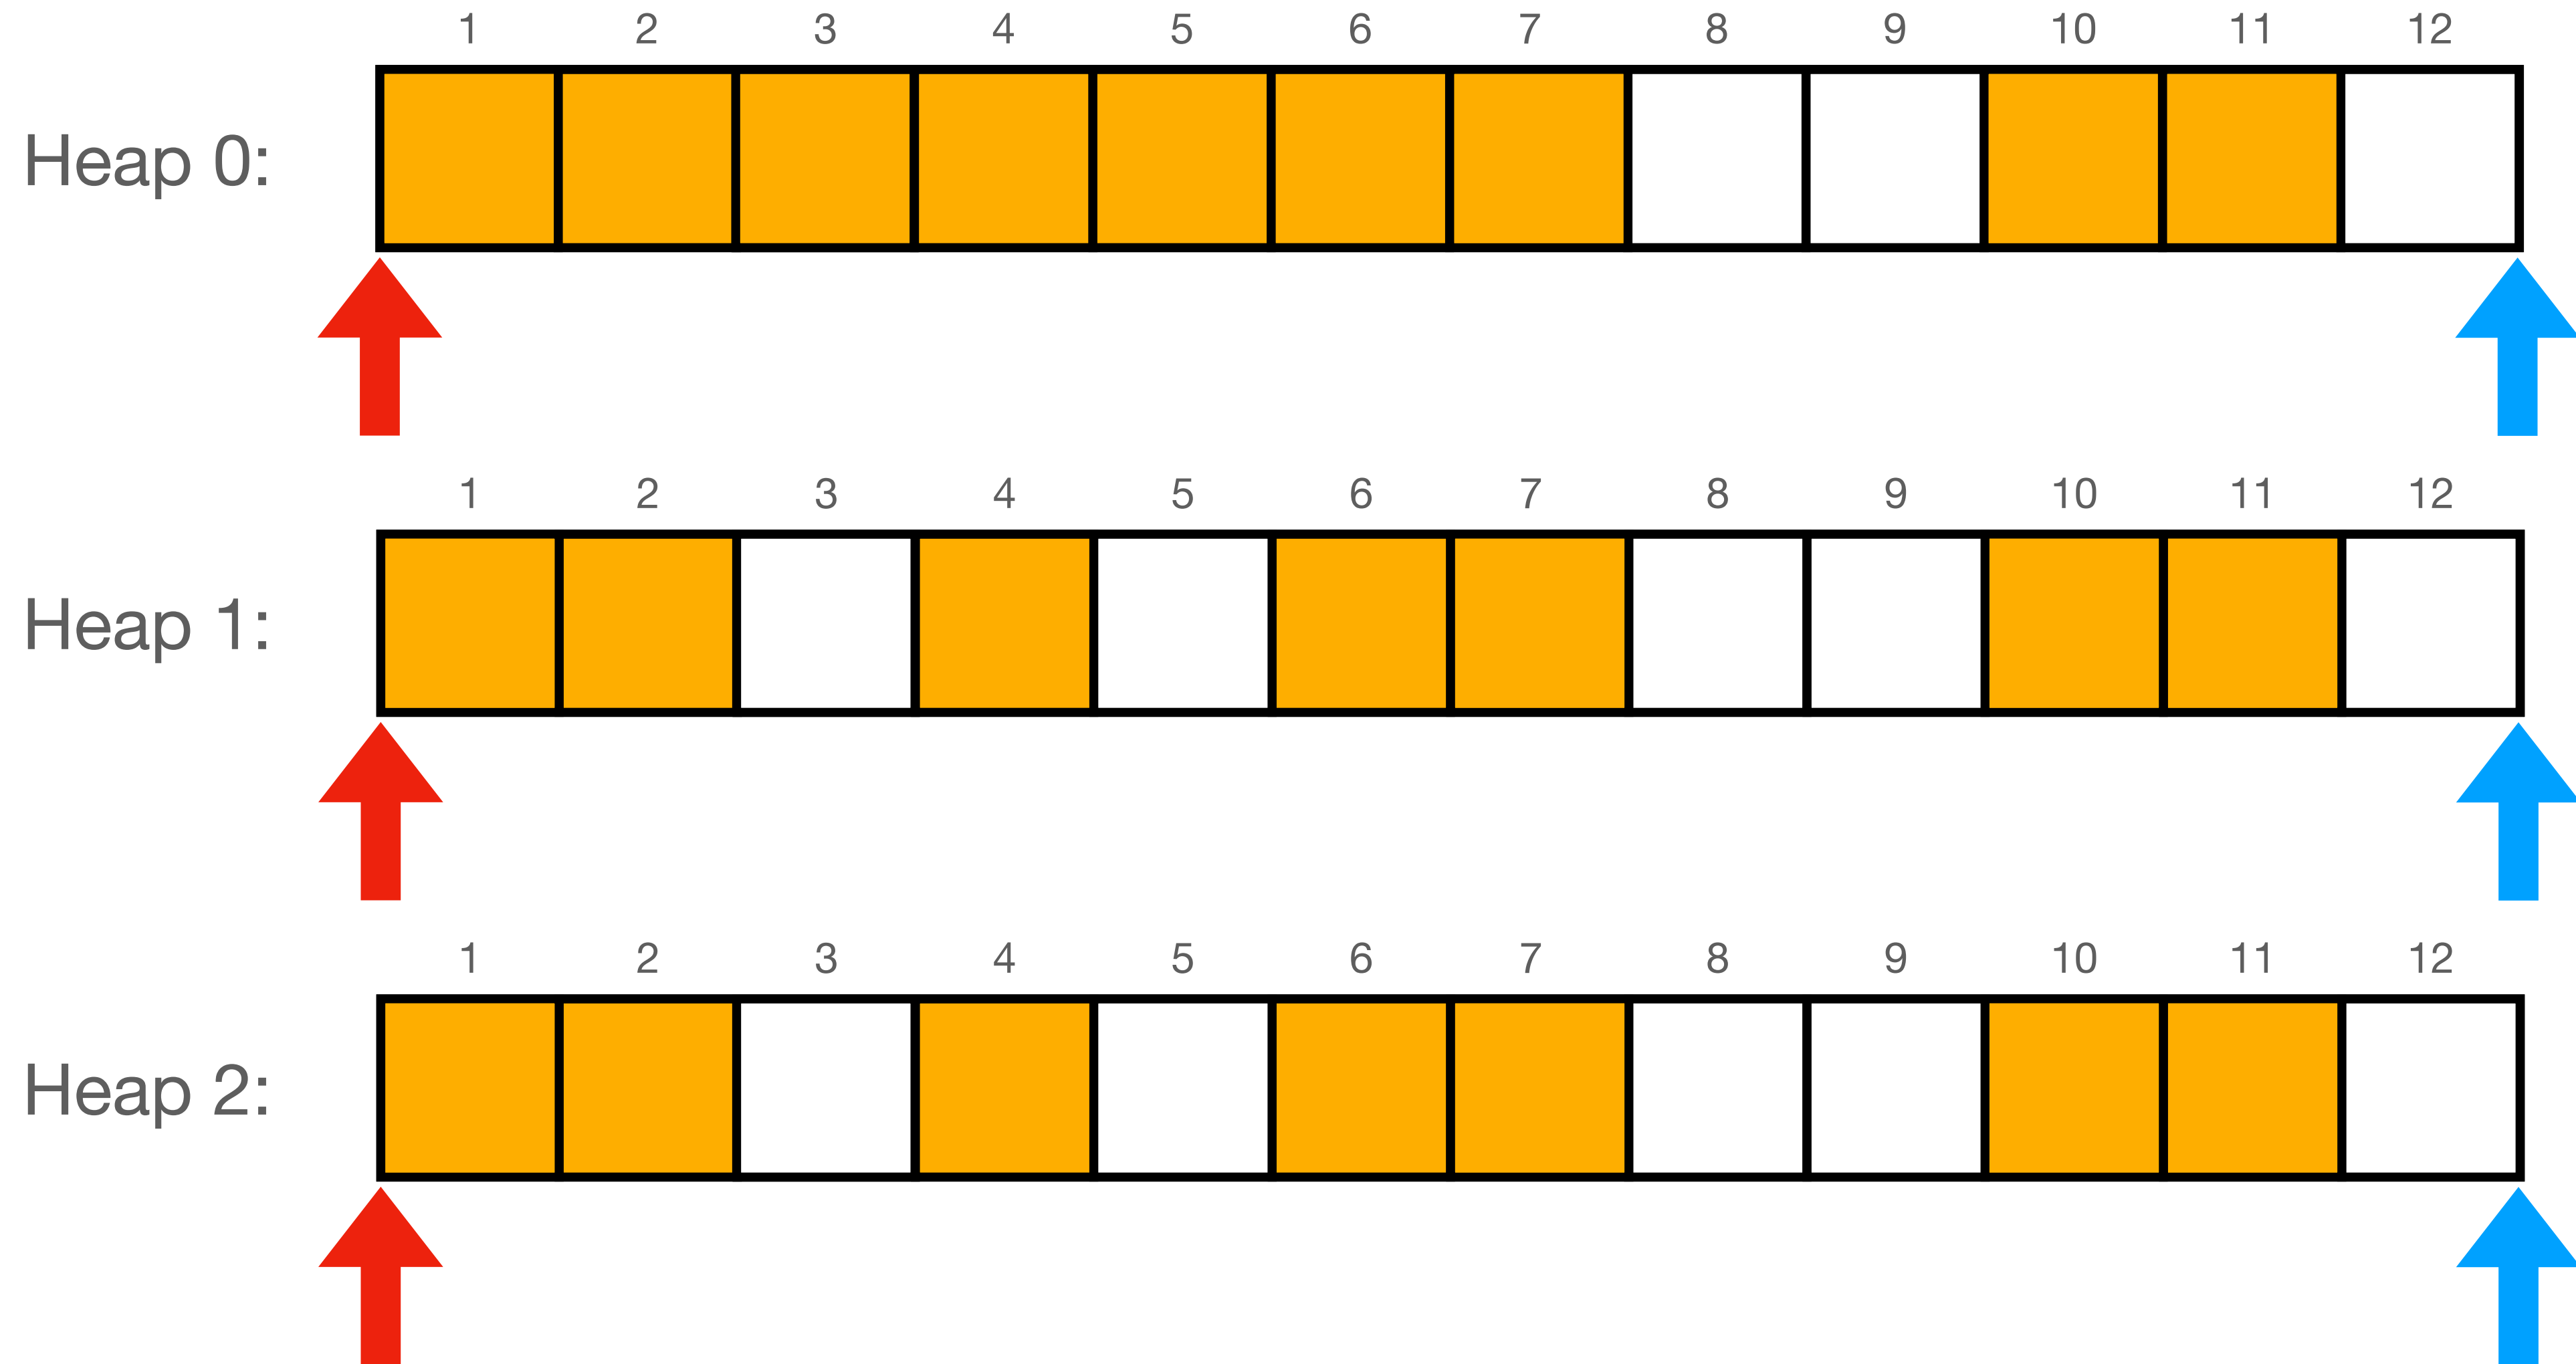
The Problem - Uncompacted Size Pools



The Problem - Compacted Size Pools



The Problem



`objspace→flags.during_compacting = false;`

The Problem - Attempt 1

Compact each heap independently

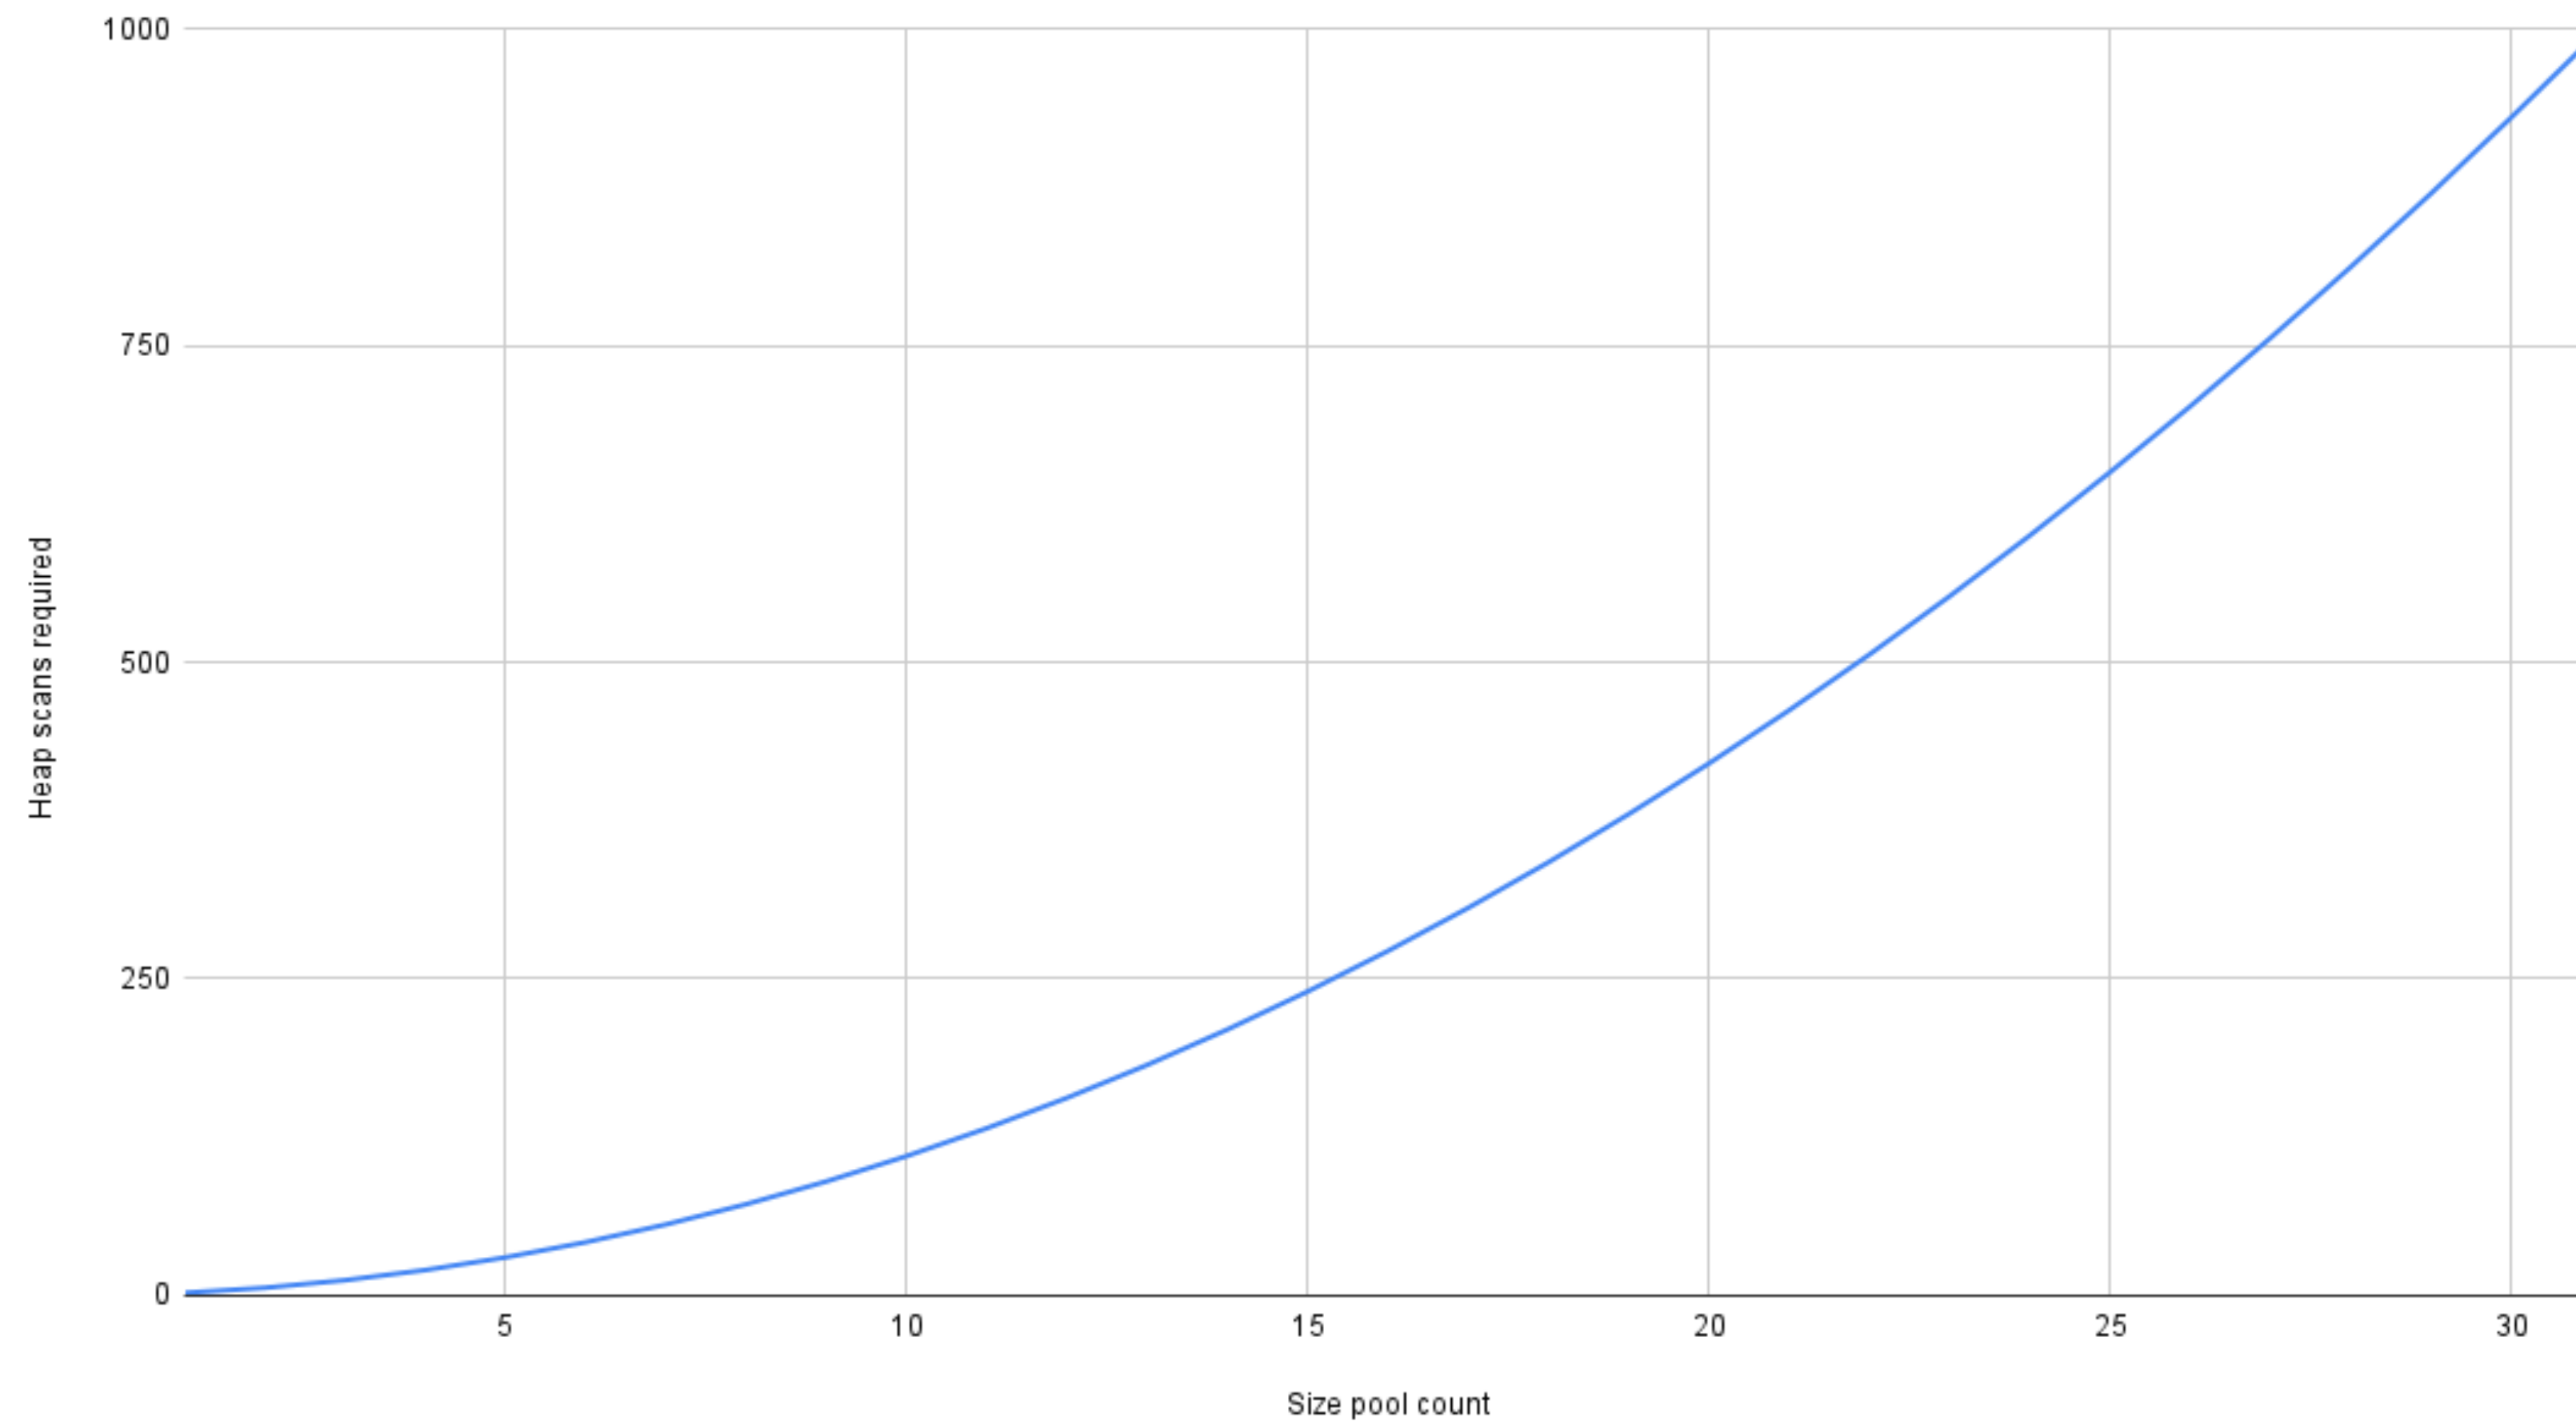
 Simple

Scope all compaction context into the size pool. Each heap has it's own compaction status flag.

The Problem - Attempt 1

Compact each heap independently


Heap scans required vs Size pool count



The Problem - Attempt 1

Compact each heap independently

 Simple

 Unacceptably slow due to
excessive heap scans

Scope all compaction context into
the size pool. Each heap has it's
own compaction status flag.

An Aside

Resizing Objects

✓ hash = { a: true, b: false }

✓ string = "Hello, RubyKaigi"

✓ array = [1, 2, 3, a, b, c]

An Aside

Resizing Objects

✓ `hash = { a: true, b: false }`

✓ `string = "Hello, RubyKaigi"`

✓ `array = [1, 2, 3, a, b, c]`

✗ `array = []`
`array.push(1)`
`array.push(2)`

✗ `string << ", pleased to meet you"`

An Aside

Resizing Objects

```
string = "Hello"  
string << " RubyKaigi, thanks for having me"
```

string

Meta-data	ary: "Hello"
EMBED: 1	

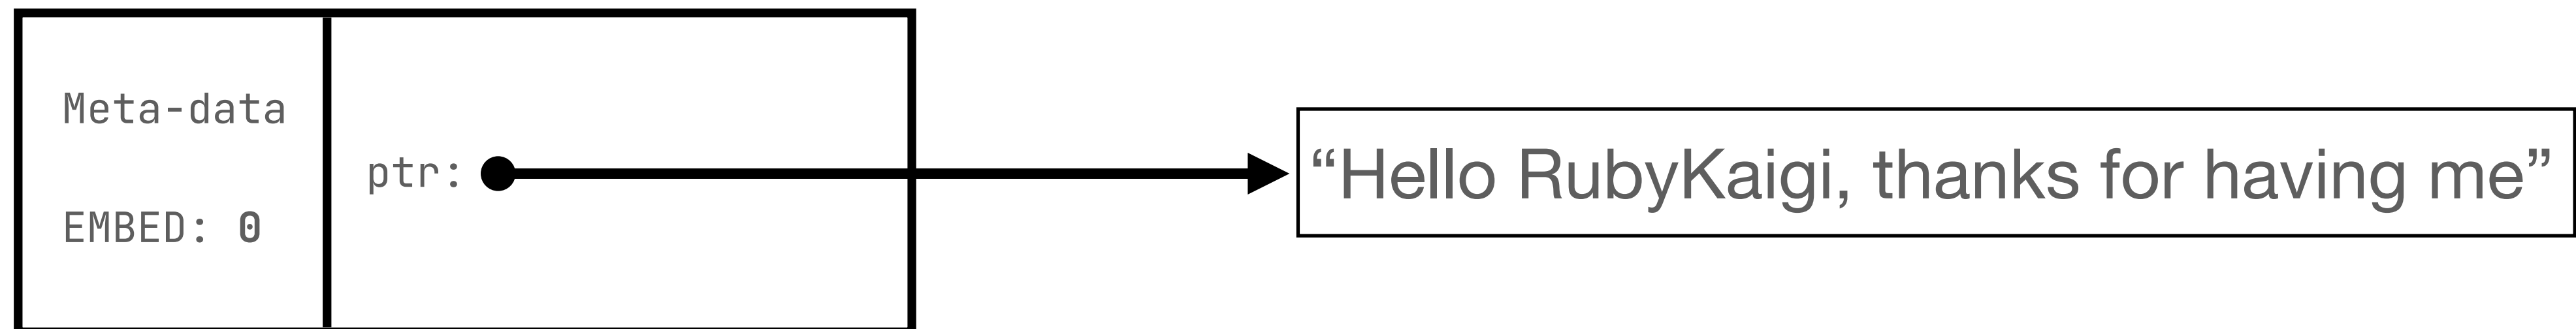


An Aside

Resizing Objects

```
string = "Hello"  
string << " RubyKaigi, thanks for having me"
```

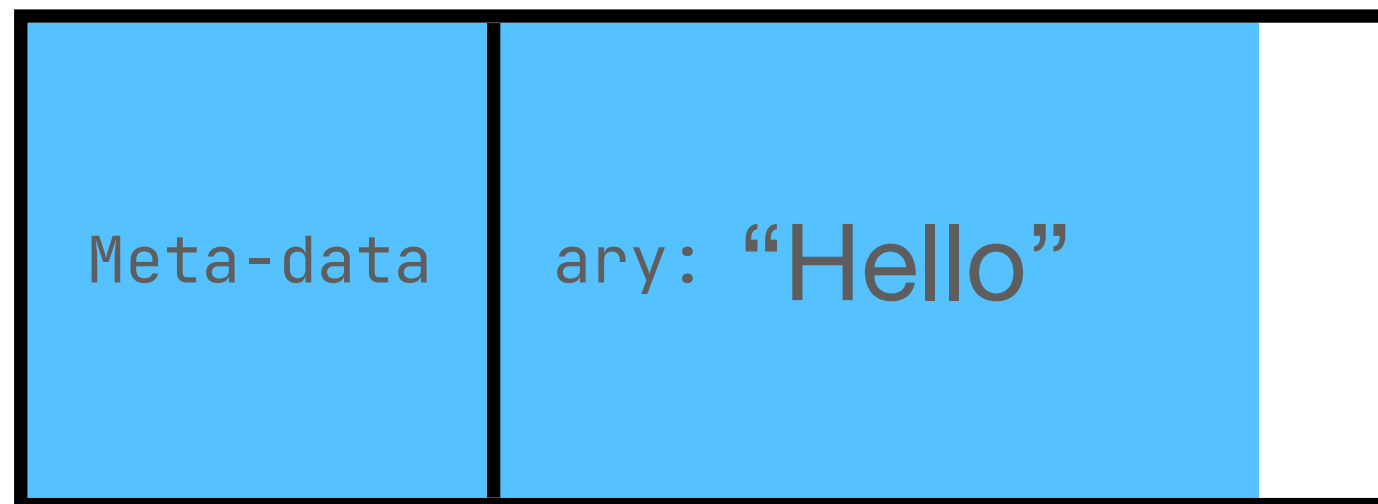
string



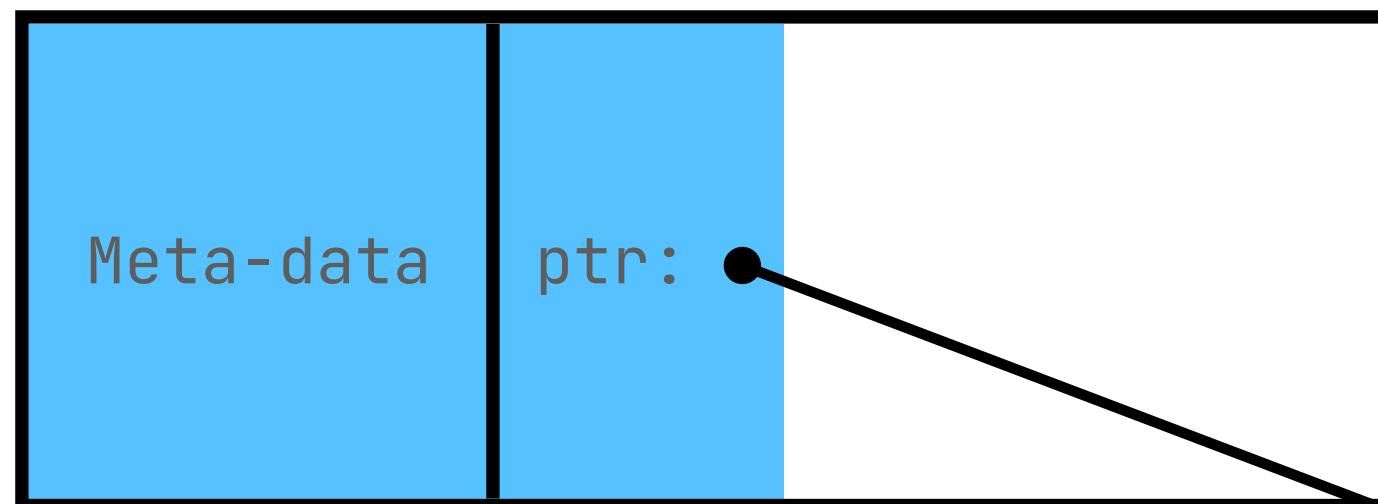
An Aside

Resizing Objects

RString



RString




“Hello RubyKaigi, thanks for having me”

The Problem - Attempt 1

Compact each heap independently

Scope all compaction context into the size pool. Each heap has its own compaction status flag.

 Simple

 Unacceptably slow due to excessive heap scans

 Potentially wasteful of memory

The Problem - Attempt 1

Compact each heap independently

Scope all compaction context into the size pool. Each heap has its own compaction status flag.

✓ Simple

✗ Unacceptably slow due to excessive heap scans

✗ Potentially wasteful of memory

✗ Doesn't allow for object movement between size pools

The Problem - Attempt 2

Seperate Object Movement from Reference Updating

Perform object movement on each heap in turn.

When object movement is complete for all heaps, then update references in a single step

The Problem - Attempt 2

Seperate Object Movement from Reference Updating

Perform object movement on each heap in turn.

When object movement is complete for all heaps, then update references in a single step

✓ Faster due to fewer heap scans. References are only updated once

The Problem - Attempt 2

Seperate Object Movement from Reference Updating

Perform object movement on each heap in turn.

When object movement is complete for all heaps, then update references in a single step

- ✓ Faster due to fewer heap scans. References are only updated once
- ✓ Allows some limited object movement between pools

The Problem - Attempt 2

Seperate Object Movement from Reference Updating

Perform object movement on each heap in turn.

When object movement is complete for all heaps, then update references in a single step

- ✓ Faster due to fewer heap scans. References are only updated once
- ✓ Allows some limited object movement between pools
- ✗ More complex algorithm

The Problem - Attempt 2

Seperate Object Movement from Reference Updating

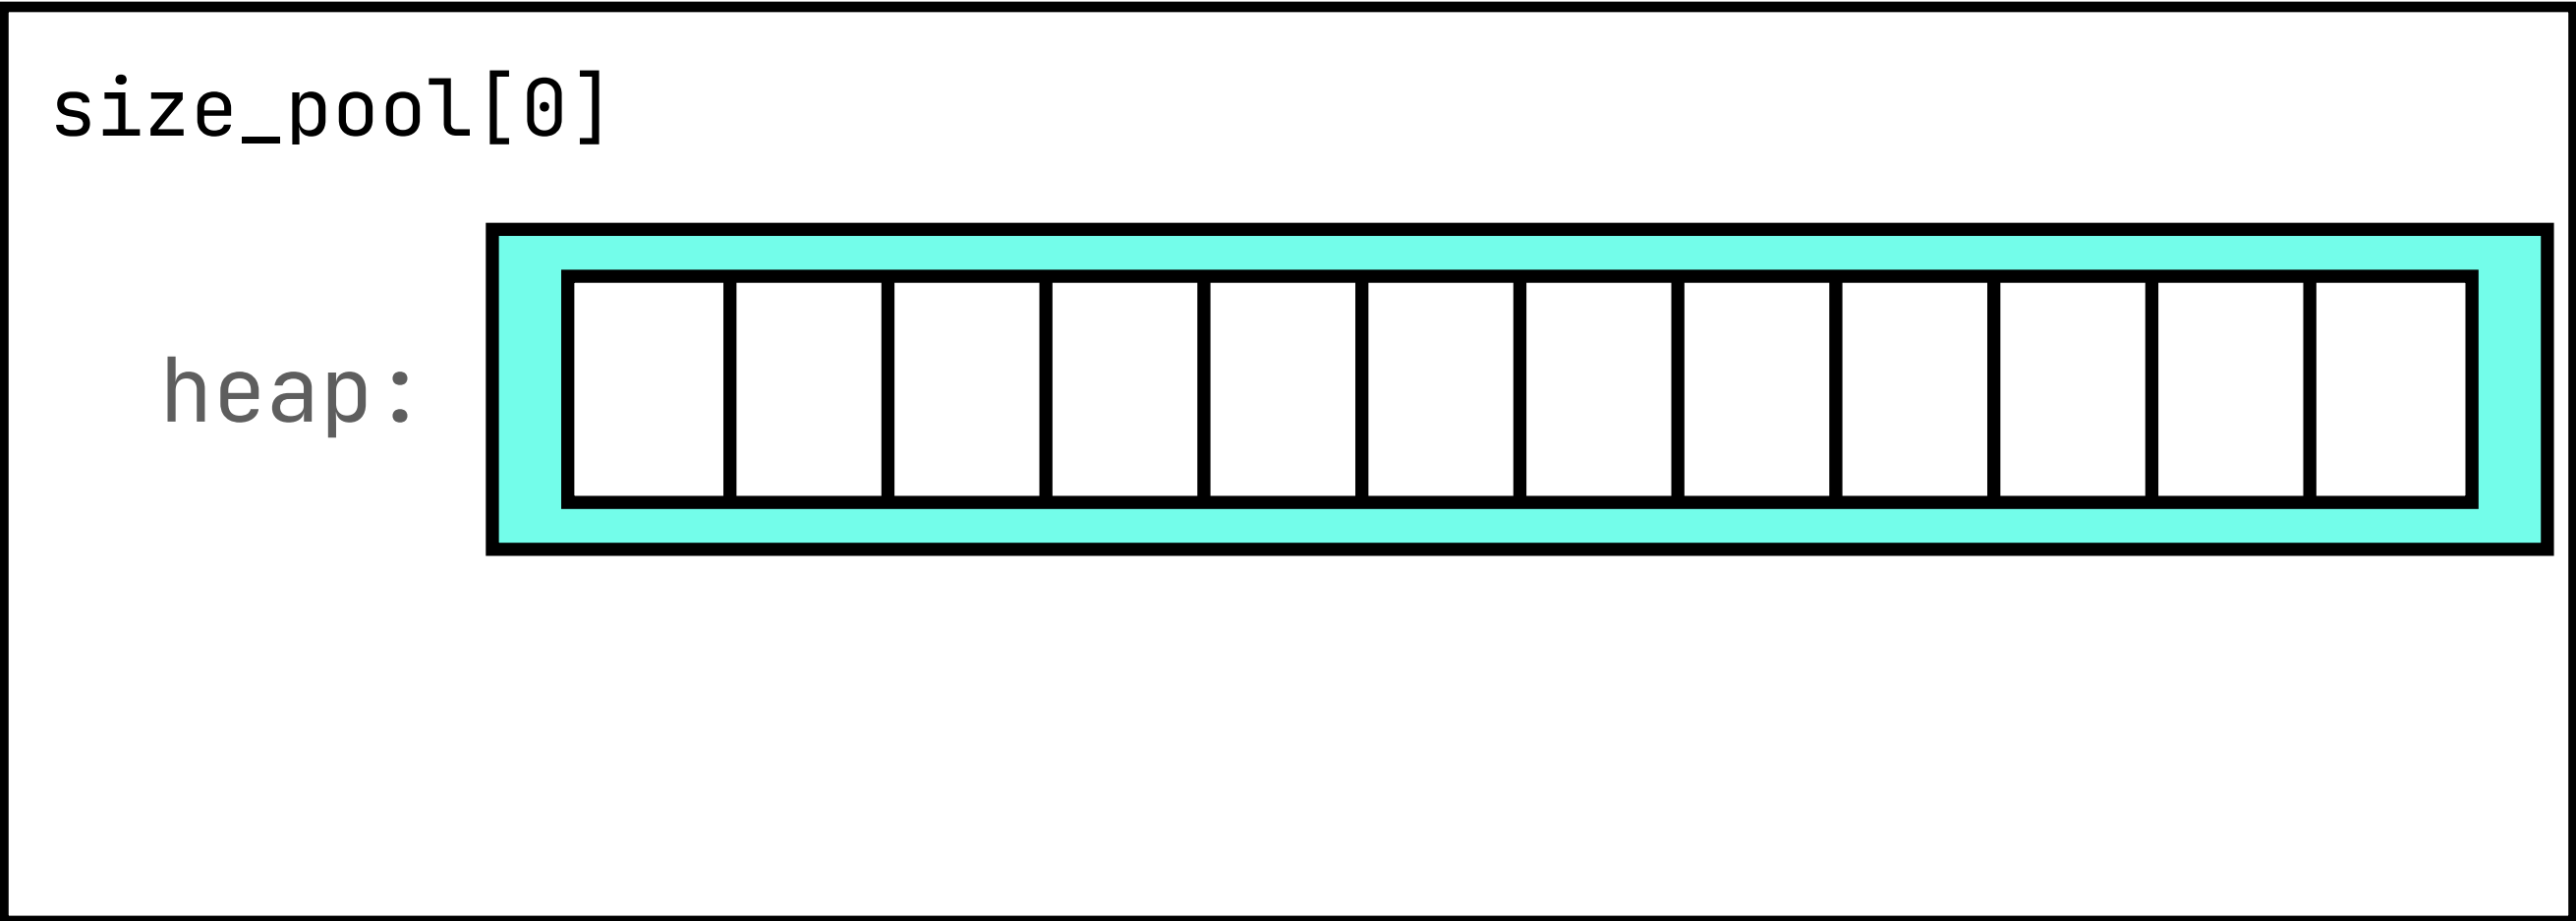
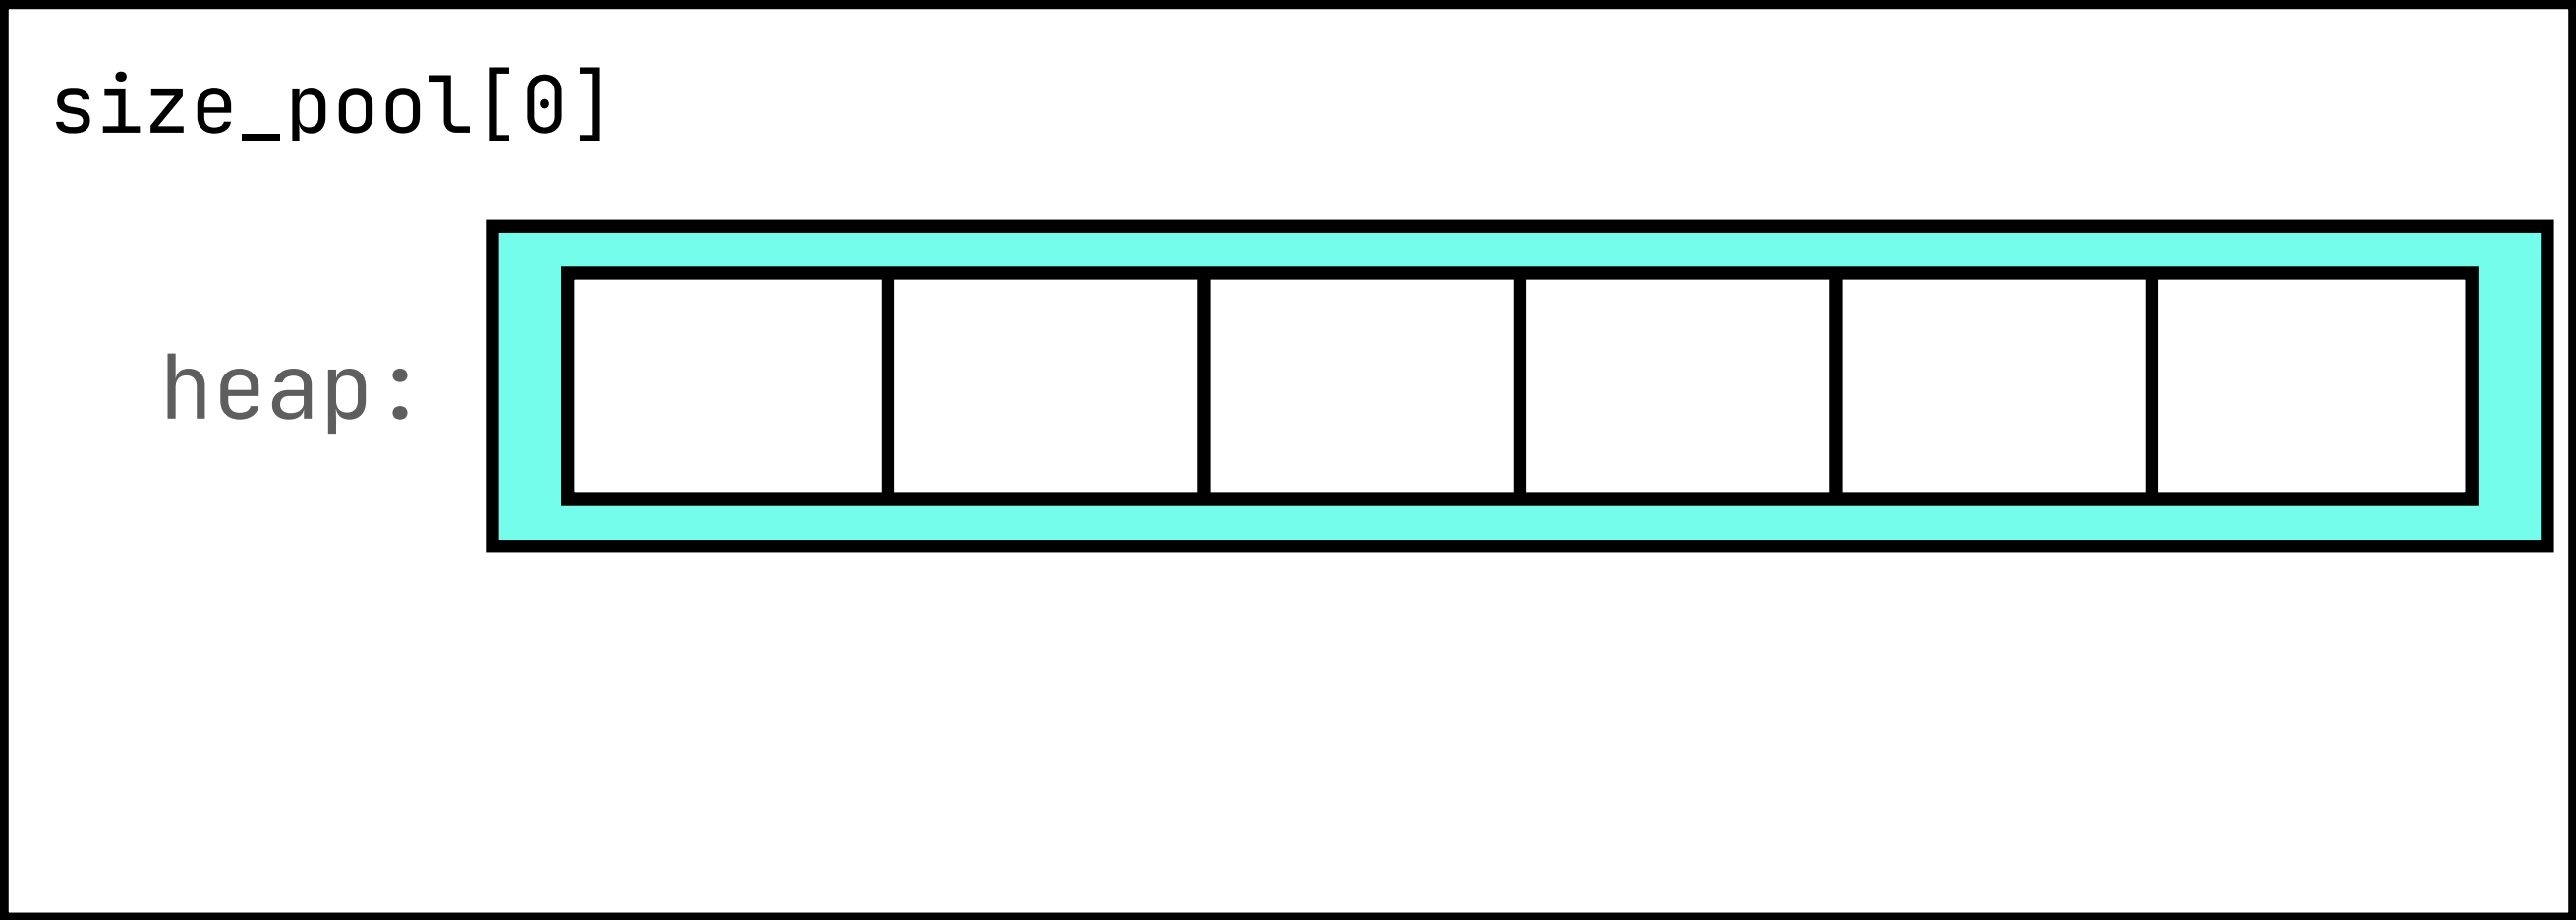
Perform object movement on each heap in turn.

When object movement is complete for all heaps, then update references in a single step

- ✓ Faster due to fewer heap scans. References are only updated once
- ✓ Allows some limited object movement between pools
- ✗ More complex algorithm
- ✗ Object Movement only one way

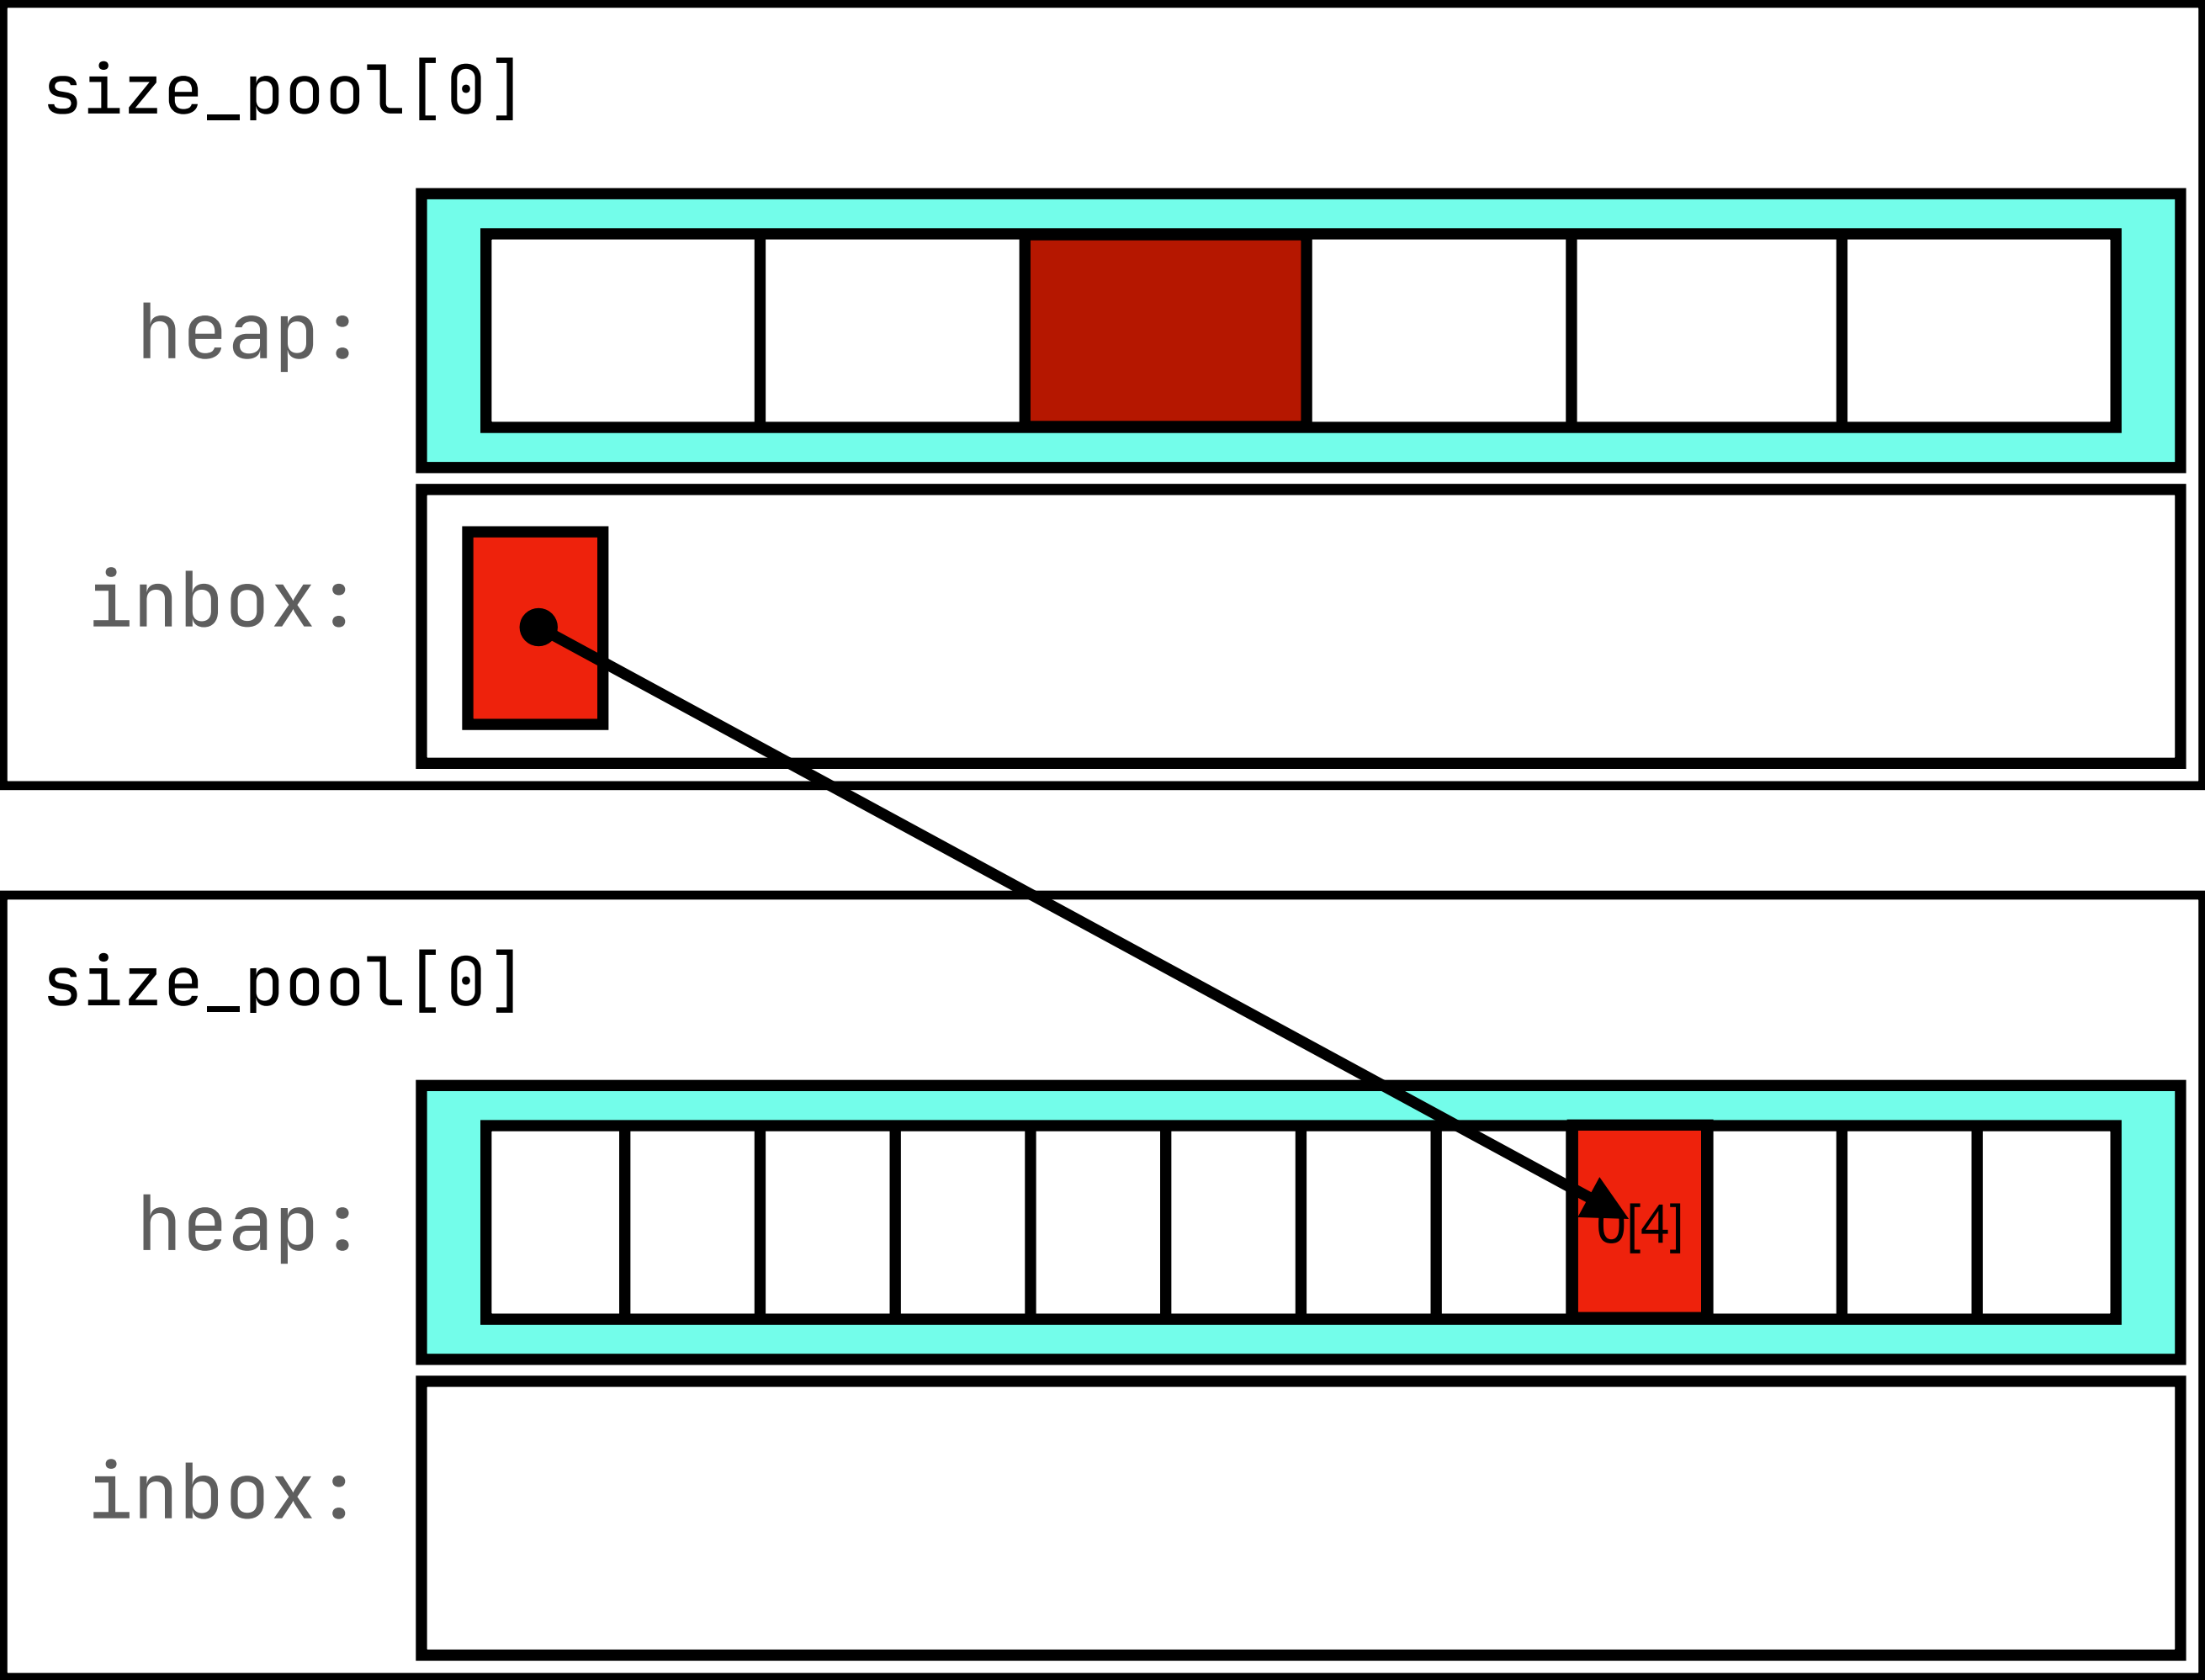
The Problem - Attempt 3

Size Pool Inboxes



The Problem - Attempt 3

Size Pool Inboxes



The Problem - Attempt 3

Size Pool Inboxes

Assign each size pool it's own separate data structure; an "inbox".

When an object is mutated, calculate the best size pool for it and push a pointer to it into the inbox for that pool.

When compacting the pool, first empty the inbox before moving from the compact cursor

✓ As fast as attempt 2

The Problem - Attempt 3

Size Pool Inboxes

Assign each size pool it's own separate data structure; an "inbox".

When an object is mutated, calculate the best size pool for it and push a pointer to it into the inbox for that pool.

When compacting the pool, first empty the inbox before moving from the compact cursor

- ✓ As fast as attempt 2
- ✓ Allows complete bi-directional object movement

The Problem - Attempt 3

Size Pool Inboxes

Assign each size pool it's own separate data structure; an "inbox".

When an object is mutated, calculate the best size pool for it and push a pointer to it into the inbox for that pool.

When compacting the pool, first empty the inbox before moving from the compact cursor

- ✓ As fast as attempt 2
- ✓ Allows complete bi-directional object movement
- ✗ Requires extra managed data structures

The Problem - Attempt 3

Size Pool Inboxes

Assign each size pool it's own separate data structure; an "inbox".

When an object is mutated, calculate the best size pool for it and push a pointer to it into the inbox for that pool.

When compacting the pool, first empty the inbox before moving from the compact cursor

- ✓ As fast as attempt 2
- ✓ Allows complete bi-directional object movement
- ✗ Requires extra managed data structures
- ✗ Leaks GC abstractions outside the GC

What now?

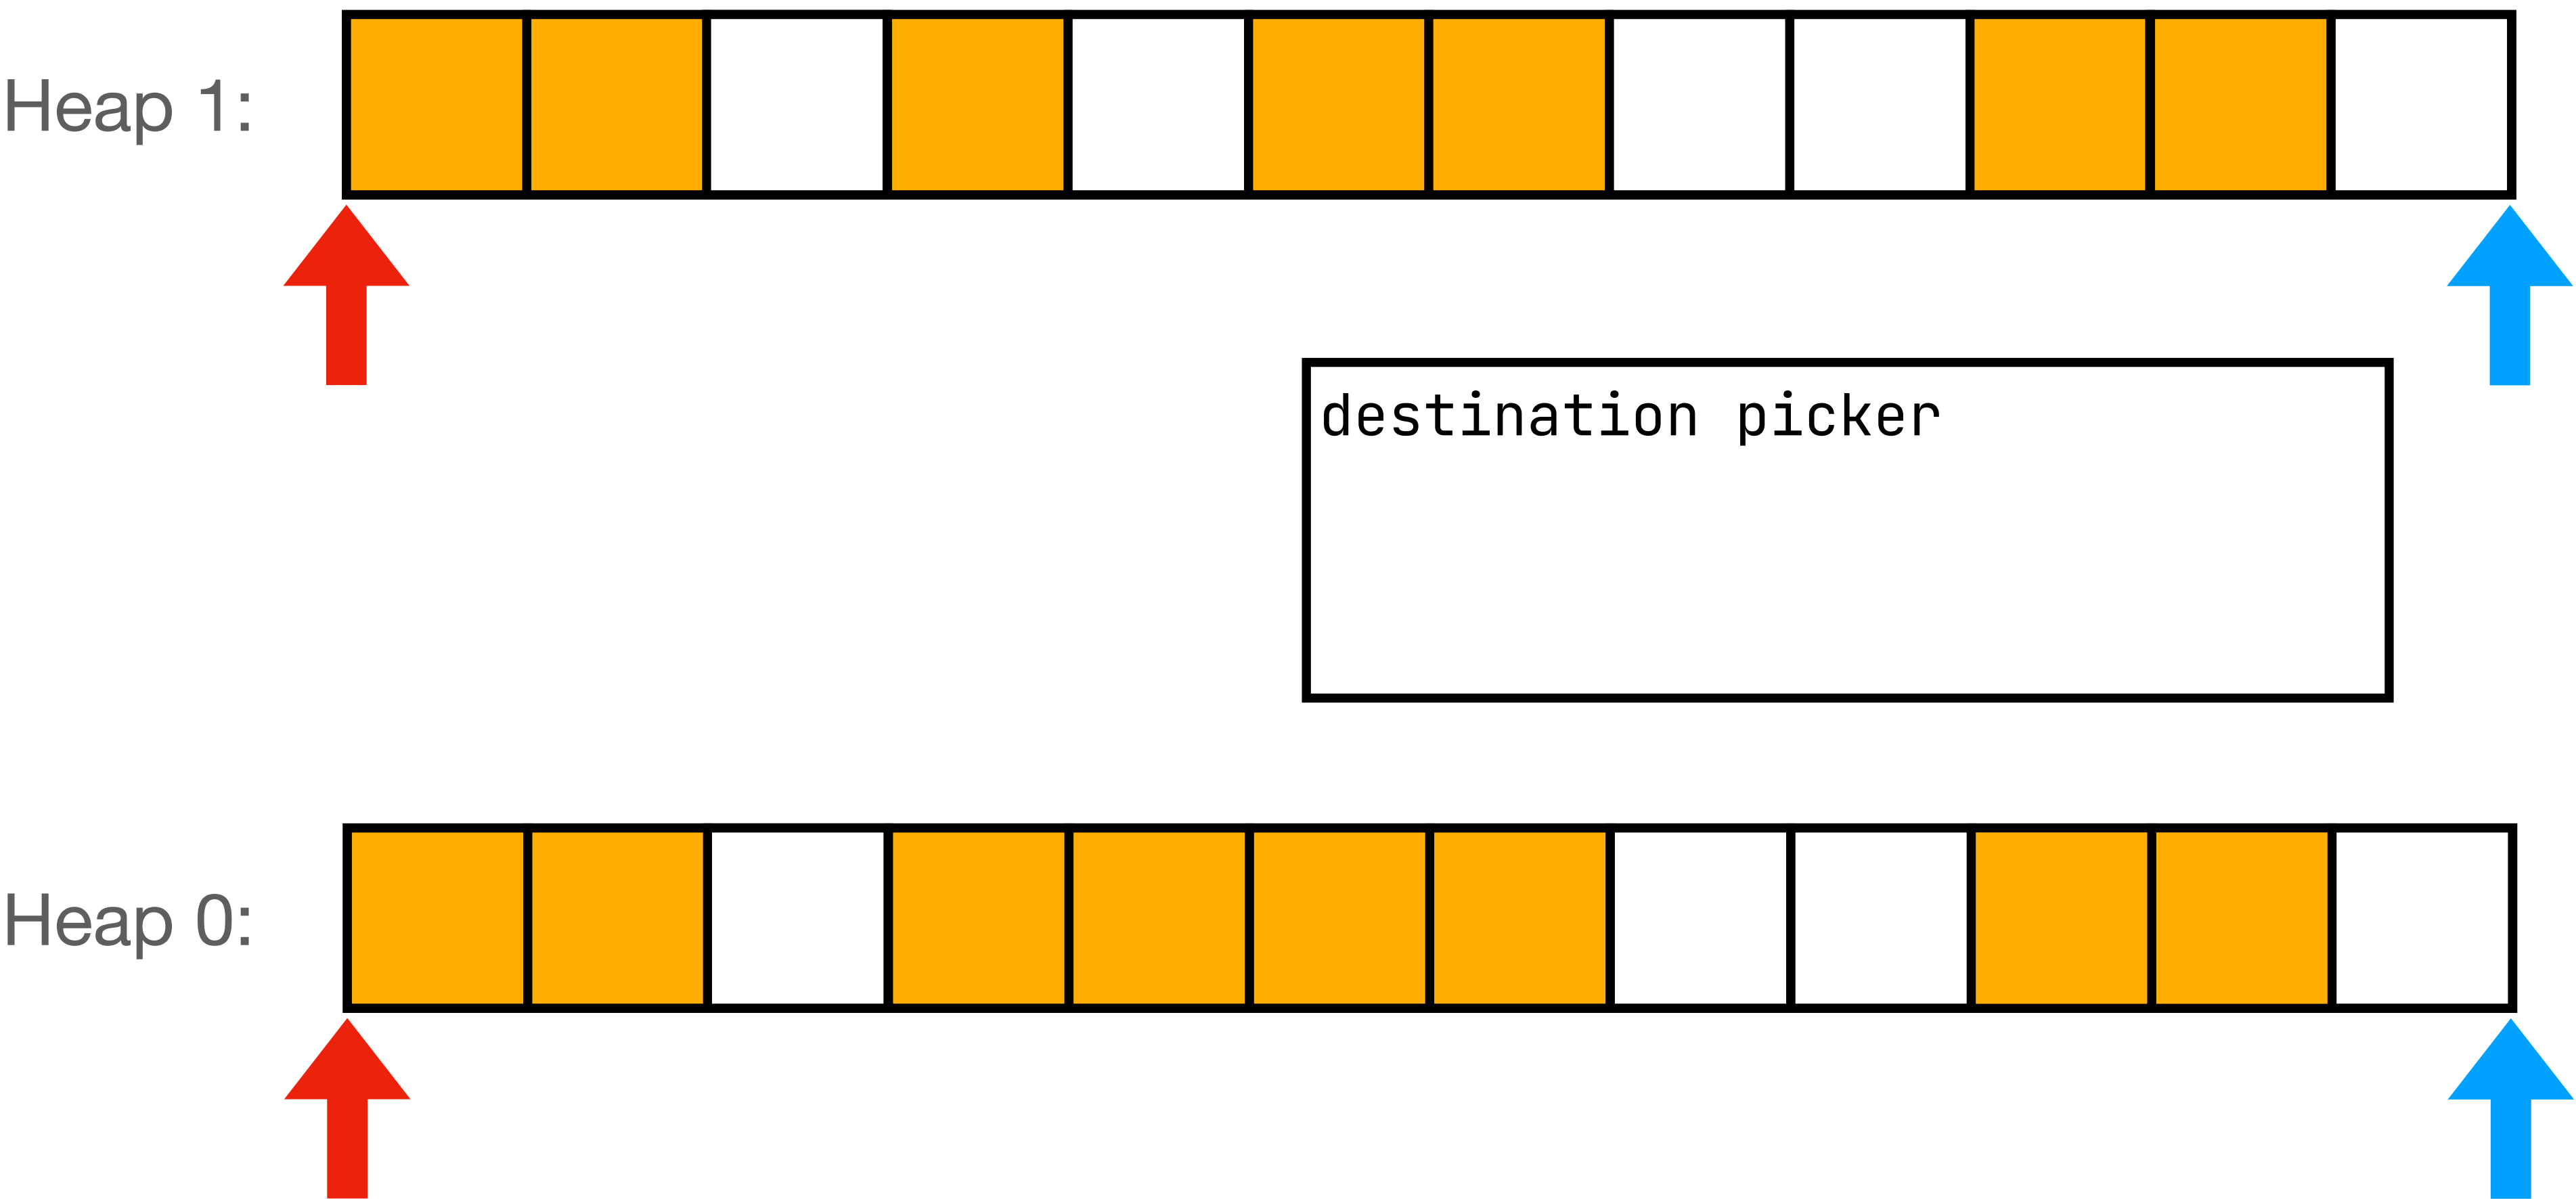


What now?

“What is the best place to put this object?”

The Problem - Attempt 4

Reversing the cursors



The Problem - Attempt 4

Seperate Object Movement from Reference Updating

Move the compact cursor down the heap until a live object is reached.

Determine the best size pool based on the object size.

Start sweeping the destination pool's heap until a free slot is found

Move the live object into the free slot

 Almost as fast as the existing compaction algorithm

The Problem - Attempt 4

Seperate Object Movement from Reference Updating

Move the compact cursor down the heap until a live object is reached.

Determine the best size pool based on the object size.

Start sweeping the destination pool's heap until a free slot is found

Move the live object into the free slot

- ✓ Almost as fast as the existing compaction algorithm
- ✓ negligible amount of extra complexity

The Problem - Attempt 4

Seperate Object Movement from Reference Updating

Move the compact cursor down the heap until a live object is reached.

Determine the best size pool based on the object size.

Start sweeping the destination pool's heap until a free slot is found

Move the live object into the free slot

- ✓ Almost as fast as the existing compaction algorithm
- ✓ negligible amount of extra complexity
- ✓ No extra data structures to manage

The Problem - Attempt 4

Seperate Object Movement from Reference Updating

Move the compact cursor down the heap until a live object is reached.

Determine the best size pool based on the object size.

Start sweeping the destination pool's heap until a free slot is found

Move the live object into the free slot

- ✓ Almost as fast as the existing compaction algorithm
- ✓ negligible amount of extra complexity
- ✓ No extra data structures to manage
- ✓ Natural bi-directional movement between pools

The Problem - Attempt 4

Home

My page

Projects

Help

Ruby »

Ruby master

Search:

Ruby master

Overview

Activity

Roadmap

Issues

New issue

Wiki

Repository

Feature #18619

CLOSED

Edit

Unwatch

Reverse the order of GC Compaction cursor movement

Added by

eightbitraptor (Matthew Valentine-House)

5 months ago

Updated

5 months ago

« Previous

|

4 of 7

|

Next »

Status:

Closed

Priority:

Normal

Assignee:

-

Target version:

-

[ruby-core:107818]

Description

Quote

Reverse the order of GC Compaction cursor movement

Github PR: <https://github.com/ruby/ruby/pull/5637>

Summary

The current compaction algorithm works by walking up the heap until it finds a slot to fill and then asking "what object can I use to fill this slot"

This PR reverses the cursor movement to walk down the heap until it gets to a moveable live object and then asking "where is the best place to put this object"

This animation shows a very simplified view of what we mean by "reversing the cursor movement"

Scan cursor

Compact cursor

Occupied slot

Empty slot

Current compaction

My custom queries

open bugs

Custom queries

Backport 2.2

Backport 2.3

Backport 2.4

Backport 2.5

Backport 2.6

Backport 2.7

Backport 3.0

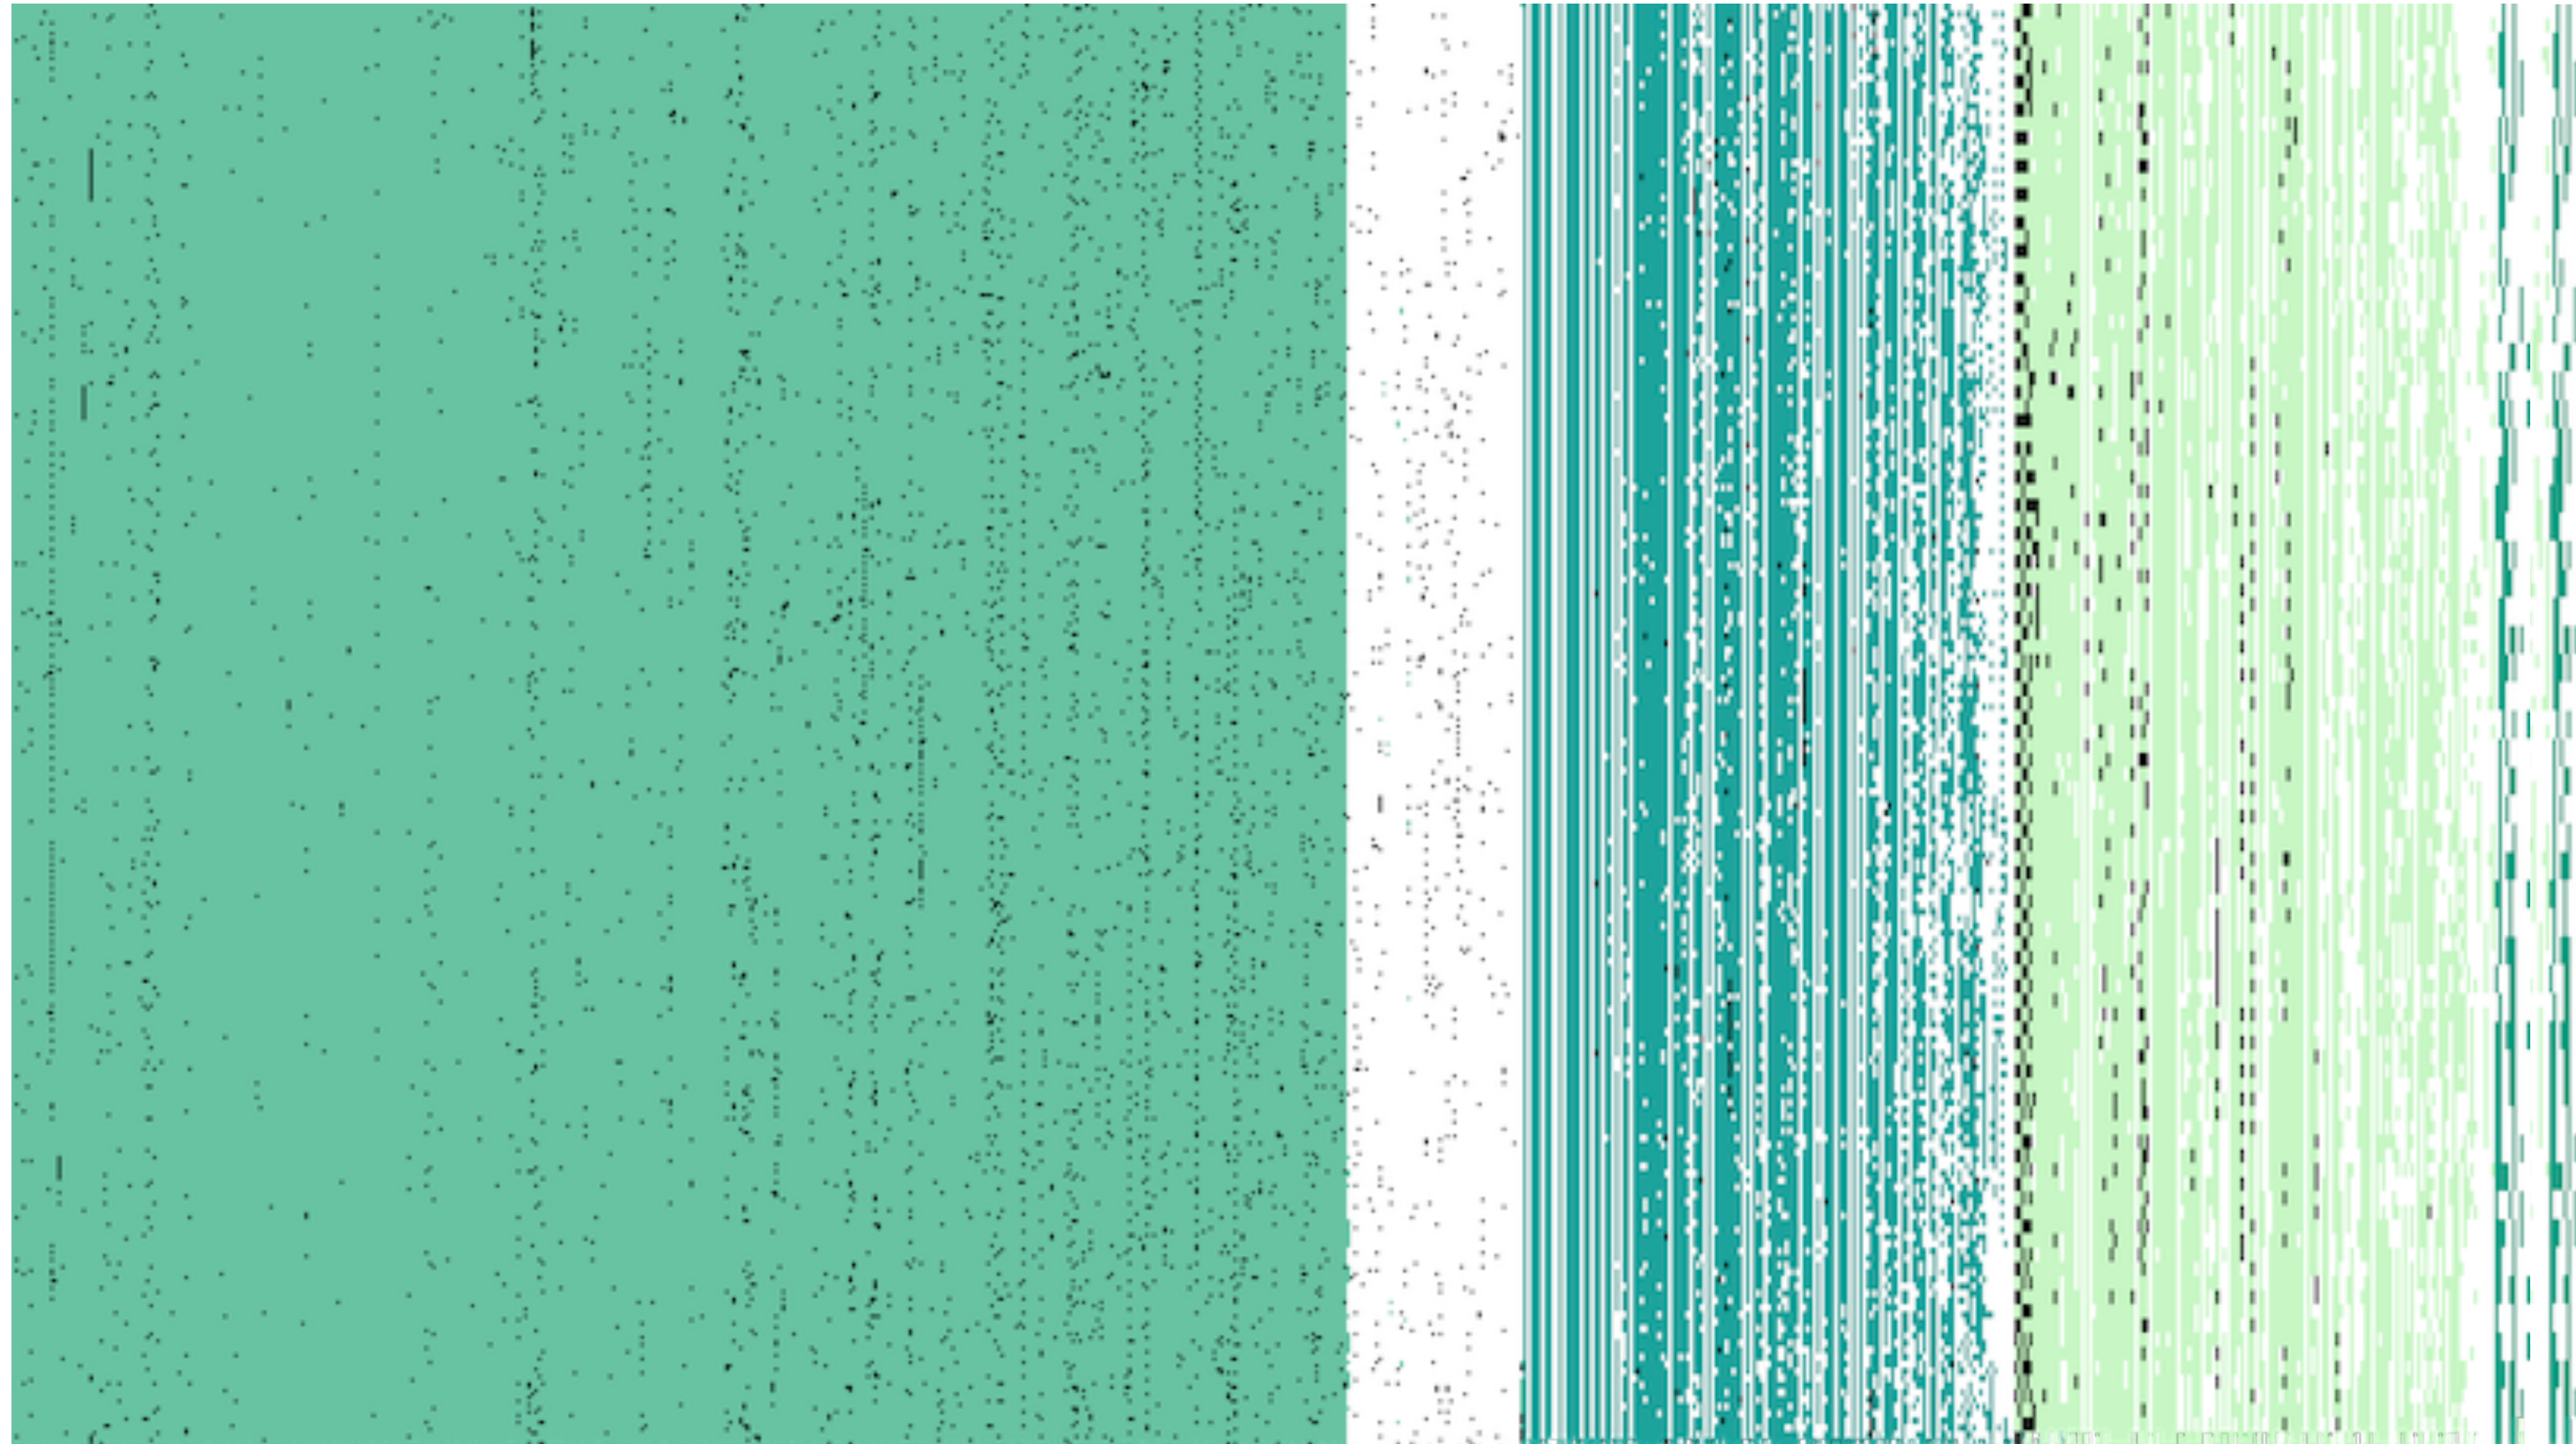
Backport 3.1

bugs: unassigned

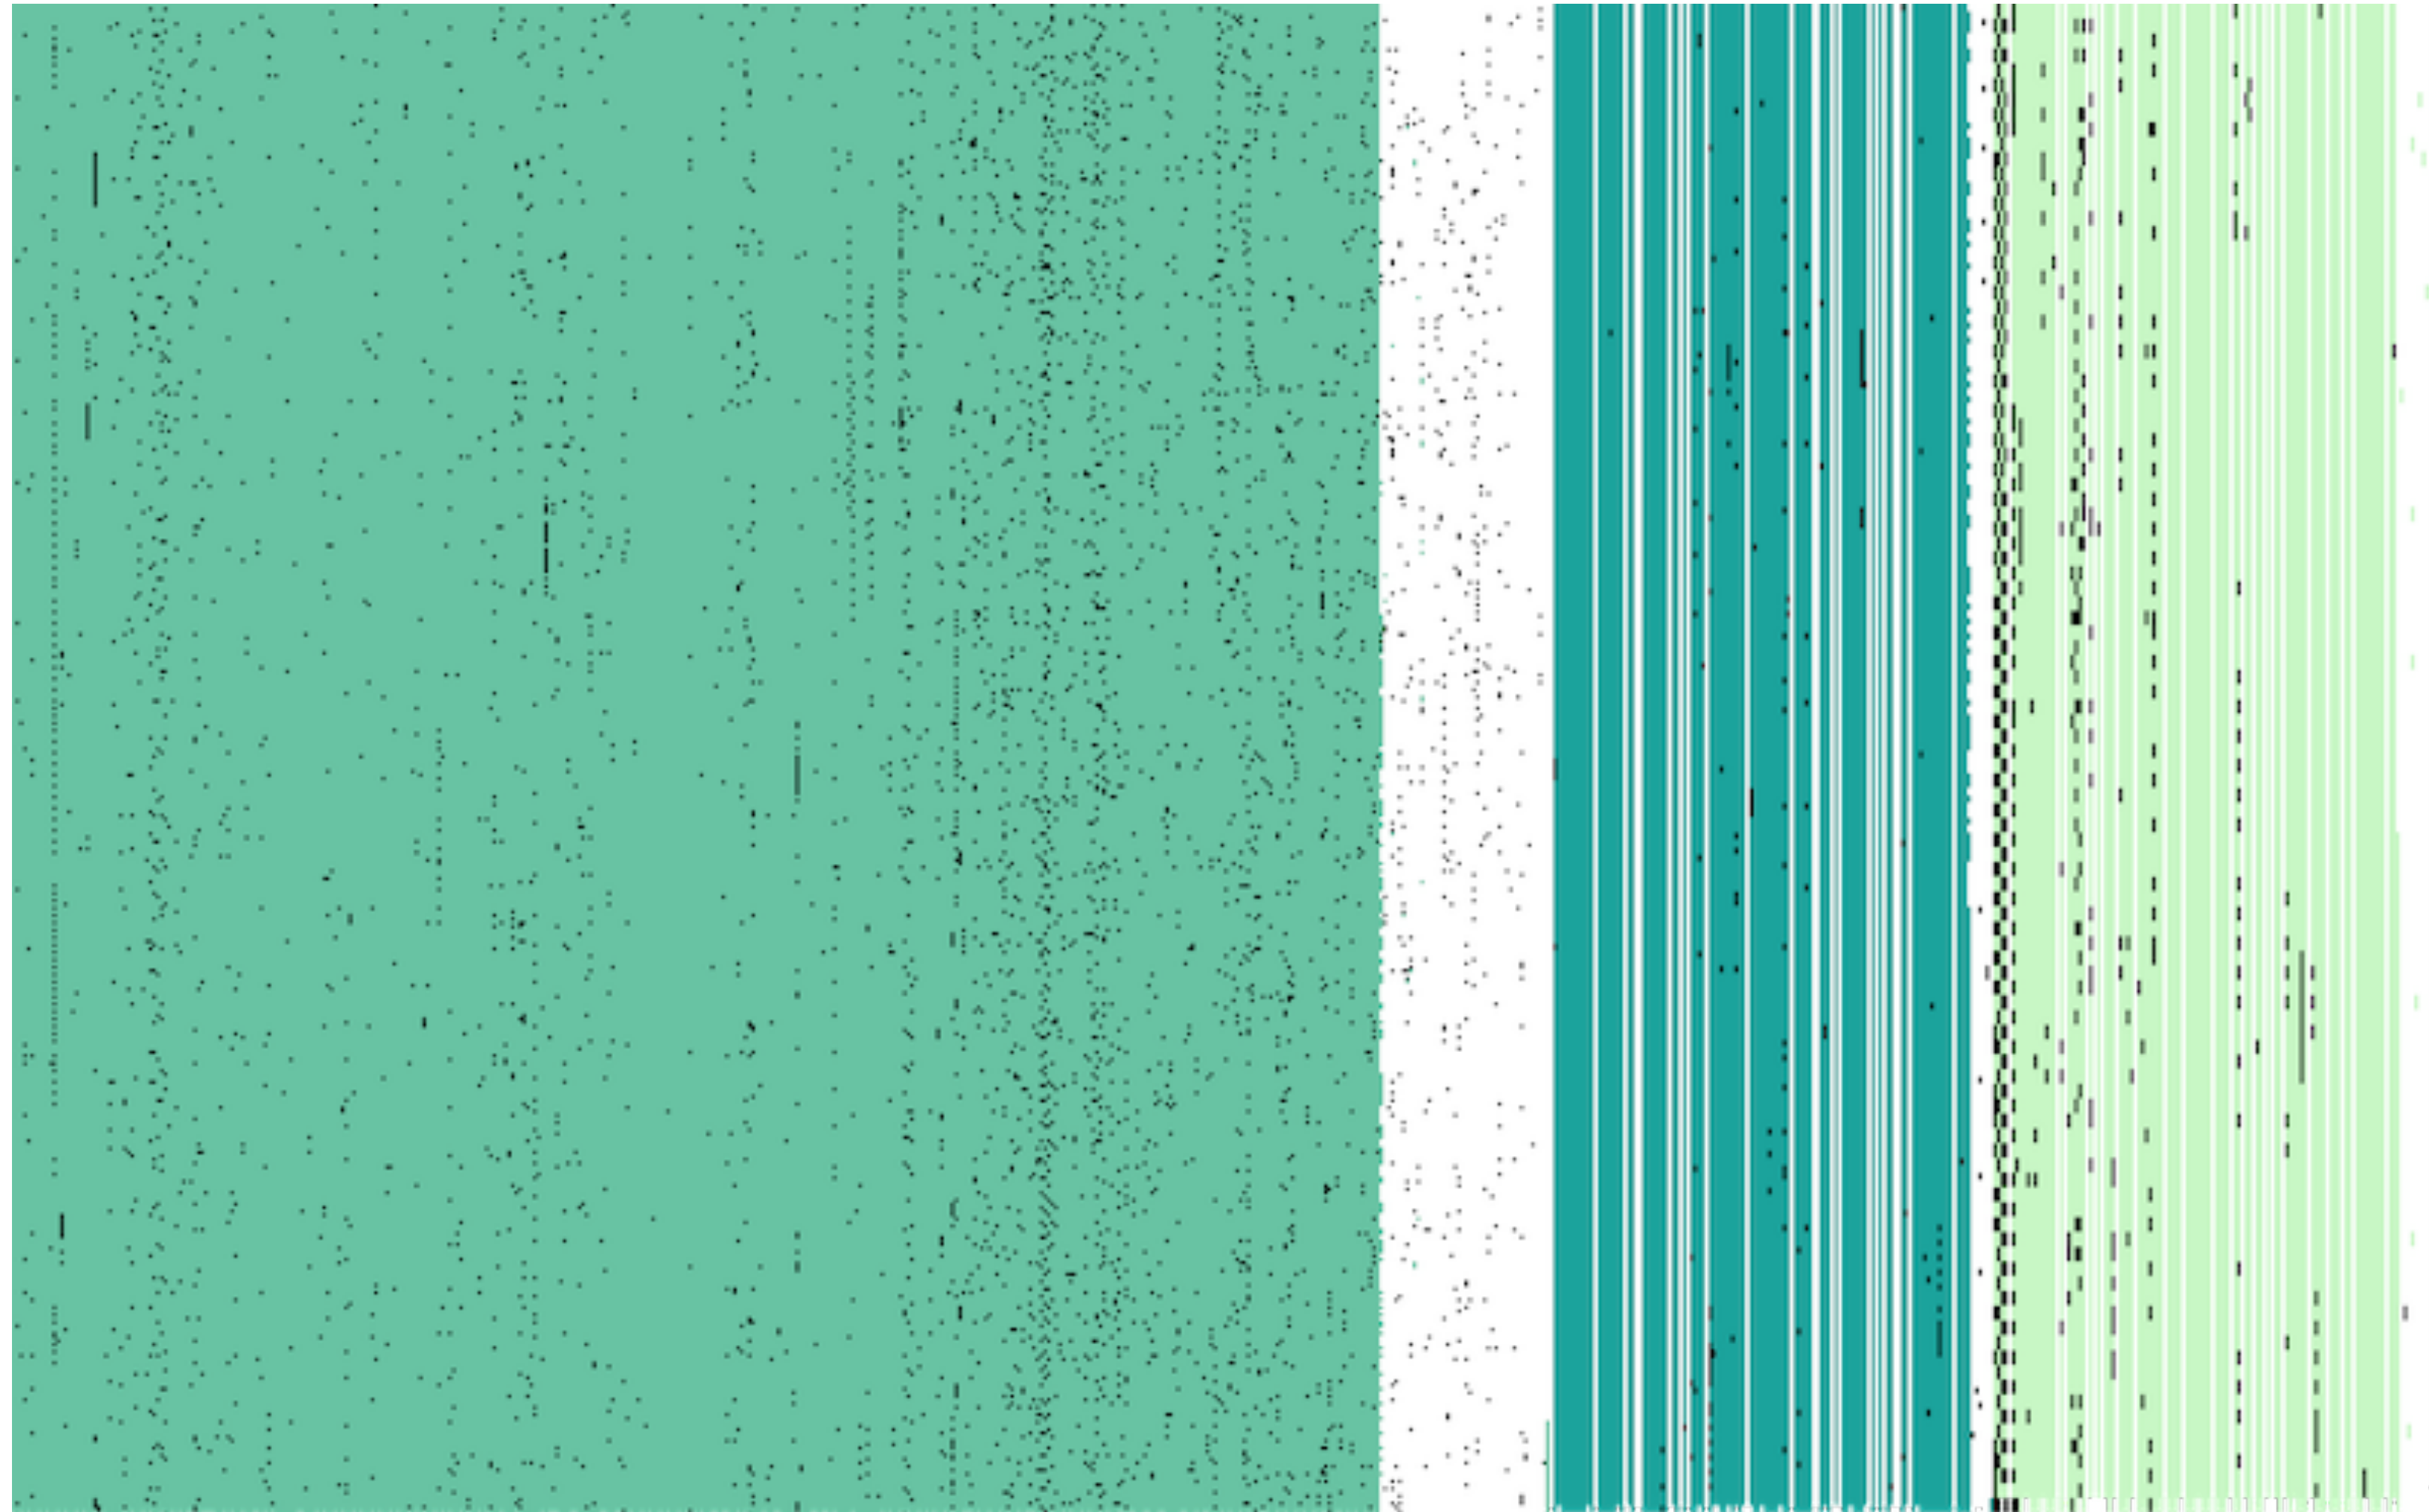
DevelopersMeeting

matz

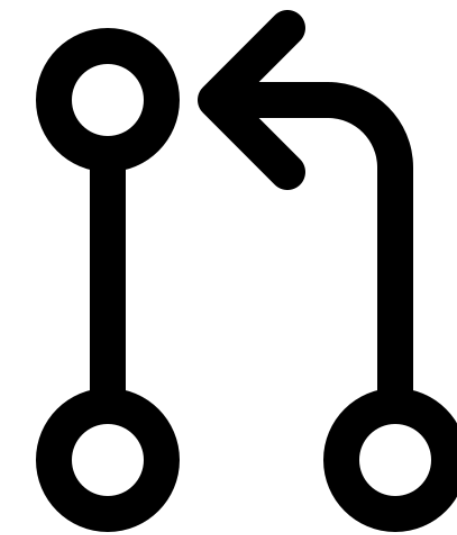
The Problem - Attempt 4



The Problem - Attempt 4



The Problem - Attempt 4



What's next?

Object Movement & Variable Width Allocation Status

String Movement

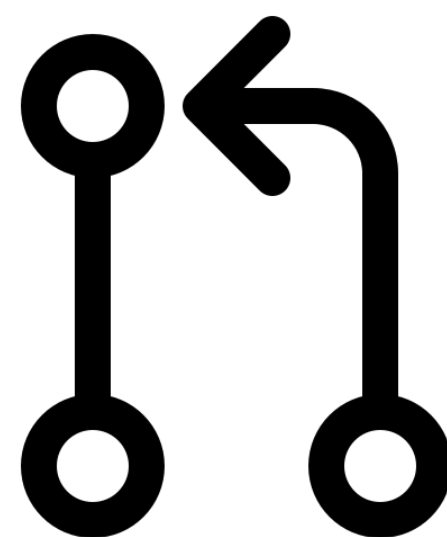
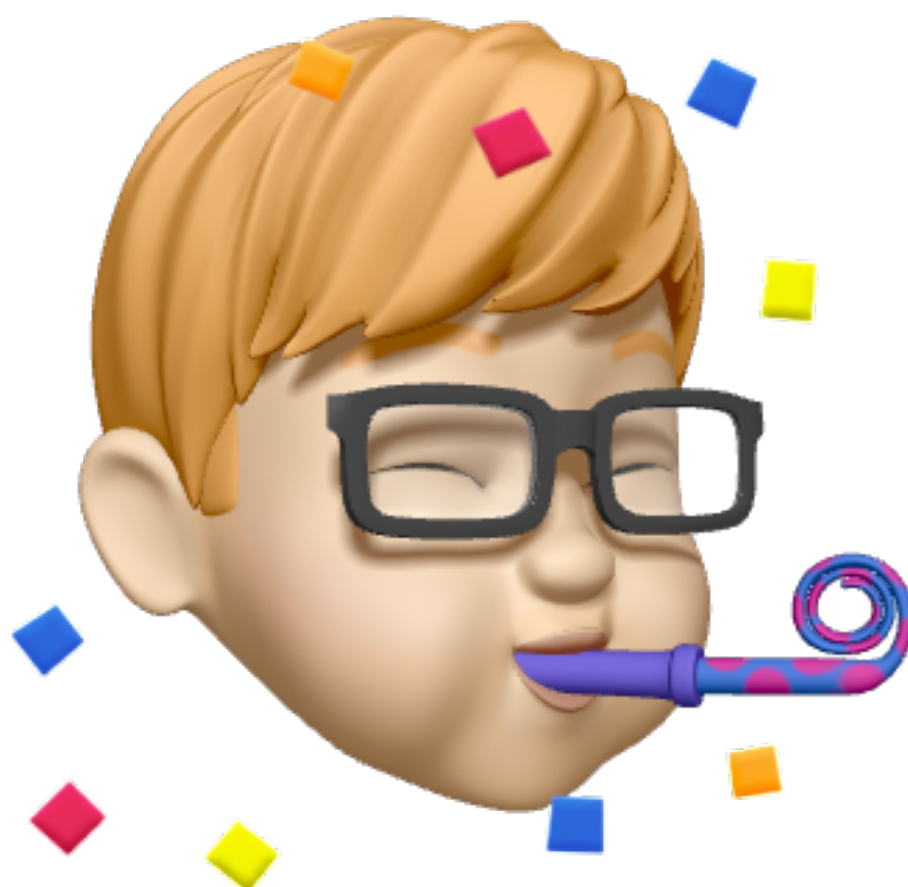
```
static rb_size_pool_t *
gc_compact_destination_pool(rb_objspace_t *objspace, rb_size_pool_t *src_pool, VALUE src)
{
    size_t obj_size;

    switch (BUILTIN_TYPE(src)) {
        case T_STRING:
            obj_size = rb_str_size_as_embedded(src);
            break;

        default:
            return src_pool;
    }

    if (rb_gc_size_allocatable_p(obj_size)){
        return &size_pools[size_pool_idx_for_size(obj_size)];
    }
    else {
        return &size_pools[0];
    }
}
```

Moving Strings



VWA Update

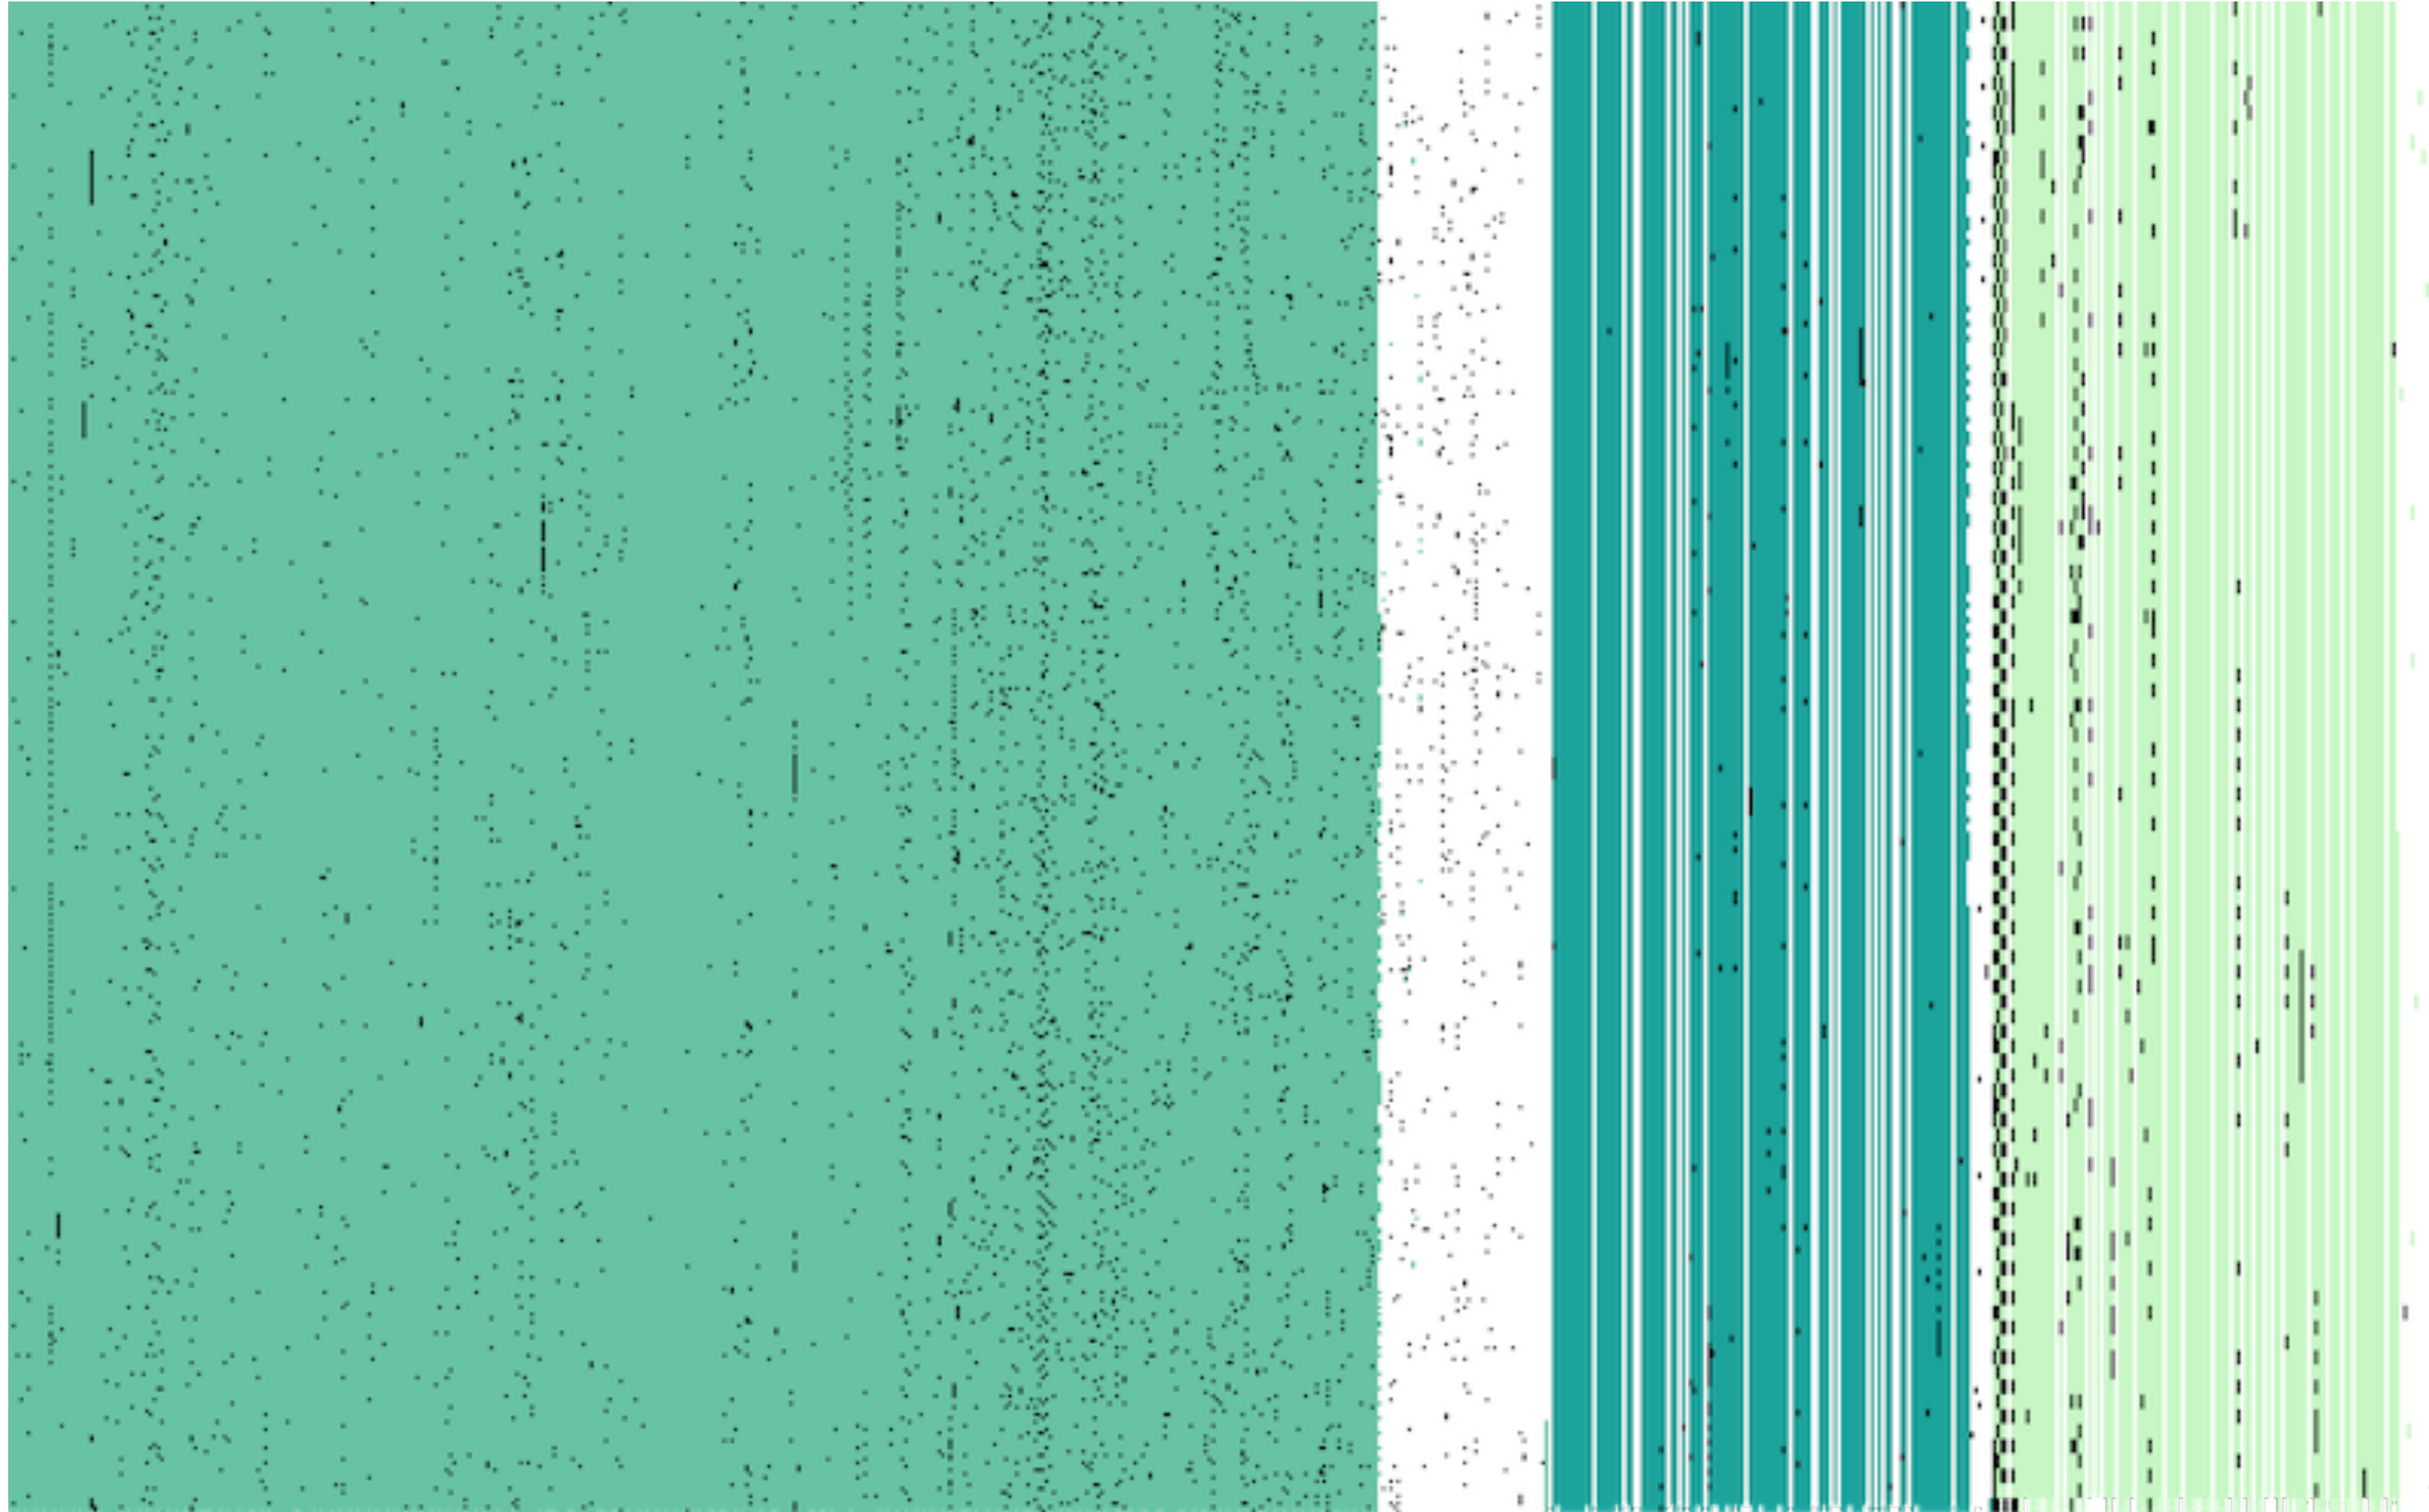
VWA is currently implemented for:

- **Classes**
- **Strings**
- **Arrays**
- **Objects**

Heapviz

A Ruby Gem for Heap Visualisation

Heapviz

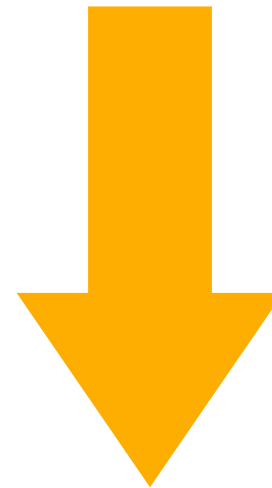


Heapviz



Visualizing your Ruby Heap

- Aaron Patterson, 2017, tenderlovemaking.com



```
> gem install heapviz
```

<https://github.com/eightbitraptor/heapviz>

Thanks RubyKaigi